



National Cheng Kung University

A large, lush green tree with a thick trunk and dense foliage, standing on a grassy area.

Chapter 24

Single-Source Shortest Paths

Sun-Yuan Hsieh

謝孫源 教授

成功大學資訊工程學系



Shortest paths

- ▶ How to find the shortest route between two points on a map.

▷ Input:

- Directed graph $G = (V, E)$
- Weight function $w : E \rightarrow \mathbf{R}$

▷ **Weight of path** $p = \langle v_0, v_1, \dots, v_k \rangle$
$$= \sum_{i=1}^k w(v_{i-1}, v_i)$$

= sum of edge weights on path p .





► **Shortest-path weight** u to v :

► $\delta(u, v) = \min\{w(p) : u \rightsquigarrow^p v\}$, if there exists a path $u \rightsquigarrow v$,
 otherwise.

► Shortest path u to v is any path p such that $w(p) = \delta(u, v)$.

► **Example:** shortest paths from s

[d values appear inside vertices. Shaded edges show shortest paths.]



- This example shows that the shortest path might not be unique.
- It also shows that when we look at shortest paths from one vertex to all other vertices, the shortest paths are organized as tree.

- ▶ Can think of weights as representing any measure that
 - ▷ accumulates linearly along a path.
 - ▷ we want to minimize.

- ▶ Example: time, cost, penalties, loss.
 - ▷ Generalization of breadth-first search to weighted graphs.



Variants

► *Single-source*

- ▷ Find shortest paths from a given source vertex $s \in V$ to every vertex $v \in V$.

► *Single-destination*

- ▷ Find shortest paths to a given destination vertex.

► *Single-pair*

- ▷ Find shortest path from u to v . No way known that's better in worst case than solving single-source.

► *All-pairs*

- ▷ Find shortest path from u to v for all $u, v \in V$.
- ▷ We'll see algorithms for all-pairs in the next chapter.



Negative-weight edges

- ▶ OK, as long as no negative-weight cycles are reachable from the source.
 - ▷ If we have a negative-weight cycle, we can just keep going around it, and get $w(s, v) = -\infty$ for all v on the cycle.
 - ▷ But OK if the negative-weight cycle is not reachable from the source.
 - ▷ Some algorithms work only if there are no negative-weight edges in the graph.

We'll be clear when they're allowed and not allowed.




Optimal substructure

► **Lemma**

Any subpath of a shortest path is a shortest path.

Proof: Cut - and - paste.




Suppose this path p is a shortest path from u to v .

Then $\delta(u, v) = w(p) = w(p_{ux}) + w(p_{xy}) + w(p_{yv})$.



Now suppose there exists a shorter path $x \overset{p'_{xy}}{\rightsquigarrow} y$.

Then $w(p'_{xy}) < w(p_{xy})$



$$\begin{aligned}
 \text{Then } w(p') &= w(p_{ux}) + w(p'_{xy}) + w(p_{yv}) \\
 &< w(p_{ux}) + w(p_{xy}) + w(p_{yv}) \\
 &= w(p)
 \end{aligned}$$

So p wasn't shortest path after all !

■ (lemma)



- ▶ Shortest paths can't contain cycles :
 - ▷ Already ruled out negative-weight cycles.
 - ▷ Positive-weight \Rightarrow we can get a shorter path by omitting the cycle.
 - ▷ Zero-weight: no reason to use them \Rightarrow assume that our solutions won't use them.



Output of single-source shortest-path algorithm

- ▶ For each vertex $v \in V$:
 - ▷ $d[v] = \delta(s, v)$.
 - Initially, $d[v] = \infty$.
 - Reduces as algorithms progress. But always maintain $d[v] \geq \delta(s, v)$.
 - Call $d[v]$ a *shortest-path estimate*.
 - ▷ $\pi[v] =$ predecessor of v on a shortest path from s .
 - If no predecessor, $\pi[v] = \text{NIL}$.
 - π induces a tree ----- *shortest-path tree*.
 - We won't prove properties of π in lecture ----- see text.



Initialization

- ▶ All the shortest-paths algorithms start with INIT-SINGLE-SOURCE.

▷ INIT-SINGLE-SOURCE (V, s)

For each $v \in V$

do $d[v] \leftarrow \infty$

$\pi[v] \leftarrow \text{NIL}$

$d[s] \leftarrow 0$



Relaxing an edge (u, v)



- ▶ Can we improve the shortest-path estimate for v by going through u and taking (u, v) ?

▷ RELAX (u, v, w)

If $d[v] > d[u] + w(u, v)$

then $d[v] \leftarrow d[u] + w(u, v)$

$\pi[v] \leftarrow u$





- ▶ For all the single-source shortest-paths algorithms we'll look at,
 - ▷ start by calling INIT-SINGLE-SOURCE,
 - ▷ then relax edges.

- ▶ The algorithms differ in the order and how many times they relax each edge.




Shortest-paths properties

- ▶ Based on calling INIT-SINGLE-SOURCE once and then calling RELAX zero or more times.
- ▶ **Triangle inequality**
 - ▷ For all $(u, v) \in E$, we have $\delta(s, v) \leq \delta(s, u) + w(u, v)$.

Proof:

Weight of shortest path $s \rightsquigarrow v$ is \leq weight of any path $s \rightsquigarrow v$.

Path $s \rightsquigarrow u \rightarrow v$ is a path $s \rightsquigarrow v$, and if we use a shortest path $s \rightsquigarrow u$, its weight is $\delta(s, u) + w(u, v)$. ■





Upper-bound property

- ▶ Always have $d[v] \geq \delta(s, v)$ for all v .

Once $d[v] = \delta(s, v)$, it never changes.

Proof: Initially true.

Suppose there exists a vertex such that $d[v] < \delta(s, v)$.

Without loss of generality, v is first vertex for which this happens.

Let u be the vertex that causes $d[v]$ to change.

Then $d[v] = d[u] + w(u, v)$.


So, $d[v] < \delta(s, v)$

$$\leq \delta(s, u) + w(u, v) \text{ (triangle inequality)}$$

$$\leq d[u] + w(u, v) \text{ (} v \text{ is first violation)}$$

$\Rightarrow d[v] < d[u] + w(u, v)$. (Contradicts $d[v] = d[u] + w(u, v)$)

Once $d[v]$ reaches $\delta(s, v)$, it never goes lower. It never goes up, since relaxations only lower shortest-path estimates. ■



No-path property

- ▶ If $\delta(s, v) = \infty$, then $d[v] = \infty$ always.

Proof:

$$d[v] \geq \delta(s, v) = \infty \Rightarrow d[v] = \infty.$$





Convergence property

- ▶ If $s \rightsquigarrow u \rightarrow v$ is a shortest path, $d[u] = \delta(s, u)$, and we call $\text{RELAX}(u, v, w)$, then $d[v] = \delta(s, v)$ afterward.

Proof:

After relaxation:

$$d[v] \leq d[u] + w(u, v) \quad (\text{RELAX code})$$

$$= \delta(s, u) + w(u, v)$$

$$= \delta(s, v) \quad (\text{lemma ----- optimal substructure})$$

Since $d[v] \geq \delta(s, v)$, must have $d[v] = \delta(s, v)$. ■



Path relaxation property

- ▶ Let $p = \langle v_0, v_1, \dots, v_k \rangle$ be a shortest path from $s = v_0$ to v_k . If we relax, *in order*, $(v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k)$, even intermixed with other relaxations, then $d[v_k] = \delta(s, v_k)$.

Proof:

Induction to show that $d[v_i] = \delta(s, v_i)$ after (v_{i-1}, v_i) is relaxed.

Basis:


$i = 0$. Initially, $d[v_0] = 0 = \delta(s, v_0) = \delta(s, s)$.

Inductive step:

Assume $d[v_{i-1}] = \delta(s, v_{i-1})$.

Relax (v_{i-1}, v_i) .

By Convergence Property, $d[v_i] = \delta(s, v_i)$ afterward and $d[v_i]$ never changes.



$$\blacktriangleright d(v_1) = \delta(s, v_0) + w(v_0, v_1) = \delta(s, v_1)$$

$$d(v_2) = \delta(s, v_1) + w(v_1, v_2) = \delta(s, v_2)$$

$$d(v_3) = \delta(s, v_2) + w(v_2, v_3) = \delta(s, v_3)$$

○

○

$$d(v_k) = \delta(s, v_{k-1}) + w(v_{k-1}, v_k) = \delta(s, v_k)$$



The Bellman-Ford algorithm

- ▶ Allows negative-weight edges.
- ▶ Computes $d[v]$ and $\pi[v]$ for all $v \in V$.
- ▶ Returns TRUE if no negative-weight cycles reachable from s , FALSE otherwise.

► BELLMAN-FORD (V, E, w, s)
INIT-SINGLE-SOURCE (V, s)
for $i \leftarrow 1$ to $|V| - 1$
 do for each edge $(u, v) \in E$
 do RELAX (u, v, w)
for each edge $(u, v) \in E$
 do if $d[v] > d[u] + w(u, v)$
 then return FALSE
return TRUE

Core: The first **for** loop relaxes all edges $|V| - 1$ times.

Time: $\Theta(VE)$.

► $i = 1$




Figure 24.4 The execution of the Bellman-Ford algorithm. The source is vertex s . The d values are shown within the vertices, and shaded edges indicate predecessor values: if edge (u, v) is shaded, then $\pi[v] = u$. In this particular example, each pass relaxes the edges in the order (t, x) , (t, y) , (t, z) , (x, t) , (y, x) , (y, z) , (z, x) , (z, s) , (s, t) , (s, y) . (a) The situation just before the first pass over the edges. (b)–(e) The situation after each successive pass over the edges. The d and π values in part (e) are the final values. The Bellman-Ford algorithm returns TRUE in this example.

► $i = 2$




Figure 24.4 The execution of the Bellman-Ford algorithm. The source is vertex s . The d values are shown within the vertices, and shaded edges indicate predecessor values: if edge (u, v) is shaded, then $\pi[v] = u$. In this particular example, each pass relaxes the edges in the order (t, x) , (t, y) , (t, z) , (x, t) , (y, x) , (y, z) , (z, x) , (z, s) , (s, t) , (s, y) . (a) The situation just before the first pass over the edges. (b)–(e) The situation after each successive pass over the edges. The d and π values in part (e) are the final values. The Bellman-Ford algorithm returns TRUE in this example.

► $i = 3$




Figure 24.4 The execution of the Bellman-Ford algorithm. The source is vertex s . The d values are shown within the vertices, and shaded edges indicate predecessor values: if edge (u, v) is shaded, then $\pi[v] = u$. In this particular example, each pass relaxes the edges in the order (t, x) , (t, y) , (t, z) , (x, t) , (y, x) , (y, z) , (z, x) , (z, s) , (s, t) , (s, y) . (a) The situation just before the first pass over the edges. (b)–(e) The situation after each successive pass over the edges. The d and π values in part (e) are the final values. The Bellman-Ford algorithm returns TRUE in this example.

► $i = 4$





Figure 24.4 The execution of the Bellman-Ford algorithm. The source is vertex s . The d values are shown within the vertices, and shaded edges indicate predecessor values: if edge (u, v) is shaded, then $\pi[v] = u$. In this particular example, each pass relaxes the edges in the order (t, x) , (t, y) , (t, z) , (x, t) , (y, x) , (y, z) , (z, x) , (z, s) , (s, t) , (s, y) . (a) The situation just before the first pass over the edges. (b)–(e) The situation after each successive pass over the edges. The d and π values in part (e) are the final values. The Bellman-Ford algorithm returns TRUE in this example.

► **Example:**



- ▷ Values you get on each pass and how quickly it converges depends on order of relaxation.
- ▷ But guaranteed to converge after $|V| - 1$ passes, assuming no negative-weight cycles.



► **Proof:** Use path-relaxation property.

Let v be reachable from s , and let $p = \langle v_0, v_1, \dots, v_k \rangle$ be a shortest path from s to v , where $v_0 = s$ and $v_k = v$.

Since p is acyclic, it has $\leq |V| - 1$ edges, so $k \leq |V| - 1$.

Each iteration of the **for** loop relaxes all edges:

- First iteration relaxes (v_0, v_1) .
- Second iteration relaxes (v_1, v_2) .
- k th iteration relaxes (v_{k-1}, v_k) .

By the path-relaxation property, $d[v] = d[v_k] = \delta(s, v_k) = \delta(s, v)$.





- ▶ How about the TRUE/FALSE return value?
 - ▷ Suppose there is no negative-weight cycle reachable from s .
At termination, for all $(u, v) \in E$,
 $d[v] = \delta(s, v)$
 $\leq \delta(s, u) + w(u, v)$ (triangle inequality)
 $= d[u] + w(u, v)$
So BELLMAN-FORD returns TRUE.

▶ ▶ Now suppose there exists negative-weight cycle
 $c = \langle v_0, v_1, \dots, v_k \rangle$ where $v_0 = v_k$, reachable from s .

Then $\sum_{i=1}^k w(v_{i-1}, v_i) < 0$

Suppose (for contradiction) that BELLMAN-FORD returns TRUE.

Then $d[v_i] \leq d[v_{i-1}] + w(v_{i-1}, v_i)$ for $i = 1, 2, \dots, k$.

sum around c :

$$\begin{aligned} \sum_{i=1}^k d[v_i] &\leq \sum_{i=1}^k (d[v_{i-1}] + w(v_{i-1}, v_i)) \\ &= \sum_{i=1}^k d[v_{i-1}] + \sum_{i=1}^k w(v_{i-1}, v_i) \end{aligned}$$



Each vertex appears once in each summation $\sum_{i=1}^k d[v_i]$
and $\sum_{i=1}^k d[v_{i-1}]$.

$$\Rightarrow 0 \sum_{i=1}^k w(v_{i-1}, v_i).$$

This contradicts C being a negative-weight cycle! ■

Single-source shortest paths in a directed acyclic graph



成功大學

COPYRIGHT 2002 NATIONAL CHENG KUNG UNIVERSITY



- ▶ Since a dag, we're guaranteed no negative-weight cycles.

DAG-SHORTEST-PATHS (V, E, w, s)

topologically sort the vertices


INIT-SINGLE-SOURCE (V, s)

for each vertex u , taken in topologically sorted order

do for each vertex $v \in Adj[u]$

do RELAX (u, v, w)

► *Example:*



► *Time:* $\Theta(V + E)$.




Figure 24.5 The execution of the algorithm for shortest paths in a directed acyclic graph. The vertices are topologically sorted from left to right. The source vertex is s . The d values are shown within the vertices, and shaded edges indicate the π values.



- ▶ **Correctness:** Because we process vertices in topologically sorted order, edges of any path must be relaxed in order of appearance in the path.
 - ⇒ Edges on any shortest path are relaxed in order.
 - ⇒ By path-relaxation property, correct. ■





Dijkstra's algorithm

- ▶ No negative-weight *edges*.
- ▶ Essentially a weighted version of breadth-first search.
 - ▷ Instead of a **FIFO** queue, uses a priority queue.
 - ▷ Keys are shortest-path weights ($d[v]$).
- ▶ Have two sets of vertices:
 - ▷ S = vertices whose final shortest-path weights are determined.
 - ▷ Q = priority queue = $V - S$.


```
DIJKSTRA ( $V, E, w, s$ )  
INIT-SINGLE-SOURCE ( $V, s$ )  
 $S \leftarrow \emptyset$   
 $Q \leftarrow V$  // i.e., insert all vertices into  $Q$   
While  $Q \neq \emptyset$   
    do  $u \leftarrow \text{EXTRACT-MIN}(Q)$   
         $S \leftarrow S \cup \{u\}$   
        for each vertex  $v \in \text{Adj}[u]$   
            do RELAX ( $u, v, w$ )
```

- ▶ Like Prim's algorithm, but computing $d[v]$, and using shortest-path weights as keys.
- ▶ Dijkstra's algorithm can be viewed as greedy, since it always chooses the "lightest" ("closest"?) vertex in $V - S$ to add to S .





► *Example:*



Order of adding to S : s, y, z, x .

► **Correctness:**

- ▷ **Loop invariant** : At the start of each iteration of the **while** loop, $d[v] = \delta(s, v)$ for all $v \in S$.
- ▷ **Initialization**: Initially, $S = \emptyset$, so trivially true.
- ▷ **Termination**: At end, $Q = \emptyset \Rightarrow S = V \Rightarrow d[v] = \delta(s, v)$ for all $v \in V$.

► **Maintenance:**

Need to show that $d[u] = \delta(s, u)$ when u is added to S in each iteration.

Suppose there exists u such that $d[u] \neq \delta(s, u)$. Without loss of generality, let u be the first vertex for which $d[u] \neq \delta(s, u)$ when u is added to S .

Observations:

- ▷ $u \neq s$, since $d[s] = \delta(s, s) = 0$.
- ▷ Therefore, $s \in S$, so $S \neq \emptyset$.
- ▷ There must be some path $s \rightsquigarrow u$,
since otherwise $d[u] = \delta(s, u) = \infty$ by no-path property.




So, there's a path $s \rightsquigarrow u$.

This means there's a shortest path $s \overset{p}{\rightsquigarrow} u$.

Just before u is added to S , path p connects a vertex in S (i.e., s) to a vertex in $V - S$ (i.e., u).

Let y be first vertex along p that's in $V - S$, and let $x \in S$ be y 's predecessor.



Decompose p into $s \overset{p_1}{\rightsquigarrow} x \rightarrow y \overset{p_2}{\rightsquigarrow} u$.

(Could have $x = s$ or $y = u$, so that p_1 or p_2 may have no edges.)



► ***Claim***

$d[y] = \delta(s, y)$ when u is added to S .

Proof

$x \in S$ and u is the first vertex such that $d[u] \neq \delta(s, u)$ when u is added to $S \Rightarrow d[x] = \delta(s, x)$ when x is added to S .

Relaxed (x, y) at that time, so by the convergence property,

$d[y] = \delta(s, y)$. ■ (claim)



Now can get a contradiction to $d[u] \neq \delta(s, u)$:

y is on shortest path $s \rightsquigarrow u$, and all edge weights are nonnegative

$$\Rightarrow \delta(s, y) \leq \delta(s, u)$$

$$\Rightarrow d[y] = \delta(s, y)$$

$$\leq \delta(s, u)$$

$$\leq d[u] \quad (\text{upper-bound property}).$$

Also, both y and u were in Q when we chose u , so $d[u] \leq d[y]$

$$\Rightarrow d[u] = d[y].$$

Therefore, $d[y] = \delta(s, y) = \delta(s, u) = d[u]$.

Contradicts assumption that $d[u] \neq \delta(s, u)$.

Hence, Dijkstra's algorithm is correct. ■



- ▶ **Analysis:** Like Prim's algorithm, depends on implementation of priority queue.
 - ▷ If binary heap, each operation takes $O(\lg V)$ time
 $\Rightarrow O(E \lg V)$.
 - ▷ If a Fibonacci heap:
 - Each EXTRACT-MIN takes $O(1)$ amortized time.
 - There are $O(V)$ other operations, taking $O(\lg V)$ amortized time each.
 - Therefore, time is $O(V \lg V + E)$.



Difference constraints

- ▶ Given a set of inequalities of the form $x_j - x_i \leq b_k$.
 - ▷ x 's are variables, $1 \leq i, j \leq n$,
 - ▷ b 's are constants, $1 \leq k \leq m$.

Want to find a set of values for the x 's that satisfy all m inequalities, or determine that no such values exist.

Call such a set of values a ***feasible solution***.



► **Example:**


$$x_1 - x_2 \begin{matrix} \blacklozenge \\ ? \end{matrix} 5$$

$$x_1 - x_3 \begin{matrix} \blacklozenge \\ ? \end{matrix} 6$$

$$x_2 - x_4 \begin{matrix} \blacklozenge \\ ? \end{matrix} -1$$

$$x_3 - x_4 \begin{matrix} \blacklozenge \\ ? \end{matrix} -2$$

$$x_4 - x_1 \begin{matrix} \blacklozenge \\ ? \end{matrix} -3$$



Solution: $x = (0, -4, -5, -3)$

Also: $x = (5, 1, 0, 2) = [\text{above solution}] + 5$

► *Lemma*

If x is a feasible solution, then so is $x + d$ for any constant d .

Proof

x is a feasible solution

$$\Rightarrow x_j - x_i \leq b_k \text{ for all } i, j, k.$$


$$\Rightarrow (x_j + d) - (x_i + d) \leq b_k.$$

■ (lemma)

► Constraint graph

$G = (V, E)$, weighted, directed.

- ▷ $V = \langle v_0, v_1, v_2, \dots, v_n \rangle$: one vertex per variable + v_0
- ▷ $E = \{(v_i, v_j) : x_j - x_i \text{ ? } b_k \text{ is a constraint}\}$
 $\cup \{(v_0, v_1), (v_0, v_2), \dots, (v_0, v_n)\}$
- ▷ $w(v_0, v_j) = 0$ for all j
- ▷ $w(v_i, v_j) = b_k$ if $x_j - x_i \text{ ? } b_k$



► **Theorem**

Given a system of difference constraints, let $G = (V, E)$ be the corresponding Constraint graph.

1. If G has no negative-weight cycles, then

$x = (\delta(v_0, v_1), \delta(v_0, v_2), \dots, \delta(v_0, v_n))$ is a feasible solution.

2. If G has a negative-weight cycle, then there is no feasible solution.

► **Proof**

1. Show no negative-weight cycles \Rightarrow feasible solution.

Need to show that $x_j - x_i \leq b_k$ for all constraints. Use

$$x_j = \delta(v_0, v_j)$$

$$x_i = \delta(v_0, v_i)$$

$$b_k = w(v_i, v_j)$$

By the triangle inequality,

$$\delta(v_0, v_j) \leq \delta(v_0, v_i) + w(v_i, v_j)$$

$$x_j \leq x_i + b_k$$

$$x_j - x_i \leq b_k$$

Therefore, feasible.

2. Show negative-weight cycles \Rightarrow no feasible solution.

Without loss of generality, let a negative-weight cycle be $C = \langle v_1, v_2, \dots, v_k \rangle$ where $v_1 = v_k$ (v_0 can't be on C , since v_0 has no entering edges.) C corresponds to the constraints

$$x_2 - x_1 \leq w(v_1, v_2)$$

$$x_3 - x_2 \leq w(v_2, v_3)$$

M

$$x_{k-1} - x_{k-2} \leq w(v_{k-2}, v_{k-1})$$

$$x_k - x_{k-1} \leq w(v_{k-1}, v_k)$$

(The last two inequalities above are incorrect in the first three printings of the book. They were corrected in the fourth printing.)



If x is a solution satisfying these inequalities, it must satisfy their sum.

So add them up.

Each x_i is added once and subtracted once. ($v_1 = v_k$ \diamond $x_1 = x_k$)

We get $0 \leq w(C)$.

But $w(C) < 0$, since C is a negative-weight cycle.

Contradiction \Rightarrow no such feasible solution x exists.

■ (theorem)

► How to find a feasible solution

1. Form constraint graph.

- $n + 1$ vertices.
- $m + n$ edges.
- $\Theta(m + n)$ time.

2. Run BELLMAN-FORD from v_0 .

- $O((n + 1)(m + n)) = O(n^2 + nm)$ time.

3. If BELLMAN-FORD returns FALSE \Rightarrow no feasible solution.

If BELLMAN-FORD returns TRUE

\Rightarrow set $x_i = \delta(v_0, v_i)$ for all i .