

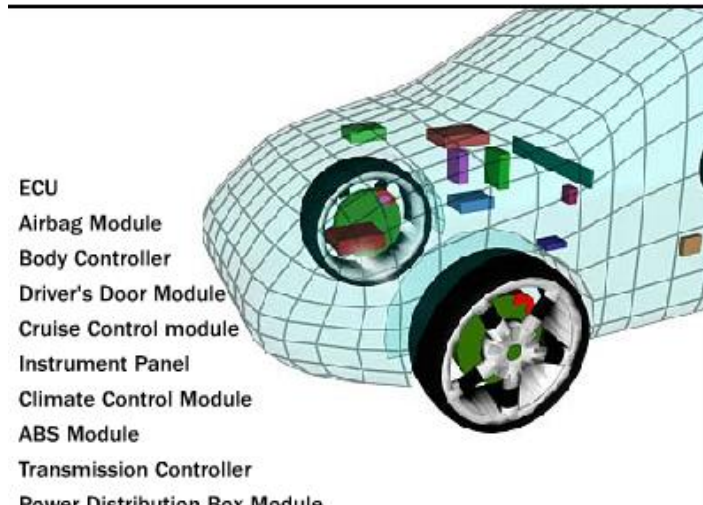
Chapter 1

Computer Abstractions and Technology

Da-Wei Chang, OSES Lab.
CSIE Dept., NCKU

The Computer Revolution

- Computers are almost everywhere
 - PC, Mobile Phone, embedded system, etc
- Makes novel applications feasible
 - Human genome project
 - World Wide Web
 - Search Engines

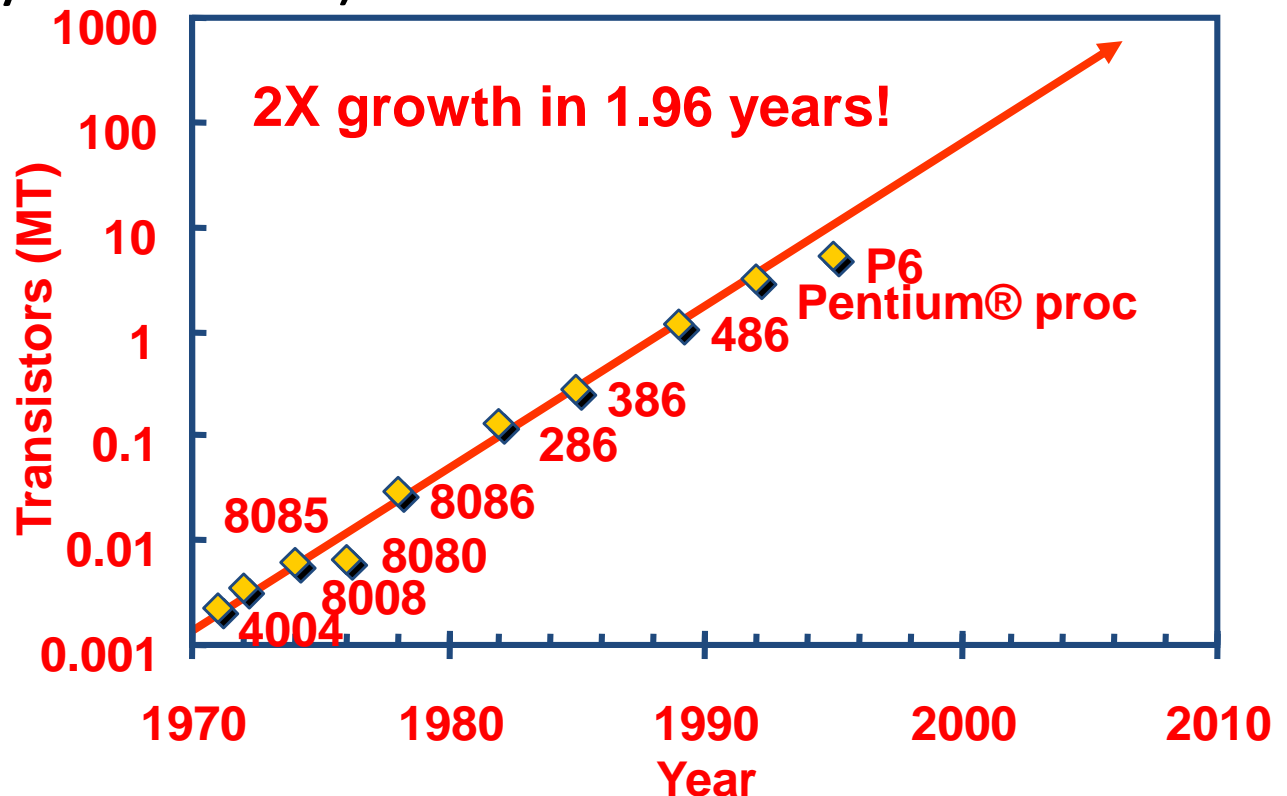


Moore's Law

- In 1965, Gordon Moore predicted that the number of transistors that can be integrated on a die would double every 18 to 24 months (i.e., grow exponentially with time).



Gordon Moore

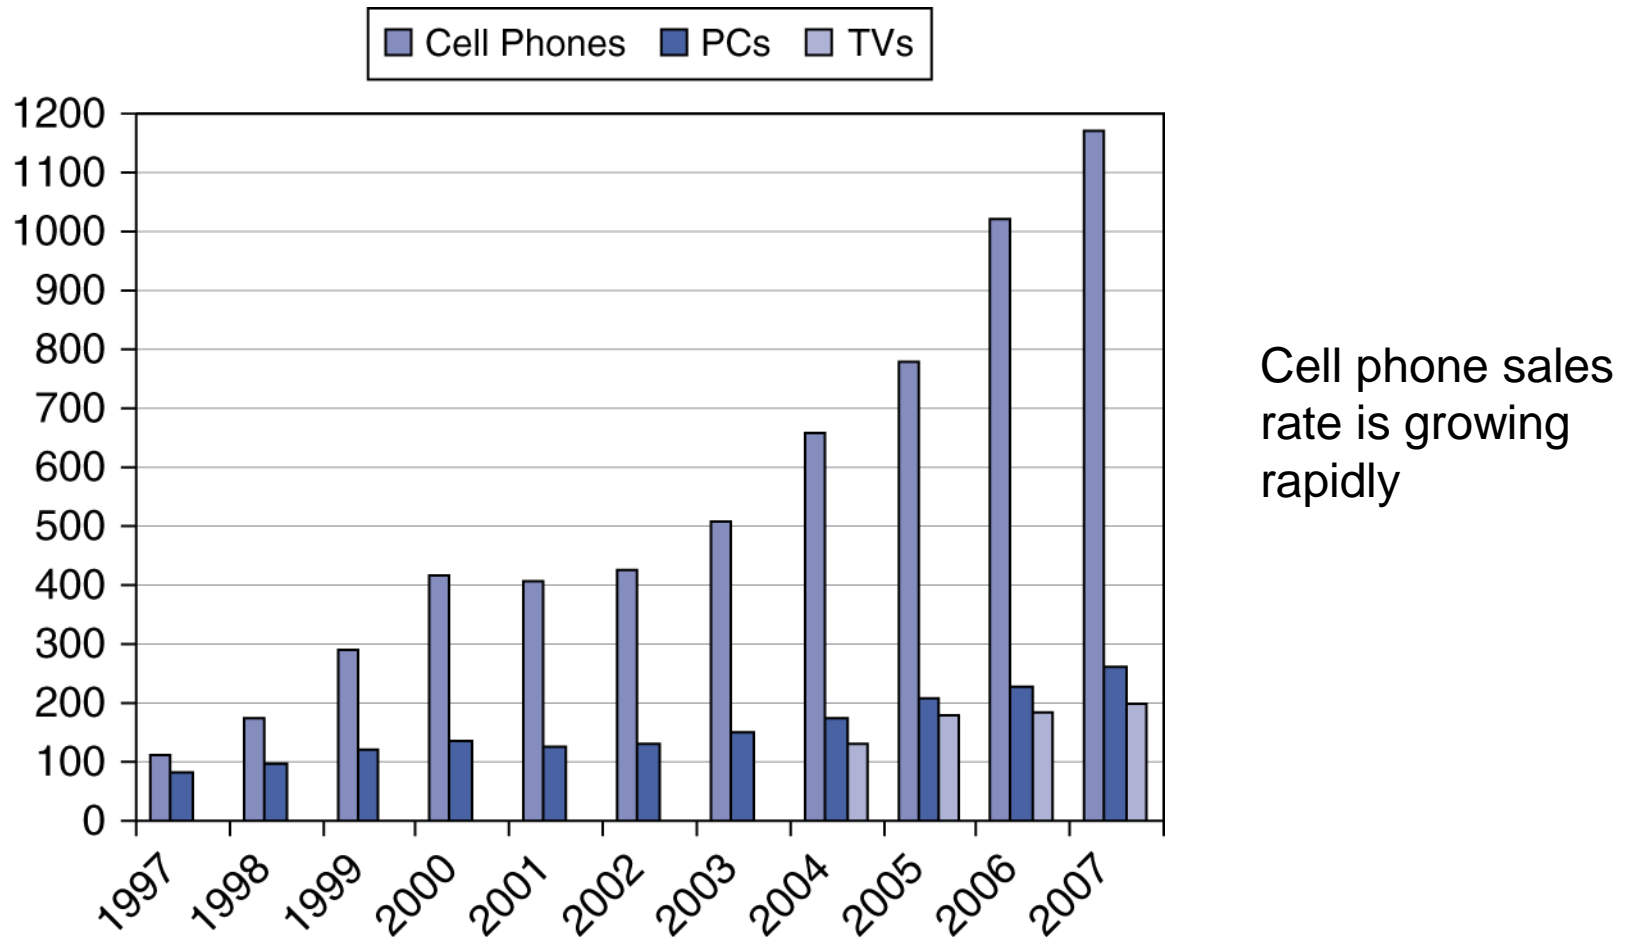


Classes of Computers

- **Desktop** computers
 - General purpose, variety of software
- **Server** computers
 - **Network** based
 - High **capacity**, performance, reliability
 - Range from small servers to building sized
- **Embedded** computers
 - Hidden as components of systems
 - Strict power/performance/cost constraints



The Processor Market

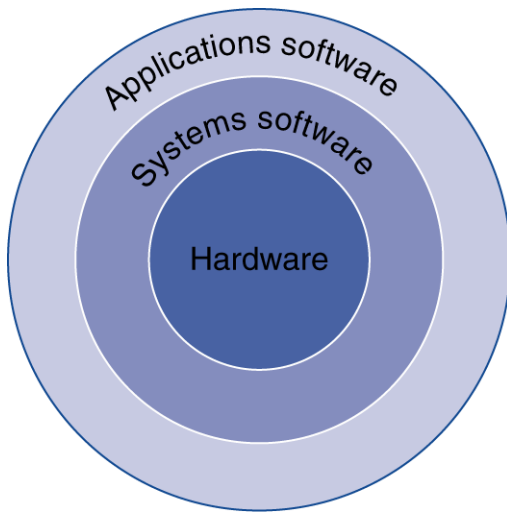


What You Will Learn in Chapter 1

- How **programs** are translated into the **machine** language
 - And how the hardware executes them
- The **hardware/software** interface
 - Instructions
- What determines program **performance**
 - How to compare the performance
- How hardware designers **improve performance**
- What is parallel processing

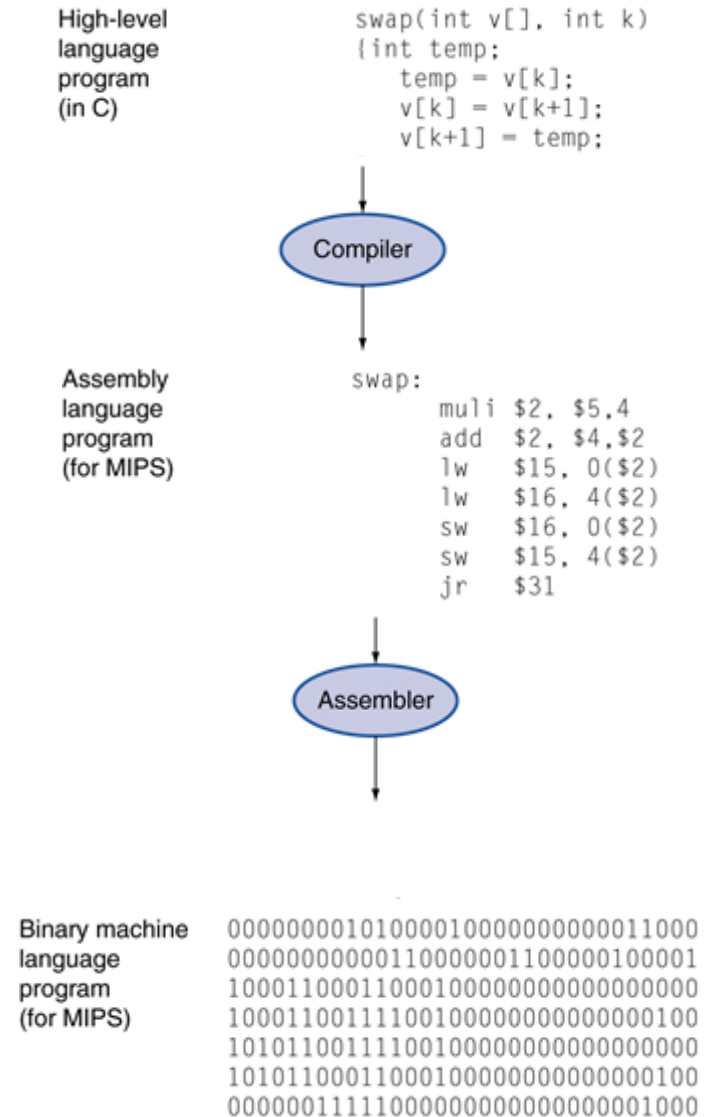
Below Your Program

- Application software
 - Written in high-level language (HLL)
- System software
 - **Compiler**: translates HLL code to machine code
 - **Operating System**: service code
 - Handling input/output
 - Managing memory and storage
 - Scheduling tasks & sharing resources
- Hardware
 - Processor, memory, I/O controllers



Levels of Program Code

- **High-level** language
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
- **Assembly** language
 - Textual representation of instructions
- **Hardware** representation (Machine code)
 - Binary digits (bits)
 - Encoded instructions and data



Abstractions

The BIG Picture

- Abstraction helps us deal with complexity
 - Hide lower-level detail
- Instruction set architecture (ISA)
 - The hardware/software interface
- Application binary interface (ABI)
 - The ISA plus system software interface
- Implementation
 - The details underlying the interface

Components of a Computer

- **Five** components

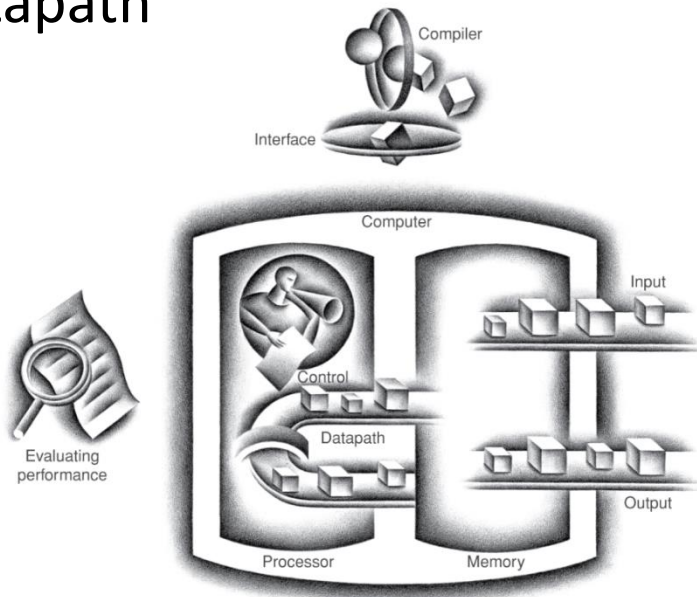
- Input
- Output
- Memory
- Control
- Datapath

- **Almost Same** components for all kinds of computer

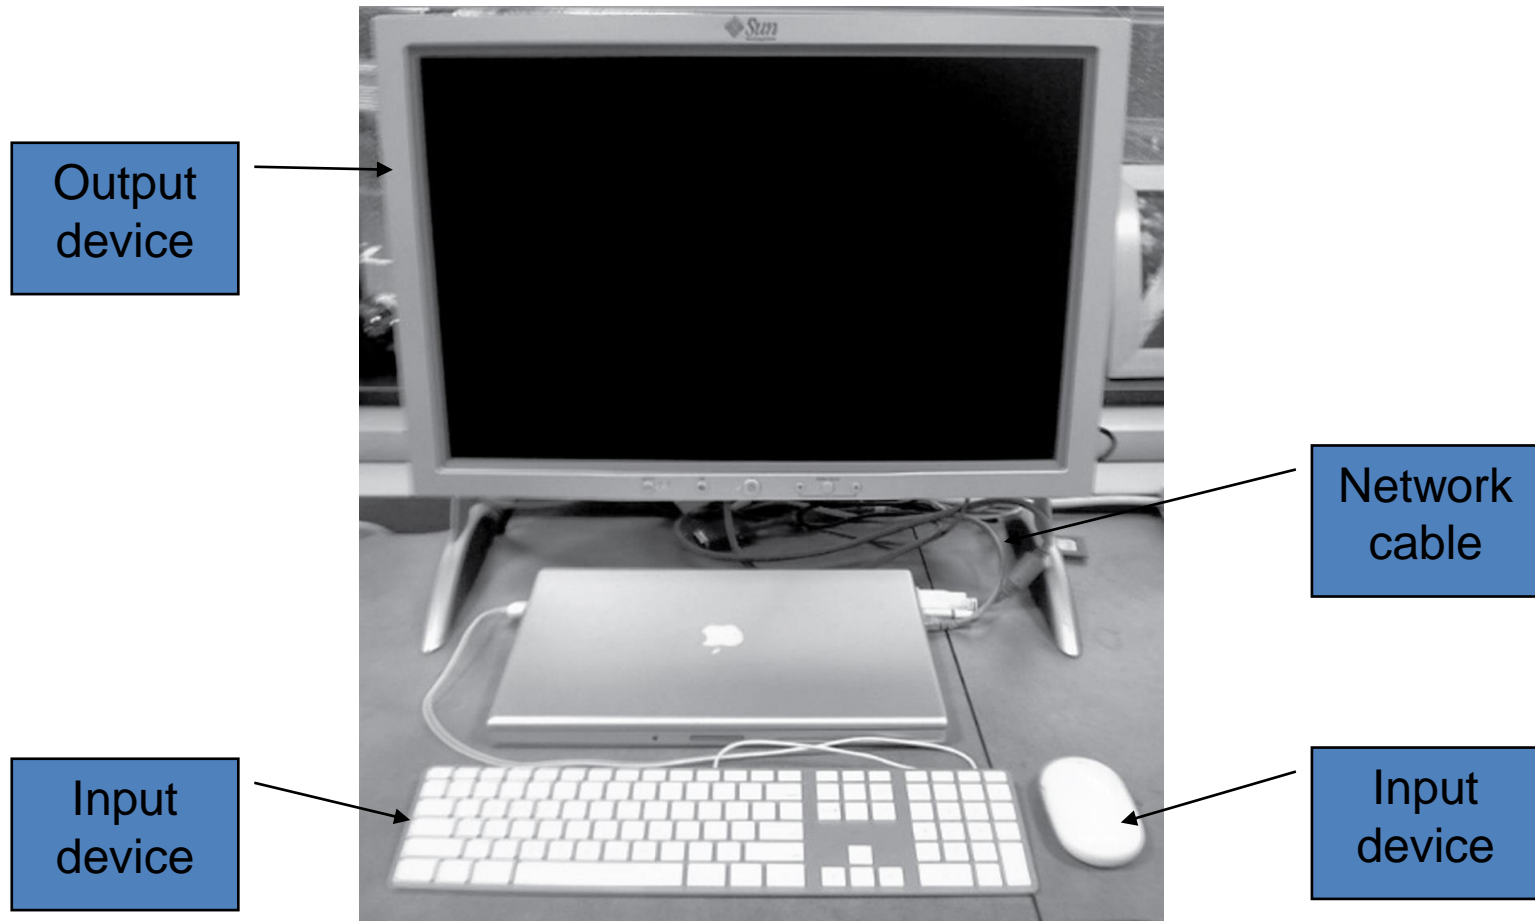
- Desktop, server, embedded

- Input/output includes

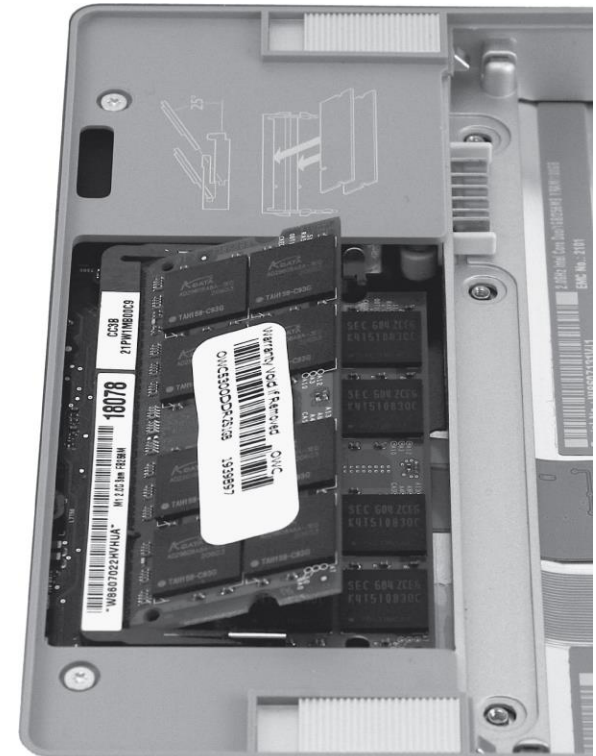
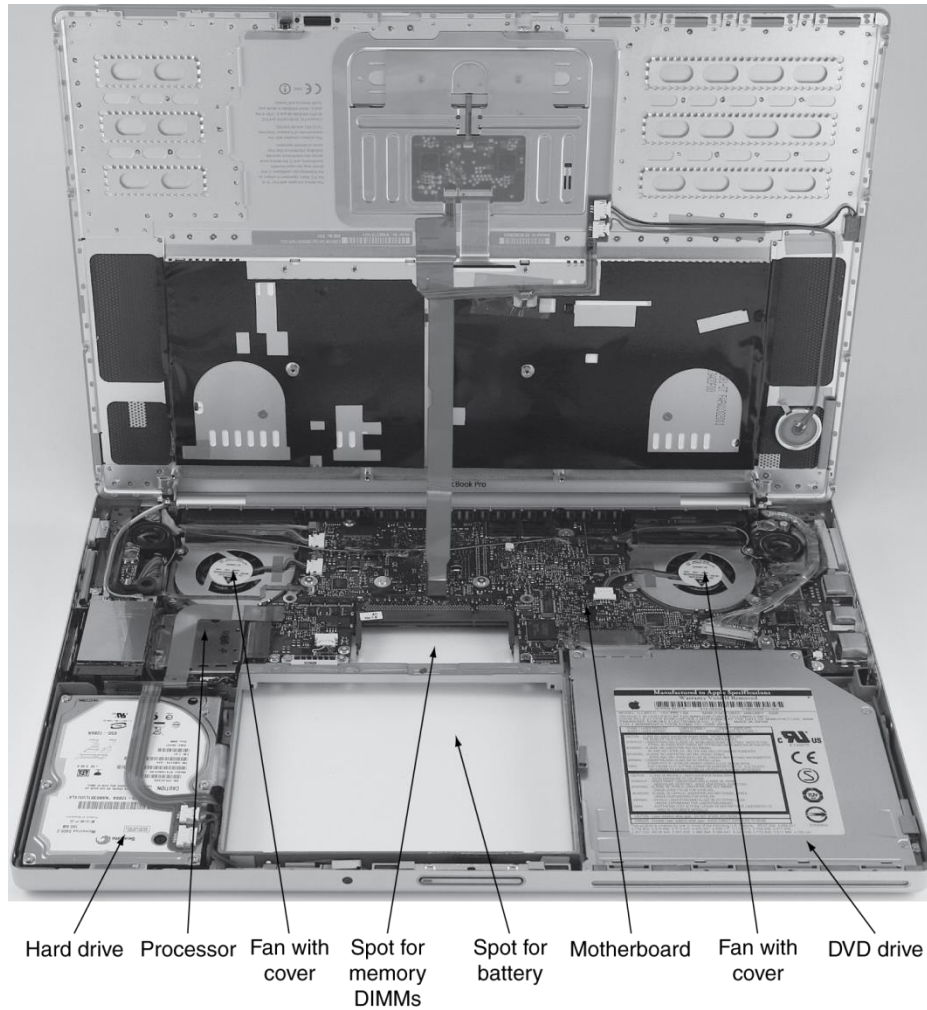
- User-interface devices
 - **Display**, keyboard, mouse
- Storage devices
 - **Hard disk**, CD/DVD, flash
- Network adapters
 - For **communicating** with other computers



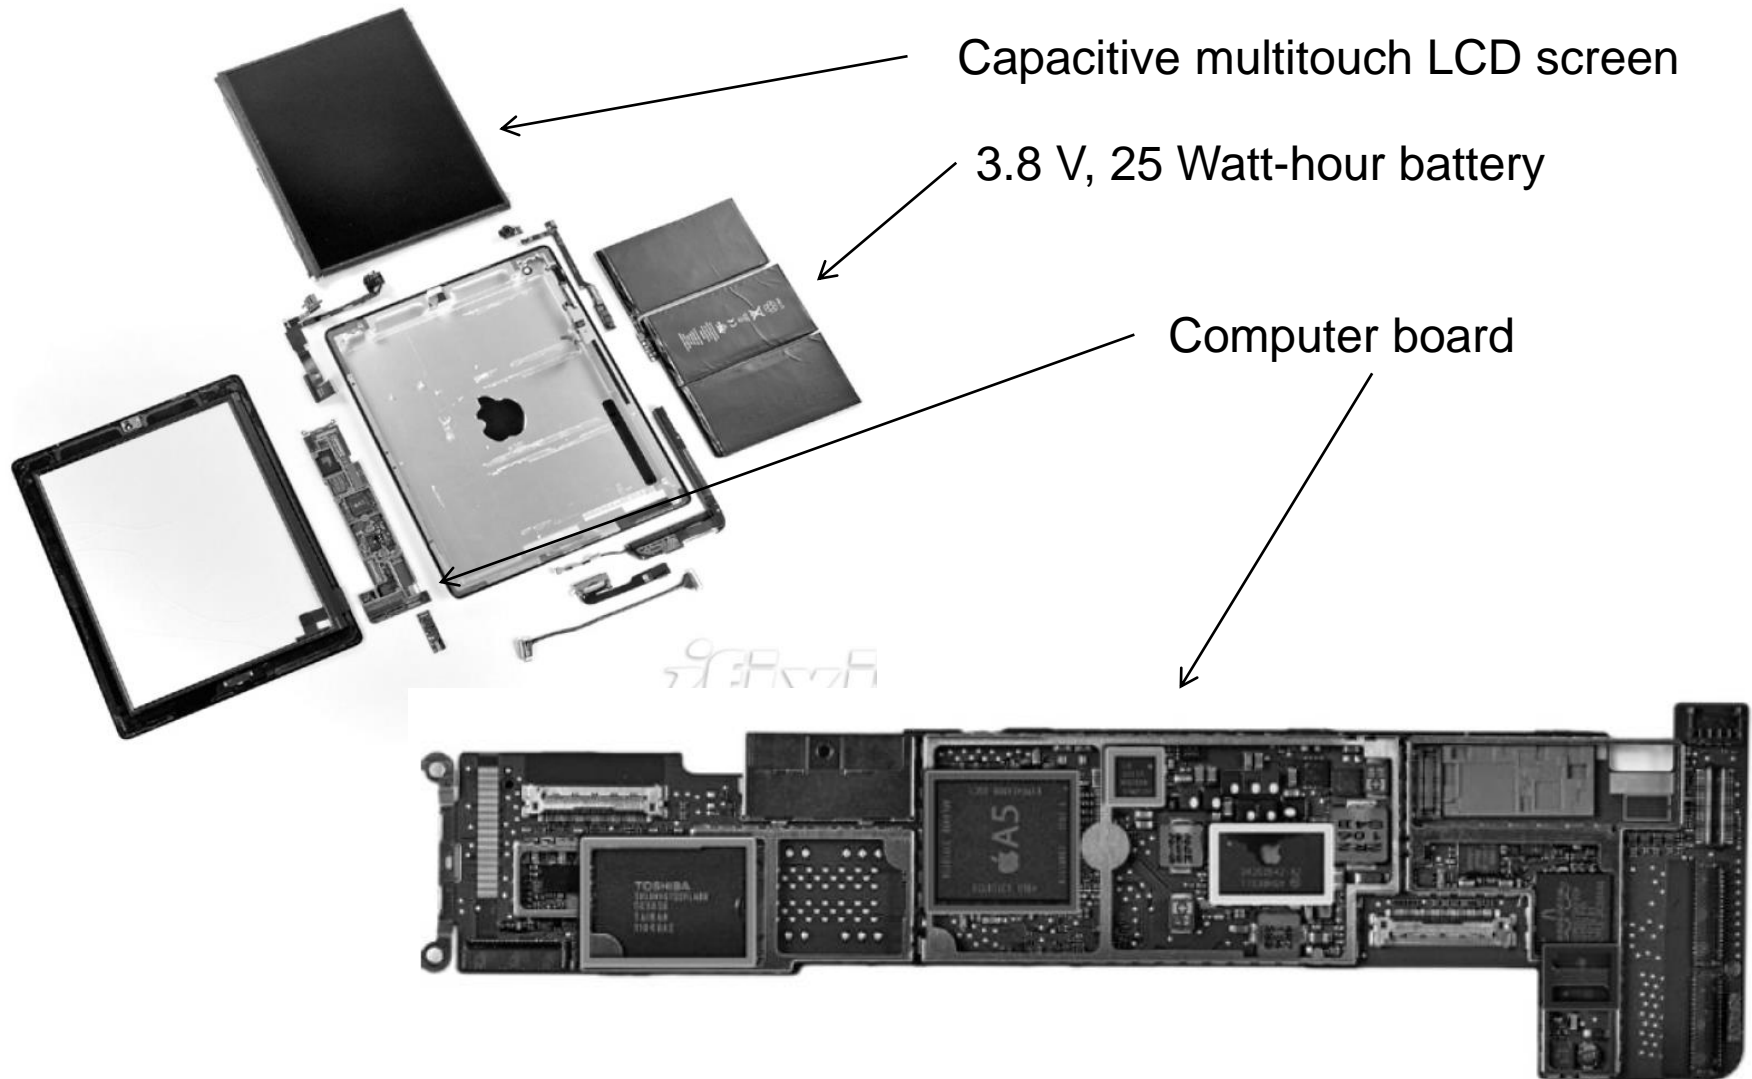
Anatomy of a Computer



Opening the Box (I)

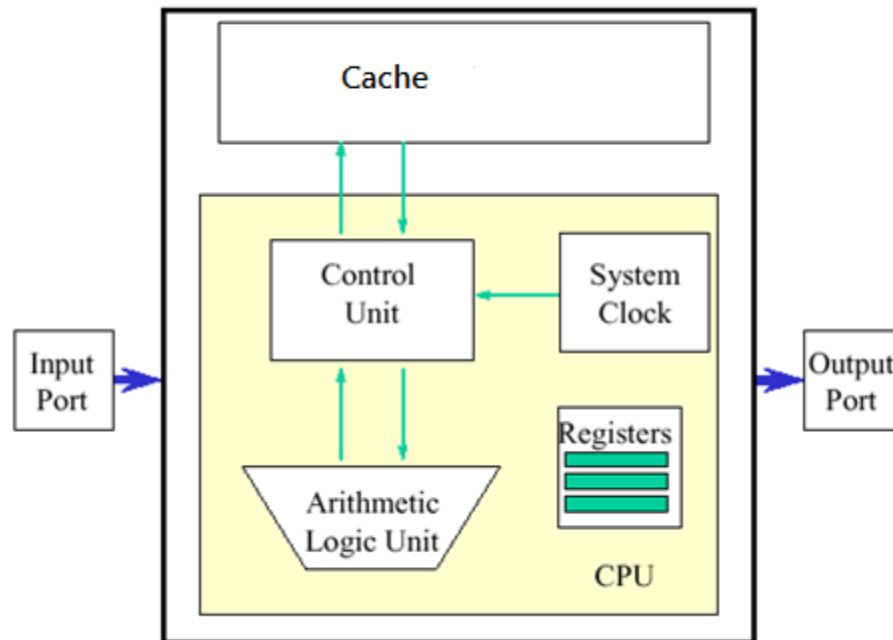


Opening the Box (II)



Inside the Processor (CPU)

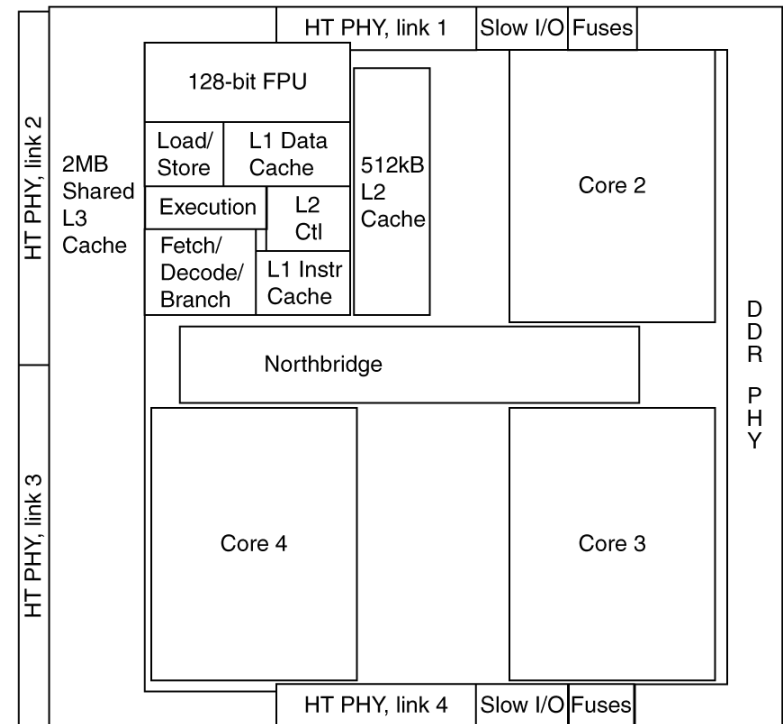
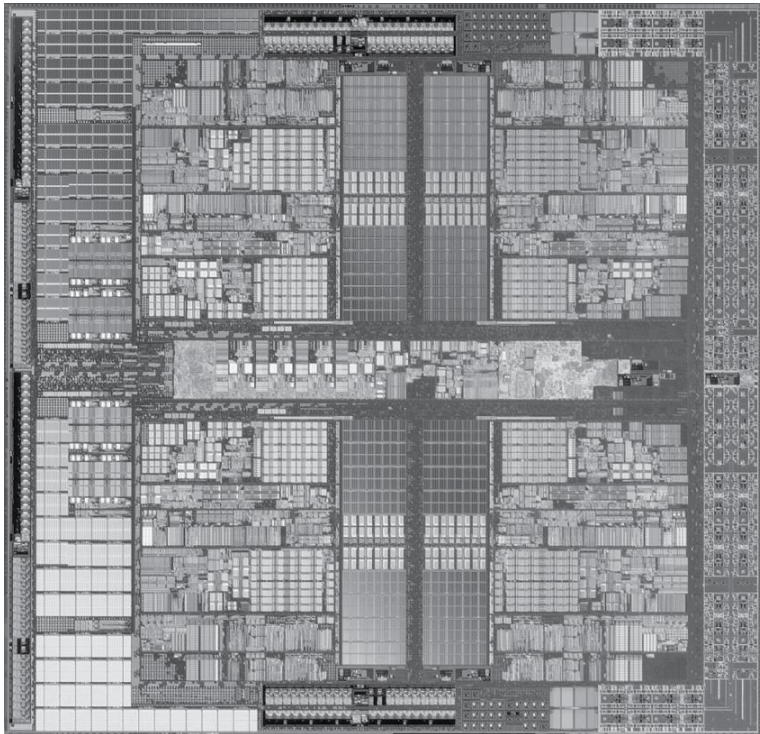
- **Datapath**: performs operations on data
- **Control**: sequences datapath, memory, ...
- **Cache** memory
 - Small fast SRAM memory for immediate access to data



More in Chapter
3, 4 and 5

Inside the Processor: Example I

- AMD Barcelona: 4 processor cores



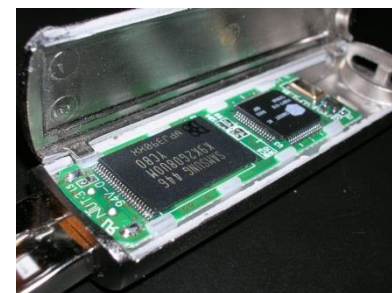
Inside the Processor: Example II

- Apple A5



A Safe Place for Data

- **Volatile** main memory (primary memory)
 - Loses instructions and data when **power off**
- **Non-volatile** secondary memory
 - Magnetic disk
 - Flash memory
 - Optical disk (CDROM, DVD)



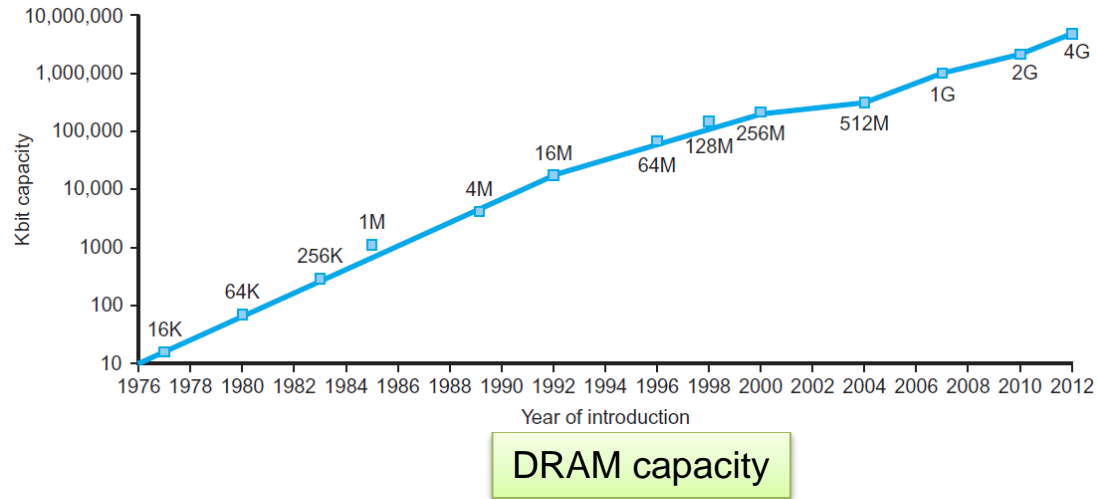
Networks

- Communication and resource sharing
- Local area network (LAN): Ethernet
 - Within a building
- Wide area network (WAN: the Internet)
- Wireless network: WiFi, Bluetooth



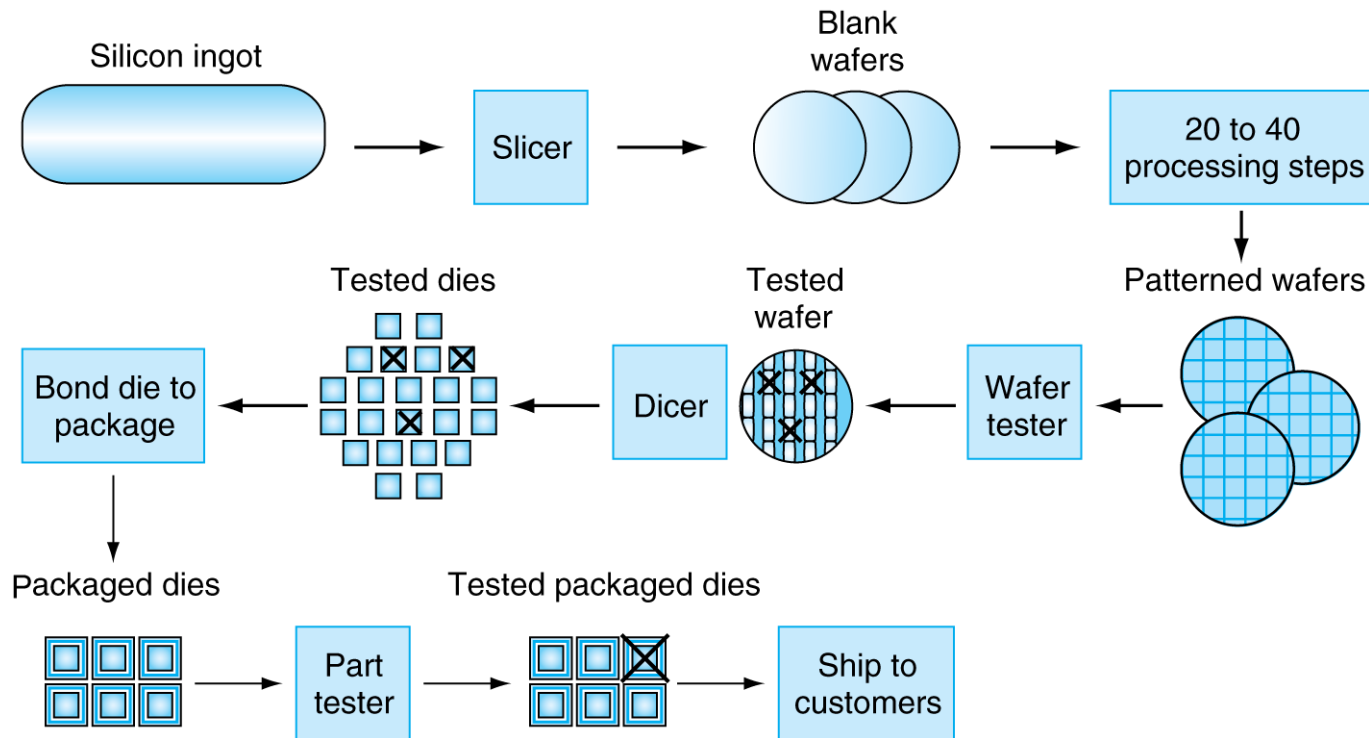
Technology Trends

- Electronics technology continues to evolve
 - Increased **capacity** and performance
 - Reduced **cost**



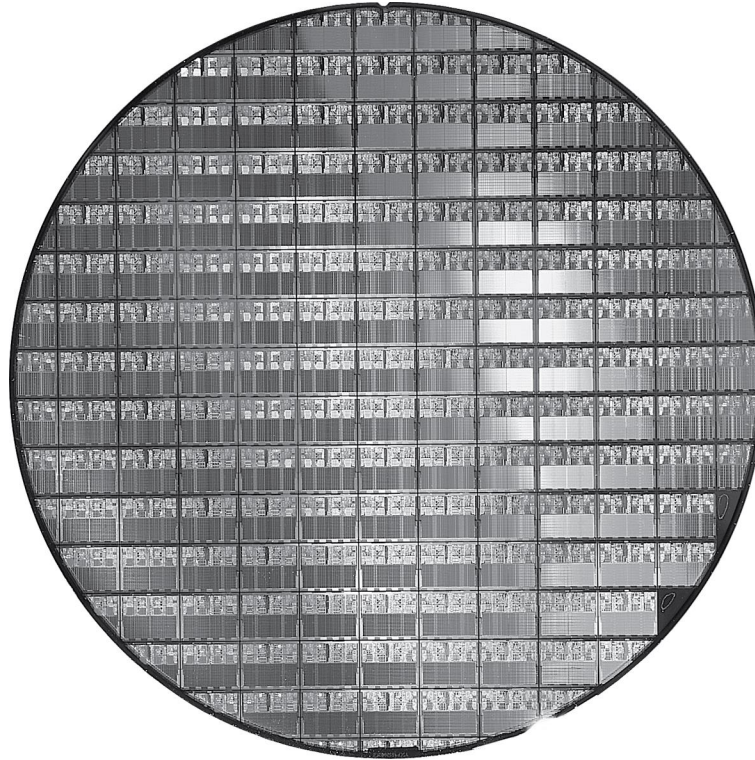
Year	Technology	Relative performance/cost
1951	Vacuum tube	1
1965	Transistor	35
1975	Integrated circuit (IC)	900
1995	Very large scale IC (VLSI)	2,400,000
2013	Ultra large scale IC	250,000,000,000

Manufacturing ICs



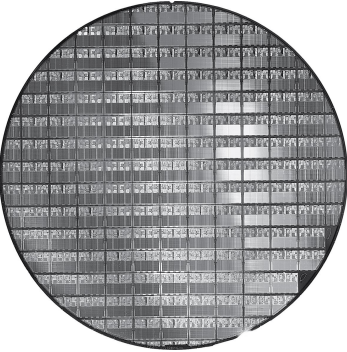
- **Yield**: proportion of working dies per wafer

AMD Opteron X2 Wafer



- X2: 300mm wafer, 117 chips, 90nm technology
- X4: 45nm technology

Integrated Circuit Cost



$$\text{Cost per die} = \frac{\text{Cost per wafer}}{\text{Dies per wafer} \times \text{Yield}}$$

$$\text{Dies per wafer} \approx \text{Wafer area} / \text{Die area}$$

$$\text{Yield} = \frac{1}{(1 + (\text{Defects per area} \times \text{Die area} / 2))^2}$$

Just an approximation

- **Wafer** cost and area are fixed
- **Defect rate** determined by manufacturing process
- **Die area** determined by **architecture** and **circuit** design

Which airplane has the best performance?

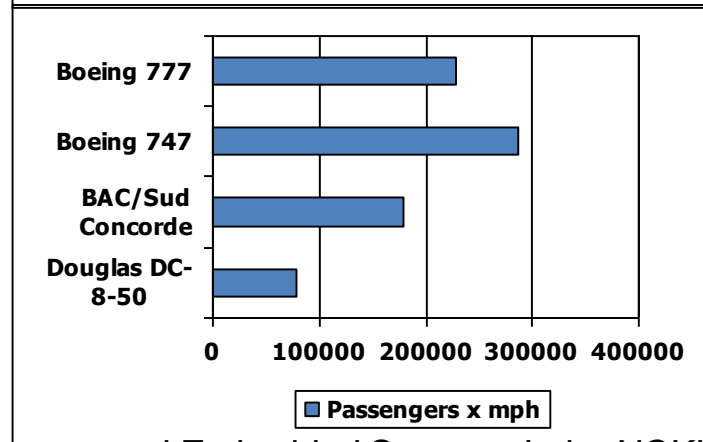
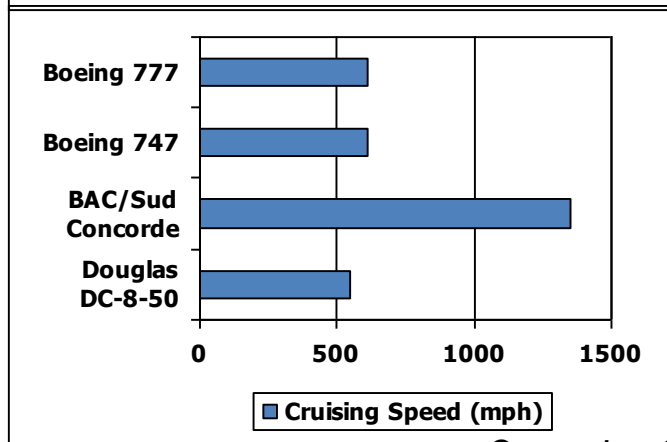
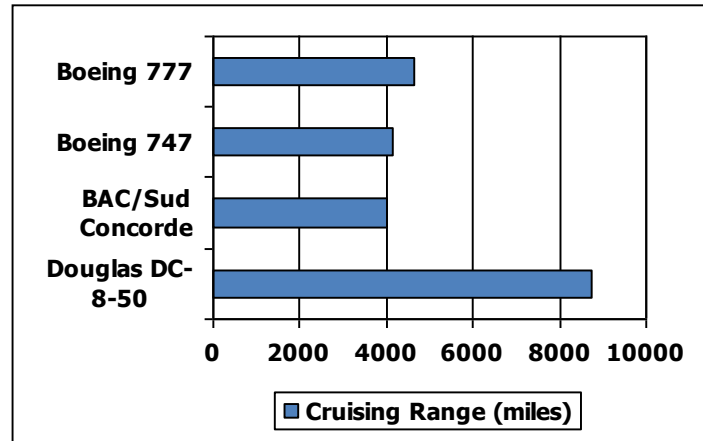
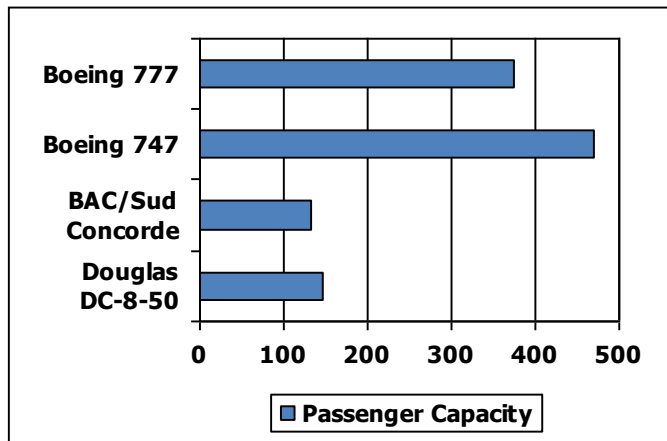
- Depend how you define performance

Largest **capacity**: 747

Highest **speed**: Concorde



Concorde



Many different ways to define performance

Response Time and Throughput

- **Response** time
 - The time between the start and completion of a task
 - e.g., 5 min to finish a task/transaction
- **Throughput**
 - Total amount of work done in a given time
 - e.g., tasks/transactions/... per hour
- We'll focus on response time for now...
- Performance of a computer X

$$Performance_X = \frac{1}{Execution\ Time_X}$$

Relative Performance

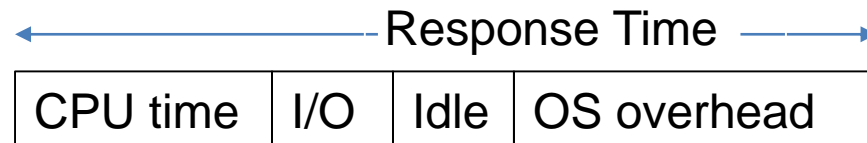
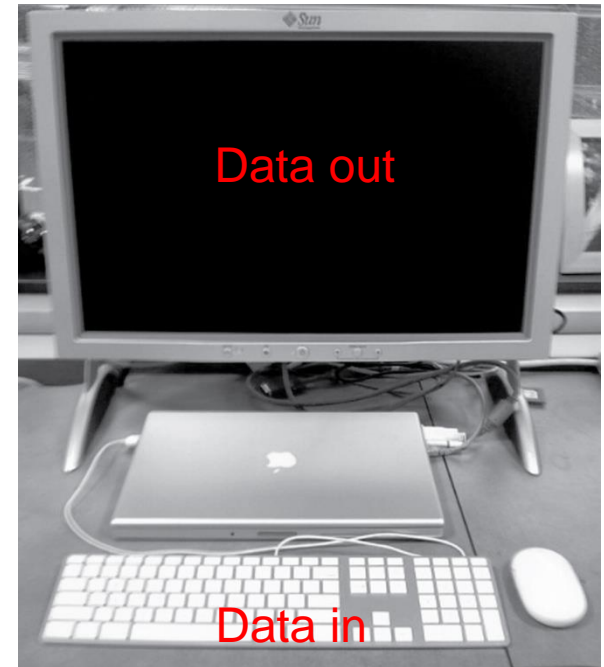
- Define Performance = 1/Execution Time
- “X is n time faster than Y”

$$\frac{Performance_X}{Performance_Y} = n = \frac{Execution\ Time_Y}{Execution\ Time_X}$$

- Example: A takes 10s and B takes 15s to run a program, how much faster is A than B?
 - $Execution\ Time_B / Execution\ Time_A$
 $= 15s / 10s = 1.5$
 - So A is 1.5 times faster than B

Response Time

- **Response** time (or called **elapsed** time)
 - **Total** response time,
 - including all aspects
 - Processing,
 - I/O,
 - OS overhead
 - idle time
 - e.g. A task takes **5** min to finish, including all processing time, I/O, and etc.
 - It is the **time** that users will experience

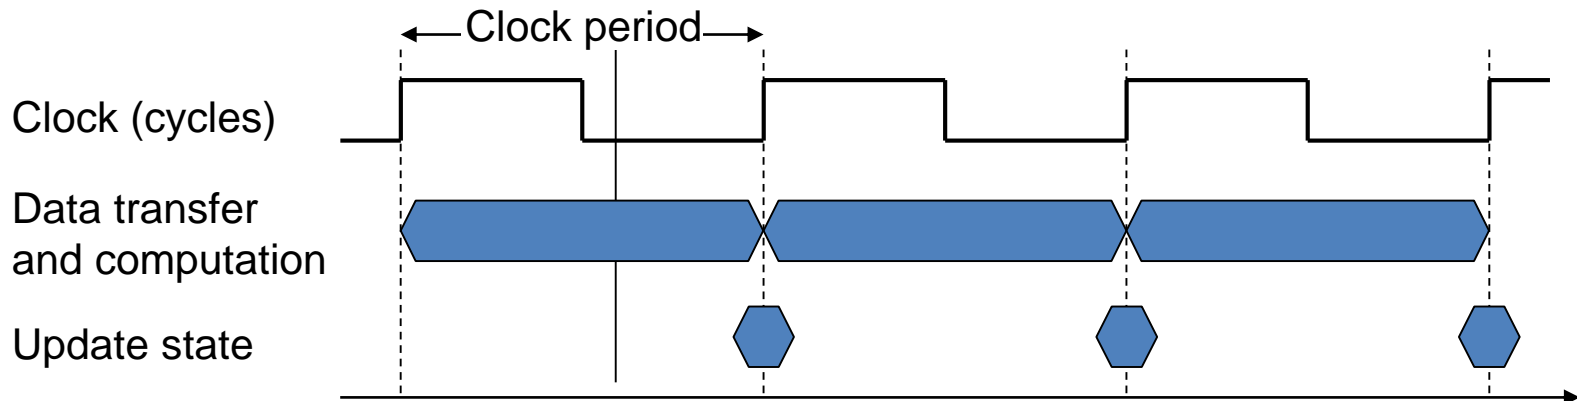


CPU Time

- CPU execution time (CPU time)
 - Time spent processing a given job (CPU may process many jobs at the same time)
 - Do not consider I/O time, other jobs' shares
 - Can be divided into user CPU time and system CPU time
 - User CPU time: CPU time spent in the program
 - System CPU time: CPU time spent in OS for this program

Review: CPU Clocking

- Operation of digital hardware governed by a constant-rate clock



- Clock **period**: duration of a clock cycle
 - e.g., $250\text{ps} = 0.25\text{ns} = 250 \times 10^{-12}\text{s}$
- Clock **frequency** (rate): cycles per second
 - e.g., $4.0\text{GHz} = 4000\text{MHz} = 4.0 \times 10^9\text{Hz}$

CPU Time

$$\begin{aligned}\text{CPU Time} &= \text{CPU Clock Cycles} \times \text{Clock Cycle Time} \\ &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}}\end{aligned}$$

- Performance improved by
 - Reducing number of clock cycles
 - Increasing clock rate
 - Hardware designer must often trade off clock rate against cycle count
- Example: a task takes 100 CPU cycles and each cycle is 1ns, what is the CPU Time?

$$\text{CPU Time} = 100 \times 1 = 100 \text{ ns}$$

CPU Time Example

- Computer A: 2GHz clock, 10s CPU time. Design Computer B. It has $1.2 \times$ clock cycles. However, CPU time is 6s because it uses faster clock. What is the clock rate of Computer B?

$$\begin{aligned}\text{Clock Cycles}_A &= \text{CPU Time}_A \times \text{Clock Rate}_A \\ &= 10\text{s} \times 2\text{GHz} = 20 \times 10^9\end{aligned}$$

$$\begin{aligned}\text{Clock Rate}_B &= \frac{\text{Clock Cycles}_B}{\text{CPU Time}_B} = \frac{1.2 \times \text{Clock Cycles}_A}{6\text{s}} \\ &= \frac{1.2 \times 20 \times 10^9}{6\text{s}} = \frac{24 \times 10^9}{6\text{s}} = 4\text{GHz}\end{aligned}$$

Instruction Count and CPI

- CPI: Cycle per instruction
 - Number of cycles each instruction takes to execute

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}}$$

- Instruction Count for a program determined by **program, ISA and compiler**
- **Redefine CPU time using CPI**

$$\text{CPU Time} = \text{Clock Cycles} \times \text{Clock Cycle Time}$$

$$= \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

$$= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$$

CPI Example

- Computer A: Cycle Time = 250ps, CPI = 2.0
- Computer B: Cycle Time = 500ps, CPI = 1.2
- Same ISA
- Which is faster, and by how much?

$$\begin{aligned}\text{CPU Time}_A &= \text{Instruction Count} \times \text{CPI}_A \times \text{Cycle Time}_A \\ &= 1 \times 2.0 \times 250\text{ps} = 1 \times 500\text{ps}\end{aligned}$$

A is faster...

$$\begin{aligned}\text{CPU Time}_B &= \text{Instruction Count} \times \text{CPI}_B \times \text{Cycle Time}_B \\ &= 1 \times 1.2 \times 500\text{ps} = 1 \times 600\text{ps}\end{aligned}$$

$$\frac{\text{CPU Time}_B}{\text{CPU Time}_A} = \frac{1 \times 600\text{ps}}{1 \times 500\text{ps}} = 1.2$$

1.2 times faster

CPI in More Detail

- If different **instruction** classes take different numbers of **cycles**

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

■ **Weighted** average CPI

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(\text{CPI}_i \times \frac{\text{Instruction Count}_i}{\text{Instruction Count}} \right)$$

Relative frequency

CPI Example

- Alternative compiled code sequences using instructions in classes A, B, C, find average **CPI for instruction sequence 1 and 2**

Instruction Class	A	B	C
CPI for class	1	2	3
Inst. Count in sequence 1	2	1	2
Inst. Count in sequence 2	4	1	1

- Sequence 1: IC = **5**

- Clock Cycles
 $= 2 \times 1 + 1 \times 2 + 2 \times 3$
 $= 10$
- Avg. CPI = $10/5 = 2.0$

- Sequence 2: IC = **6**

- Clock Cycles
 $= 4 \times 1 + 1 \times 2 + 1 \times 3$
 $= 9$
- Avg. CPI = $9/6 = 1.5$

Performance Summary

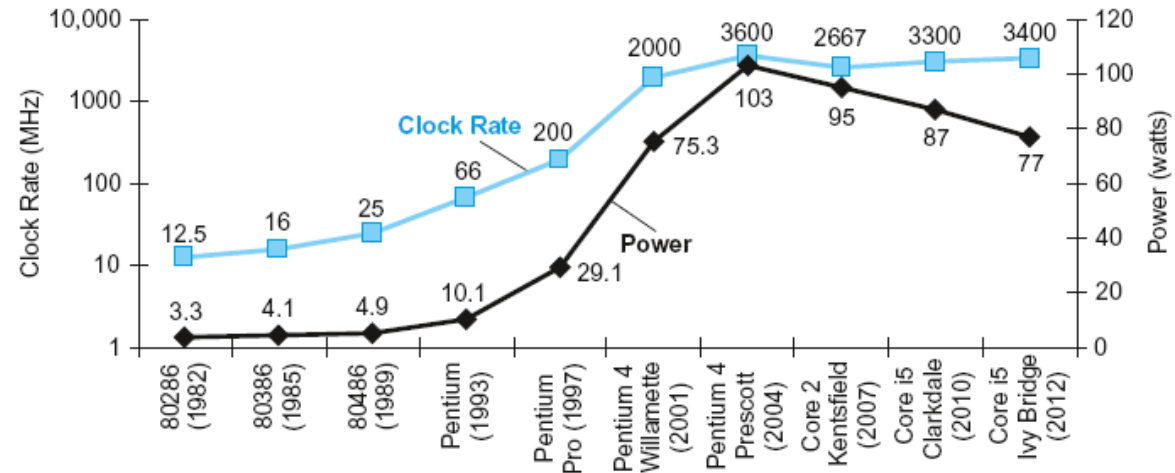
The BIG Picture

$$\text{CPU Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$

- Performance depends on
 - Algorithm: affects IC, possibly CPI
 - Programming language: affects IC, CPI
 - Compiler: affects IC, CPI
 - Instruction set architecture: affects IC, CPI, T_c

Power Trends

- CPU clock rate and power has increased rapidly for decades



- In CMOS IC technology

$$\text{Power} = \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency}$$

×30

5V → 1V

×1000

Reducing Power

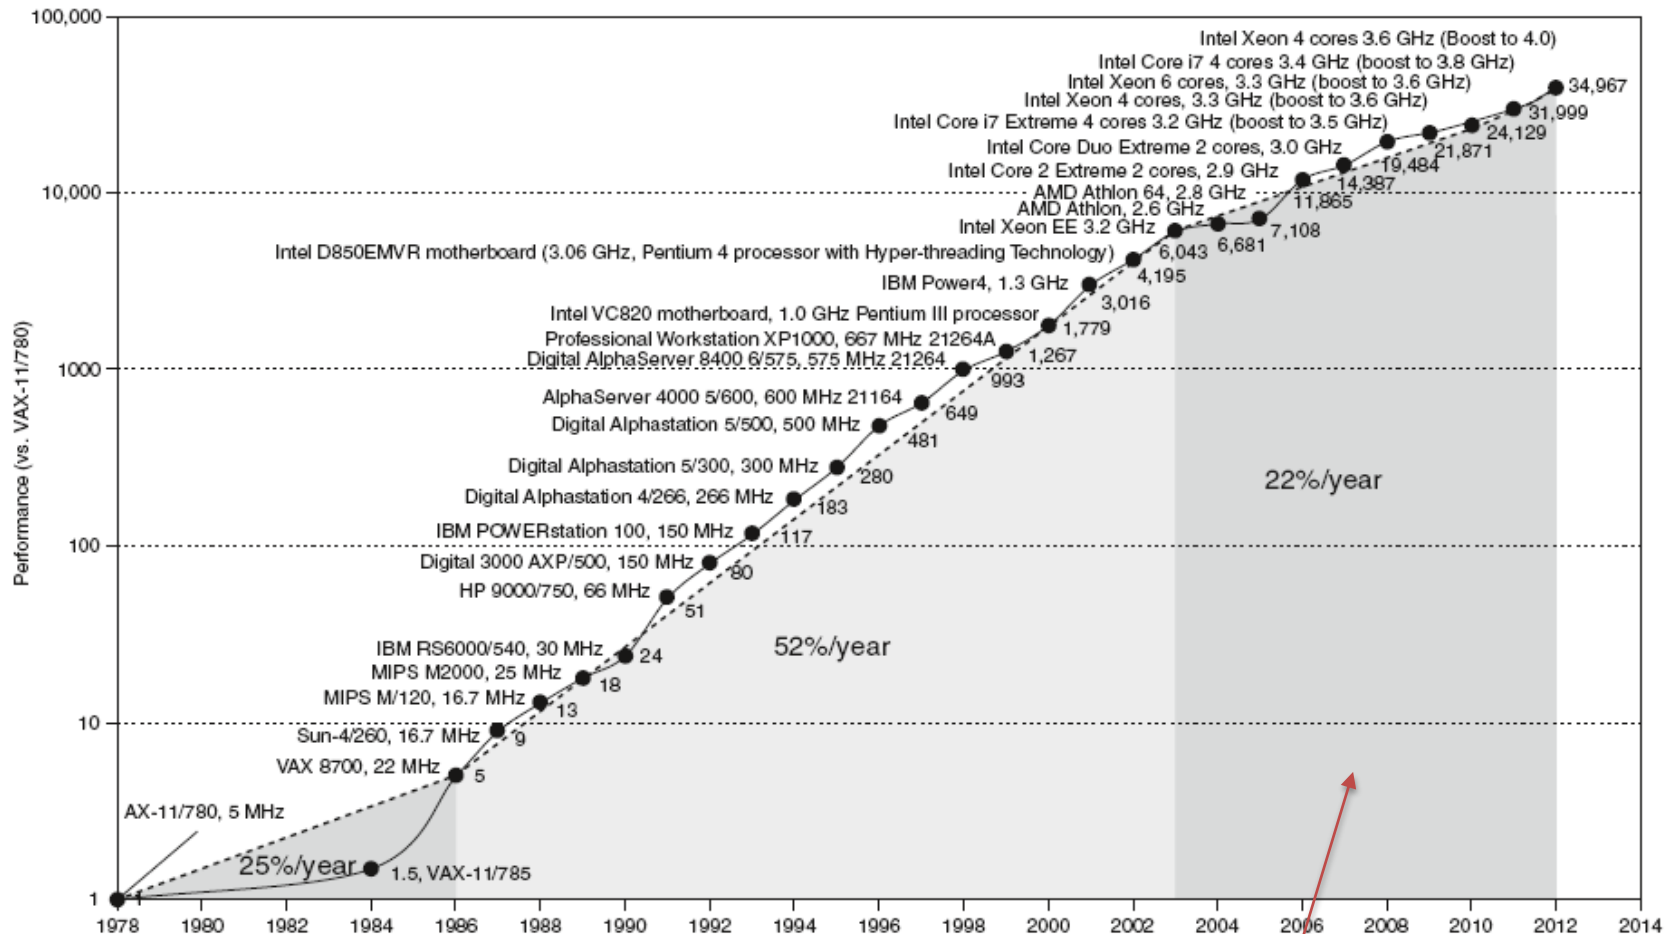
- Suppose a new CPU has **85%** of capacitive load of old CPU, and its **voltage** and **frequency** are reduced by **15%**

– What is $\frac{\text{Power}_{\text{new}}}{\text{Power}_{\text{old}}}$?

$$\frac{P_{\text{new}}}{P_{\text{old}}} = \frac{C_{\text{old}} \times 0.85 \times (V_{\text{old}} \times 0.85)^2 \times F_{\text{old}} \times 0.85}{C_{\text{old}} \times V_{\text{old}}^2 \times F_{\text{old}}} = 0.85^4 = 0.52$$

- The power wall
 - Hard to reduce voltage further
 - Hard to remove more **heat**
- How else can we improve performance?

Uniprocessor Performance



Constrained by power, instruction-level parallelism, memory latency

Multiprocessors

- **Multicore** microprocessors
 - **More** than one **processor** per chip
- Requires explicitly **parallel programming**
 - Compare with instruction level parallelism
 - Hardware executes multiple instructions at once
 - Hidden from the programmer
 - Hard to do
 - Programming for performance
 - Load balancing
 - Optimizing communication and synchronization

SPEC CPU Benchmark

- **Programs** used to measure performance
 - Supposedly **typical** of actual workload
- Standard Performance Evaluation Corp (SPEC)
 - Develops benchmarks for CPU, I/O, Web, ...
- **SPEC CPU2006**
 - **12 integer and 17 floating-point benchmarks**
 - CINT2006 (integer) and CFP2006 (floating-point)
 - Elapsed time to execute a selection of programs
 - Negligible I/O, so **focuses on CPU performance**
 - Include C compiler to chess program to quantum computer simulation
 - Normalized relative to reference machine
 - See next slide for example

CINT2006 for Opteron X4 2356

Ref time: Time required by a reference computer

Name	Description	IC×10 ⁹	CPI	Tc (ns)	Exec time	Ref time	SPEC ratio
perl	Interpreted string processing	2,118	0.75	0.40	637	9,777	15.3
bzip2	Block-sorting compression	2,389	0.85	0.40	817	9,650	11.8
gcc	GNU C Compiler	1,050	1.72	0.40	724	8,050	11.1
mcf	Combinatorial optimization	336	10.00	0.40	1,345	9,120	6.8
go	Go game (AI)	1,658	1.09	0.40	721	10,490	14.6
hmmer	Search gene sequence	2,783	0.80	0.40	890	9,330	10.5
sjeng	Chess game (AI)	2,176	0.96	0.40	837	12,100	14.5
libquantum	Quantum computer simulation	1,623	1.61	0.40	1,047	20,720	19.8
h264avc	Video compression	3,102	0.80	0.40	993	22,130	22.3
omnetpp	Discrete event simulation	587	2.94	0.40	690	6,250	9.1
astar	Games/path finding	1,082	1.79	0.40	773	7,020	9.1
xalancbmk	XML parsing	1,058	2.70	0.40	1,143	6,900	6.0
Geometric mean							11.7

15.3 times faster

CPU Time=Instruction Count×CPI ×Clock Cycle Time

SPEC Power Benchmark

- Power consumption of server at different workload levels
 - Performance: ssj_ops/sec
 - Power: Watts (Joules/sec)

$$\text{Overall ssj_ops per Watt} = \left(\sum_{i=0}^{10} \text{ssj_ops}_i \right) / \left(\sum_{i=0}^{10} \text{power}_i \right)$$

SPECpower_ssj2008 for X4

Target Load %	Performance (ssj_ops/sec)	Average Power (Watts)
100%	231,867	295
90%	211,282	286
80%	185,803	275
70%	163,427	265
60%	140,160	256
50%	118,324	246
40%	920,35	233
30%	70,500	222
20%	47,126	206
10%	23,066	180
0%	0	141
Overall sum	1,283,590	2,605
$\Sigma \text{ssj_ops} / \Sigma \text{power}$		493

Pitfall

- Improving an aspect of a computer and expecting a proportional improvement in **overall** performance
 - Example: **multiplication** accounts for 80s of 100s. Improve **multiplication** by **5x** and expect overall performance to improve **5x**, which is **20s**
 - => Wrong
- Amdahl's Law: Possible improvement is limited by the amount that the improved feature is used

$$T_{\text{improved}} = \frac{T_{\text{affected}}}{\text{improvement}} + T_{\text{unaffected}}$$

Amdahl's Law

- Amdahl's Law: Possible improved is limited by the amount that the improved feature is used

$$T_{\text{improved}} = \frac{T_{\text{affected}}}{\text{improvement factor}} + T_{\text{unaffected}}$$

- Example: multiplication accounts for 80s of 100s. How much improvement in multiplication performance to get 4x overall?

$$25 = \frac{80}{n} + 20$$

$$n = 16$$

multiplication needs to be 16x faster

Pitfall: use MIPS as a Performance Metric

- MIPS: Millions of Instructions Per Second

$$\begin{aligned}\text{MIPS} &= \frac{\text{Instruction count}}{\text{Execution time} \times 10^6} \\ &= \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}} \times 10^6 = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}\end{aligned}$$

- Question: A and B has different instruction set. A runs 5 MIPS, and B run 4 MIPS, which one is faster?

We don't know because we don't know how much time each instruction takes

Pitfall: MIPS as a Performance Metric

- MIPS: Millions of Instructions Per Second
 - Doesn't account for
 - Differences in ISAs between computers
 - Differences in complexity between instructions
- Note: CPI varies between programs on a given CPU, between various instruction sets
- See an example in the next slide

Example

Measurement	Computer A	Computer B
Instr. Count	10 billion	8 billion
Clock rate	4 GHz	4GHz
CPI	1.0	1.1

$$\text{MIPS} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}$$

- Which computer has **higher MIPS**?

$$\text{MIPS}_A = 4 / (1.0 \times 10^6)$$

$$\text{MIPS}_B = 4 / (1.1 \times 10^6)$$

A has **larger** MIPS

- Which is **faster**?

$$\text{CPU time}_A = 10 \times 10^9 \times 1.0 / 4 \times 10^9 = 2.5$$

$$\text{CPU time}_B = 8 \times 10^9 \times 1.1 / 4 \times 10^9 = 2.2$$

B is **faster**

Concluding Remarks

- **Cost/performance** is improving
 - Due to underlying technology development
- **Hierarchical** layers of abstraction
 - In both hardware and software
- **Instruction set** architecture
 - The hardware/software interface
- **Execution** time: the best performance measure
- **Power** is a limiting factor
 - Use parallelism to improve performance