

Computer Organization

計算機組織

Cache

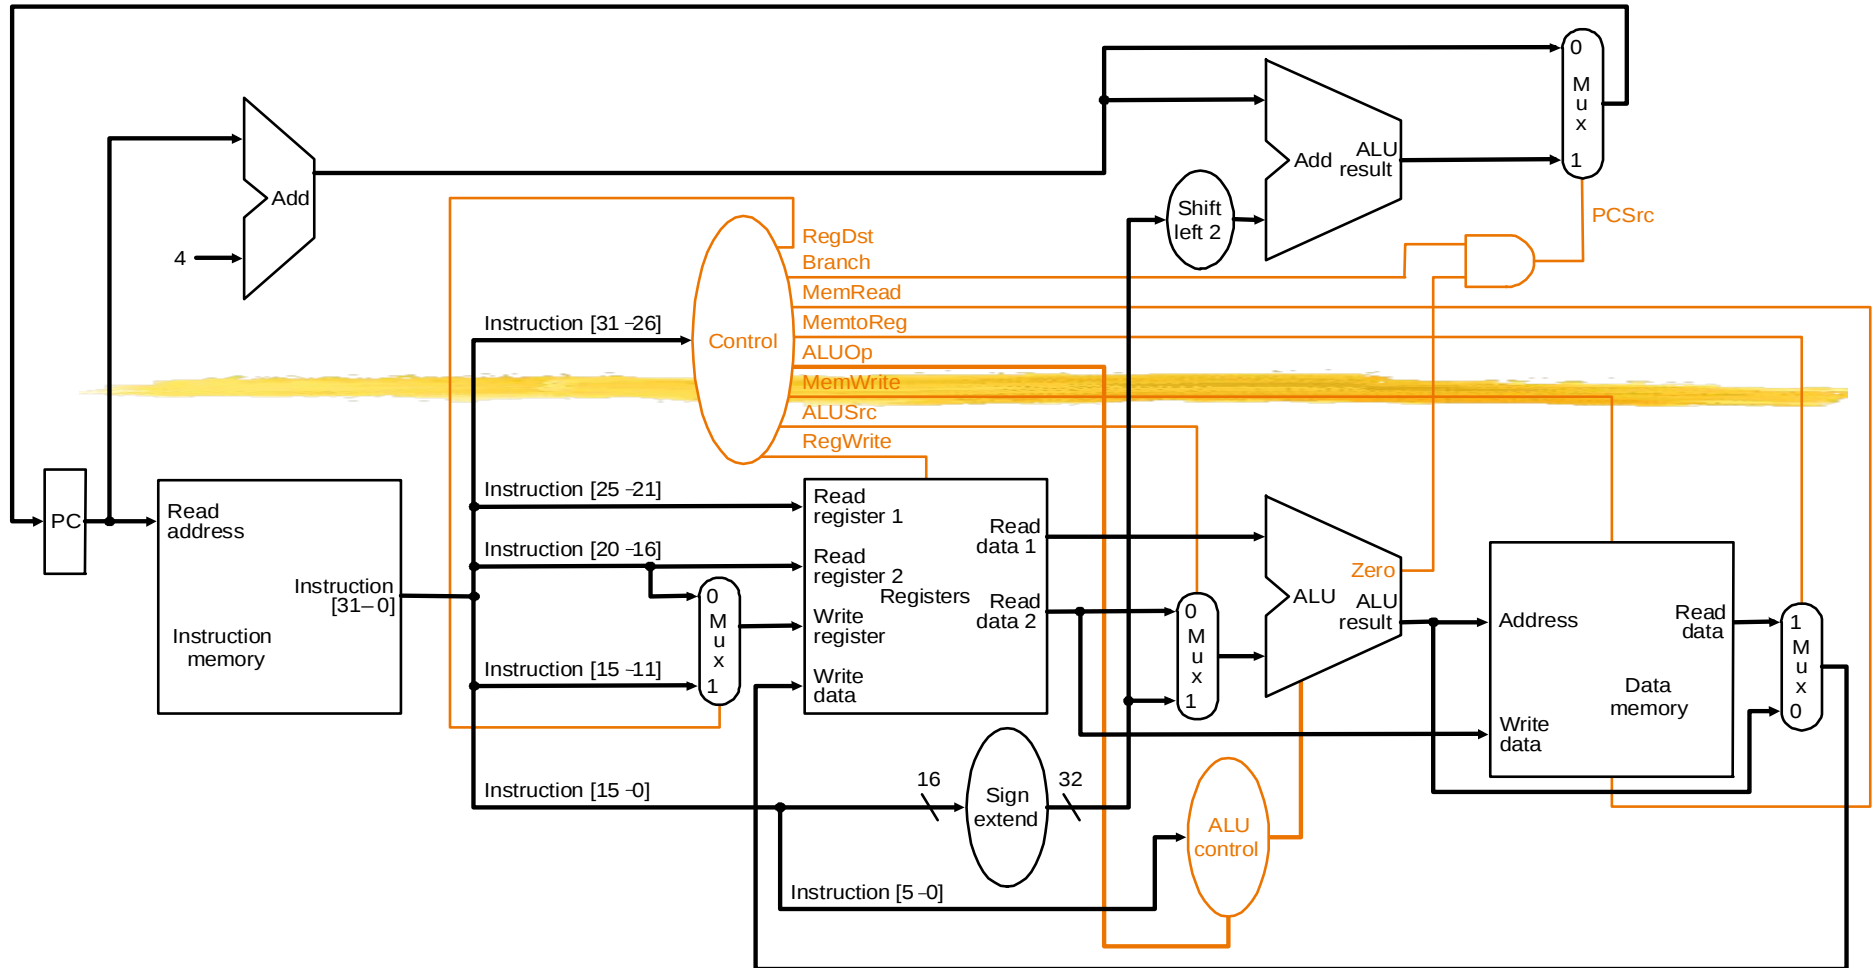


國立成功大學資訊工程學系
105 年度第二學期

What to Learn

- ◆ **Memory hierarchy**
- ◆ **Caching**
- ◆ **Cache organization**
 - **Direct mapped**
 - **Fully associative**
 - **Set associative**

Single-Cycle MIPS



Memory References Workload

Workload or
Benchmark
programs

Processor

reference stream

$\langle \text{op}, \text{addr} \rangle, \langle \text{op}, \text{addr} \rangle, \langle \text{op}, \text{addr} \rangle, \langle \text{op}, \text{addr} \rangle, \dots$

op: **instruction fetch**, read (**lw**), and write (**sw**)

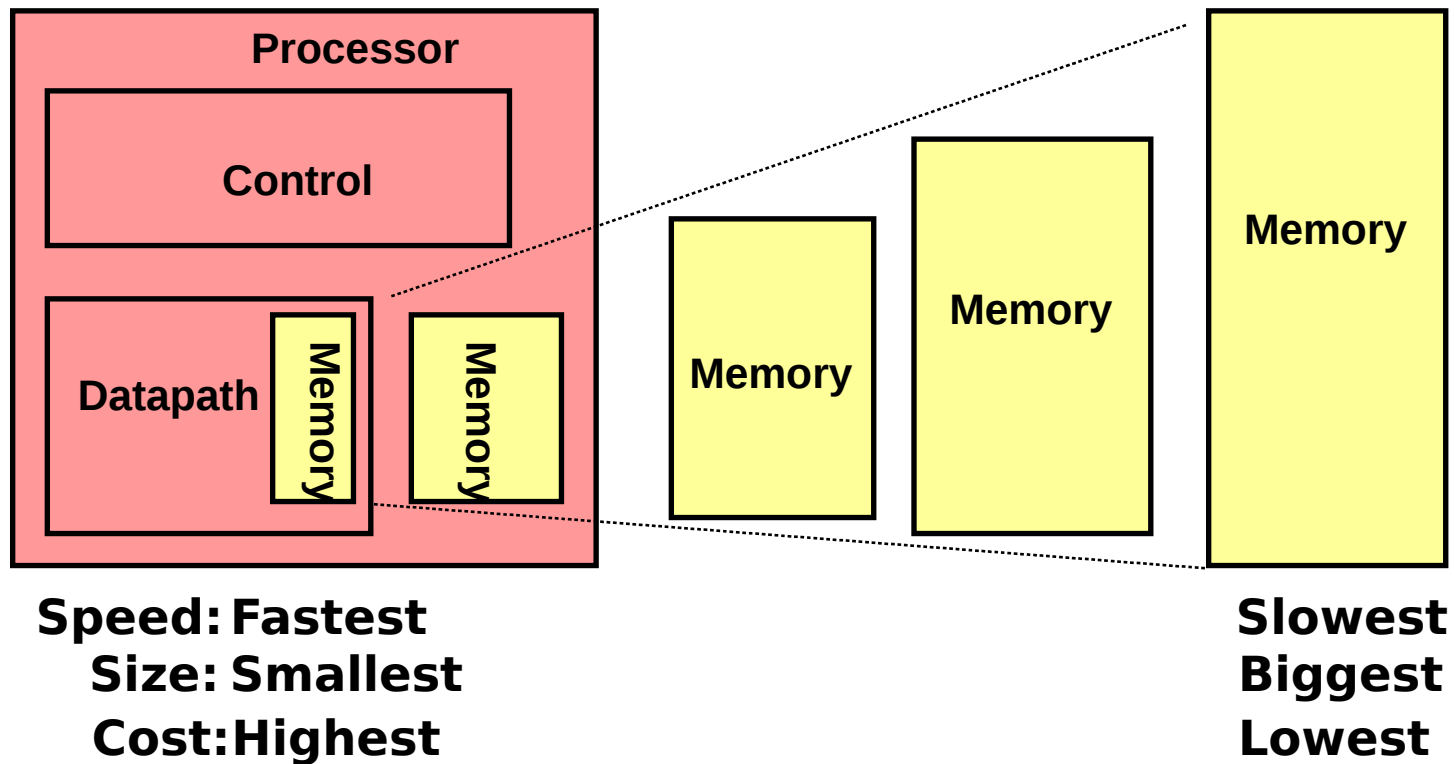
Memory

\$

Mem


Optimize the memory system organization
to minimize the average memory access time
for typical workloads

Memory Hierarchy: Concept



Memory Hierarchy: Terminology

- ◆ Hit: data appears in upper level (Block X)
 - **Hit rate (ratio)**: fraction of memory access found in the upper level
 - **Hit time**: time to access the upper level
 - Time to determine hit/miss + RAM access time
- ◆ Miss: data needs to be retrieved from a block in the lower level (Block Y)
 - **Miss rate** = $1 - (\text{Hit rate})$
 - **Miss penalty**: time to place (replace) a block (with a block) in the upper level + time to deliver the block to the processor
- ◆ **Hit Time** << **Miss Penalty**
- ◆ For simplicity:
 - Upper level memory: cache in processor
 - Lower level memory: main memory



Why Hierarchy Works?

◆ Principle of Locality:


- Program access a relatively small portion of the address space at any instant of time
- 90/10 rule (power law): 10% of code executes 90% of time

◆ Two types of locality:

- **Temporal locality**: if an item is referenced, it will tend to be referenced again soon
 - E.g., some counter variable in a loop
- **Spatial locality**: if an item is referenced, items whose addresses are close by tend to be referenced soon
 - E.g., nearby array data $A[i]$ and $A[i+1]$

Basics of Cache

- ◆ **Direct-mapped cache** (e.g., MIPS R2000)
 - Address mapping: modulo number of blocks

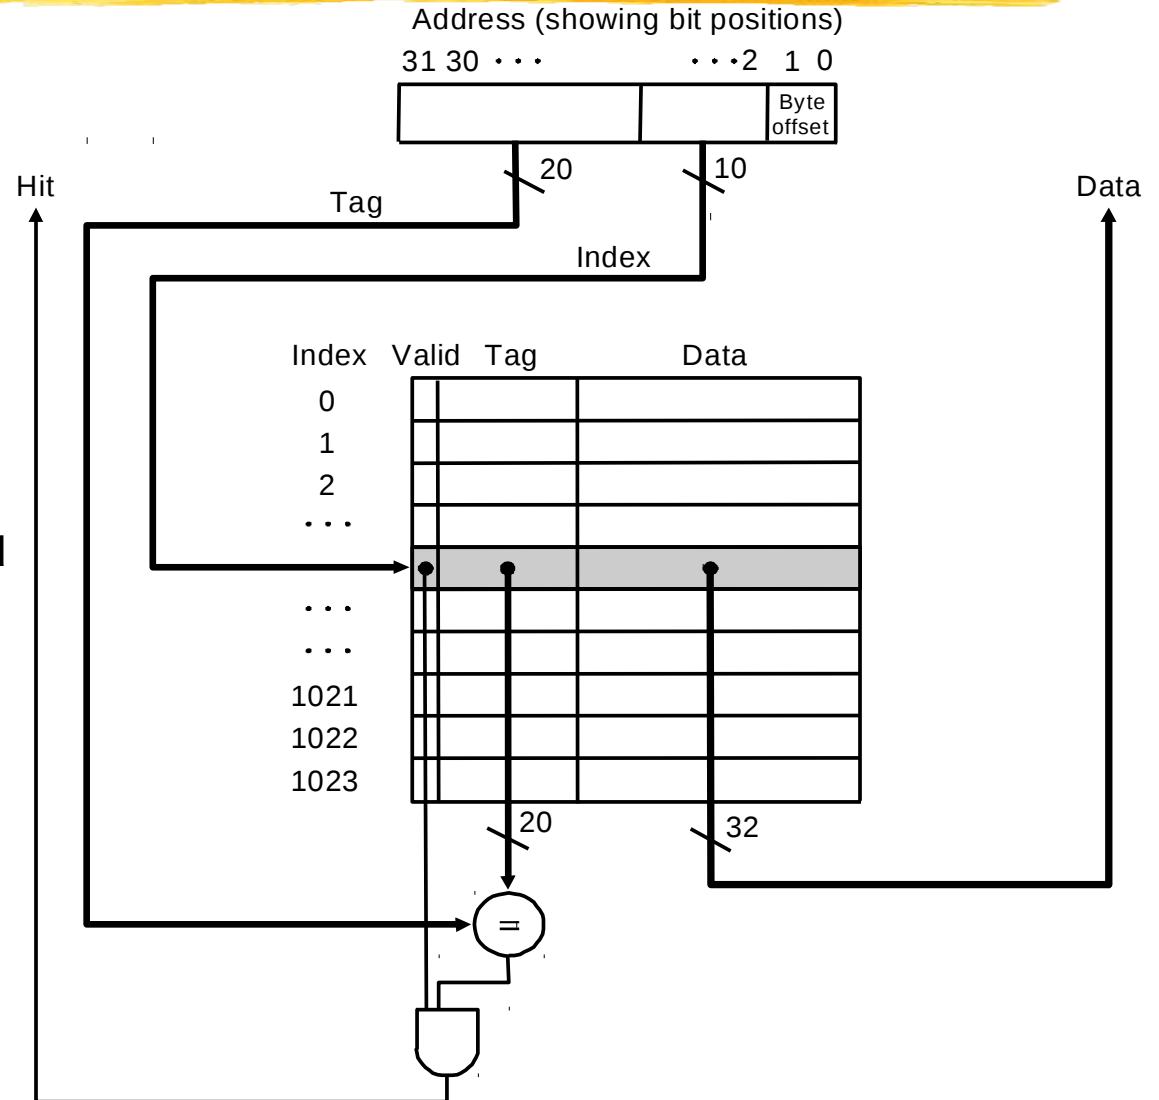


Direct-Mapped Cache

- ◆ Tag and valid bit

- ◆ Example

- 1K words, 1-word block:
- Cache index: lower 10 bits
- Cache **tag**: upper 20 bits
- **Valid bit** (when start up, valid is 0)



Example

(0) Initial state:

| Index | V | Tag | Data |
|-------|---|-----|------|
| 000 | N | | |
| 001 | N | | |
| 010 | N | | |
| 011 | N | | |
| 100 | N | | |
| 101 | N | | |
| 110 | N | | |
| 111 | N | | |

(1) Address referred 10110xx (*miss*):

| Index | V | Tag | Data |
|-------|---|-----|--------------|
| 000 | N | | |
| 001 | N | | |
| 010 | N | | |
| 011 | N | | |
| 100 | N | | |
| 101 | N | | |
| 110 | Y | 10 | Mem(10110xx) |
| 111 | N | | |

(2) Address referred 11010xx (*miss*):

| Index | V | Tag | Data |
|-------|---|-----|--------------|
| 000 | N | | |
| 001 | N | | |
| 010 | Y | 11 | Mem(11010xx) |
| 011 | N | | |
| 100 | N | | |
| 101 | N | | |
| 110 | Y | 10 | Mem(10110xx) |
| 111 | N | | |

(3) Address referred 10110xx (*hit*):

| Index | V | Tag | Data |
|-------|---|-----|--------------|
| 000 | N | | |
| 001 | N | | |
| 010 | Y | 11 | Mem(11010xx) |
| 011 | N | | |
| 100 | N | | |
| 101 | N | | |
| 110 | Y | 10 | Mem(10110xx) |
| 111 | N | | |



(4) Address referred 10010xx (*miss*):

| Index | V | Tag | Data |
|-------|---|-----|--------------|
| 000 | N | | |
| 001 | N | | |
| 010 | Y | 10 | Mem(10010xx) |
| 011 | N | | |
| 100 | N | | |
| 101 | N | | |
| 110 | Y | 10 | Mem(10110xx) |
| 111 | N | | |



Example Problem (1/2)

- ◆ How many total bits are required for a direct-mapped cache with 128 KB of data and 1-word block size, assuming a 32-bit address?
 - Cache data = 128 KB = 2^{17} bytes = 2^{15} words = 2^{15} blocks
 - Cache entry size = block data bits + tag bits + valid bit = $32 + (32 - 15 - 2) + 1 = 48$ bits
 - Therefore, cache size = $2^{15} \times 48$ bits = $2^{15} \times (1.5 \times 32)$ bits = 1.5×2^{20} bits = 1.5 Mbits

Example Problem (2/2)

- ◆ Consider a direct-mapped cache with 64 blocks and a block size of 16 bytes. What block number does byte address 1200 map to?

- As block size = 16 bytes:

byte address 1200 \Rightarrow block address $\lfloor 1200/16 \rfloor = 75$

- As cache size = 64 blocks:

block address 75 \Rightarrow cache block $(75 \bmod 64) = 11$

More on Hits

- ◆ Read hits: this is what we want!
- ◆ Write hits: keep cache/memory consistent?
 - **Write-through**: write to cache and memory at same time
=> but memory is very slow!
 - **Write-back**: write to cache only (write to memory when that block is being replaced)
 - Need a **dirty bit** for each block

More on Misses

◆ Read misses

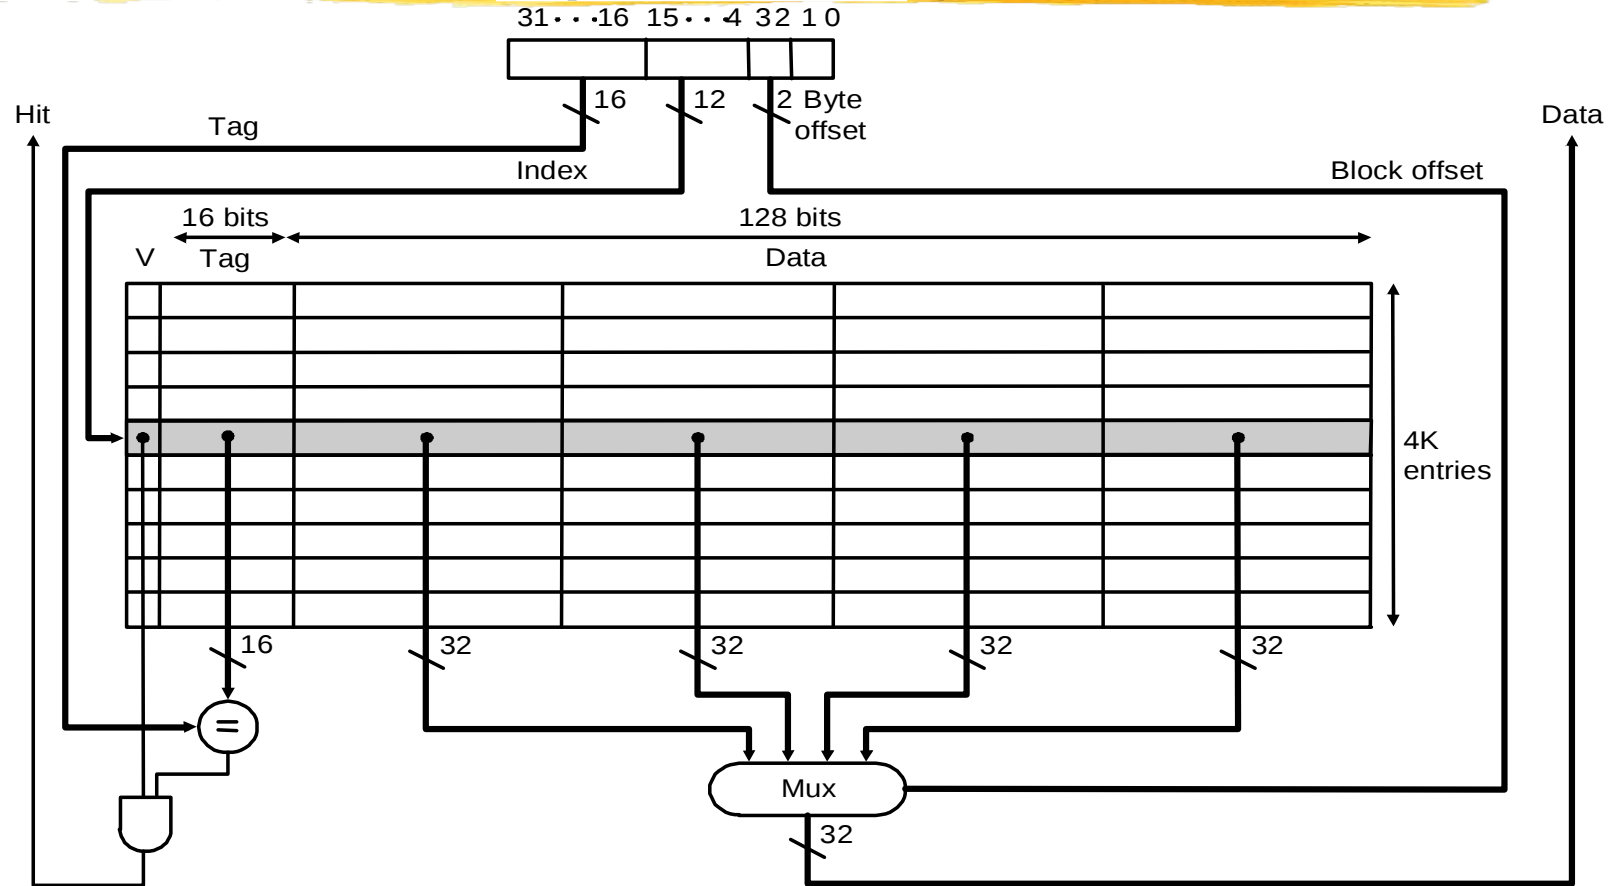
- Stall CPU, freeze register contents, fetch block from memory, deliver to cache, restart
- “**Block replacement**” may occur
 - Placement: ?
 - Replacement: ?

◆ Write misses

- Write-allocated: read block into cache, write the word
 - Block replacement may occur
- Write-non-allocate: write directly into memory

How to Exploit Spatial Locality?

◆ Increase block size for spatial locality






Cache with 4K 4-word blocks: *byte offset* (least 2 significant bits) is ignored, next 2 bits are *block offset*, and the next 12 bits are used to index into cache

Possible Cache Architectures

- ◆ **Direct mapped** (have discussed)
- ◆ **Fully associative**: each memory block can locate anywhere in cache
 - all cache entries are searched (in parallel) to locate block
- ◆ **Set associative**: each memory block can place in a unique set of cache locations (any location in the set), and if the set is of size n it is n -way set-associative
 - cache set address = memory block address *mod* number of sets in cache
 - all cache entries in the corresponding set are searched (in parallel) to locate block
 - E.g., in a 2-way 1K-word set associate cache, there are 512 sets, each having 2 words

Illustration



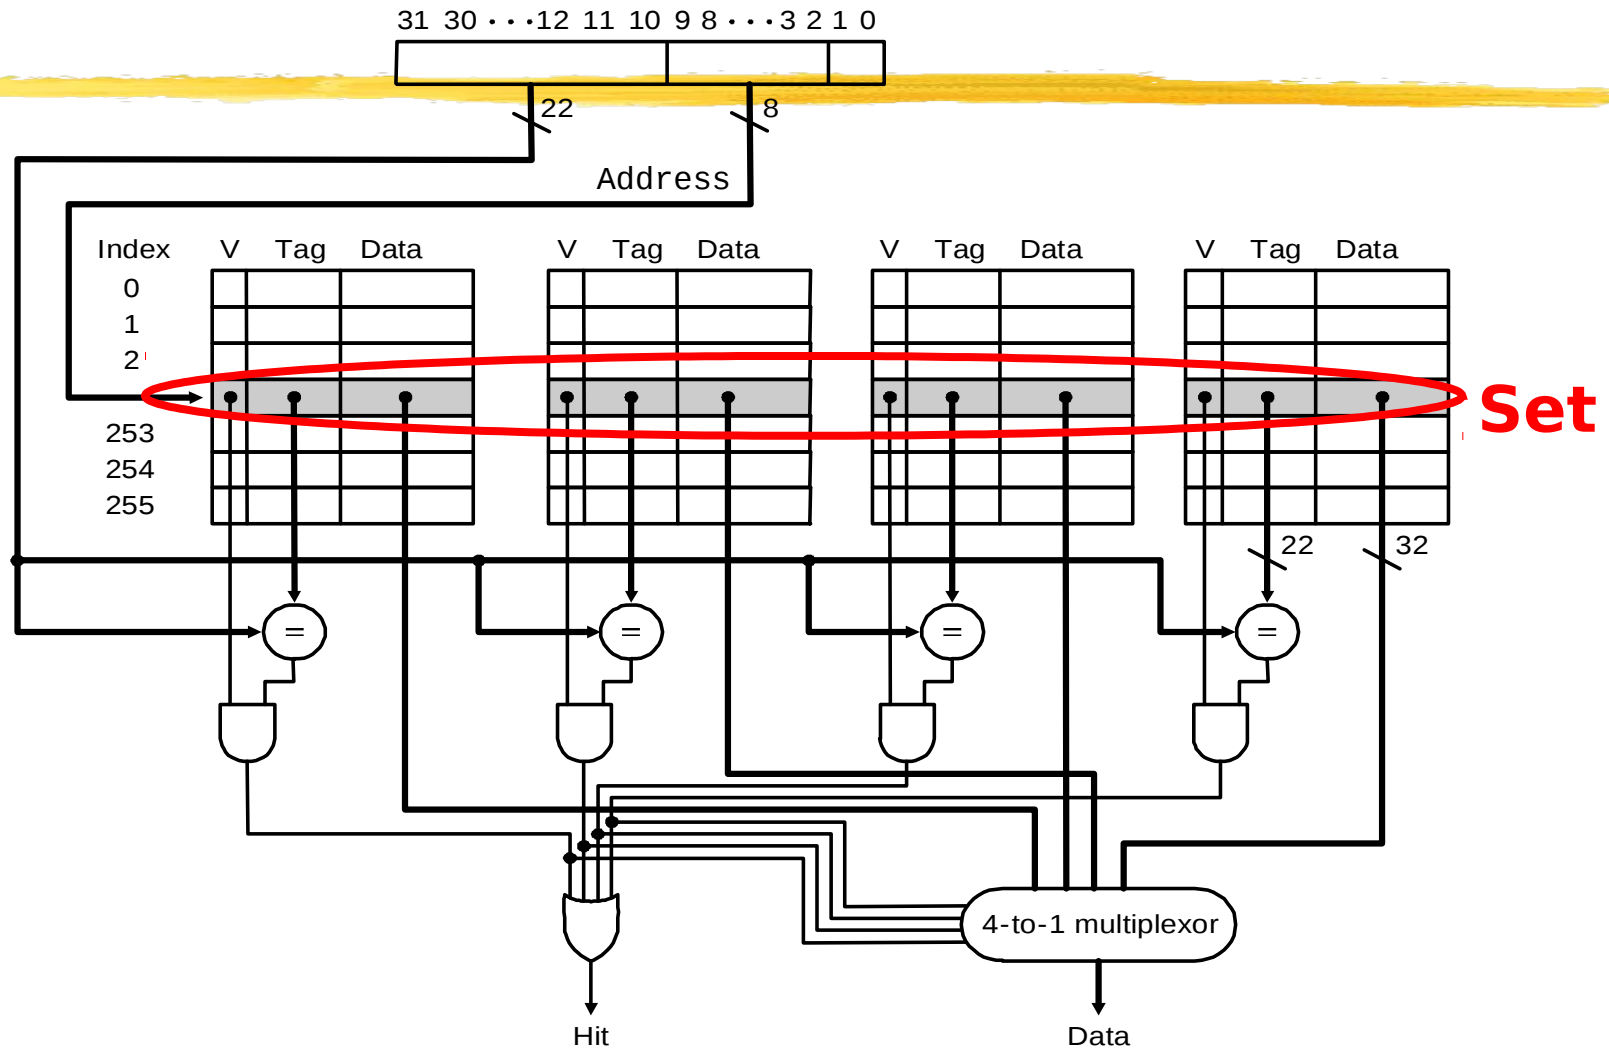
Location of a memory block with address 12 in a cache with 8 blocks with different degrees of associativity

More on Associativity

Assume N cache blocks

- ◆ K -way set associativity = K directed mapped
 - No. of sets = N / K
- ◆ Fully associativity = N directed mapped
 - No. of sets = $N / N = 1$


Implementing Set-Associative Cache



4-way set-associative cache with 4 comparators and one 4-to-1 multiplexor: size of cache is 1K blocks = 256 sets * 4-block set size

Implementing Fully Associate Cache

- ◆ Compare cache tags of all cache entries in parallel
- ◆ Ex.: Block Size = 8 words, N 27-bit comparators



Why Set Associate Cache?



- ◆ **Directed mapped may introduce too many conflict misses**
 - **Conflict miss: ≥ 1 memory blocks contend a cache line**
- ◆ **Fully associative is too expensive**
 - **Minimize the number of conflict misses**
- ◆ **Set associative intends to come up with a balance**

Example Problem

- ◆ Find the number of misses for a cache with **four 1-word blocks** given the following sequence of **word addresses** of memory block accesses: **0, 8, 0, 6 and 8**, for each of the following cache configurations
 1. direct mapped
 2. 2-way set associative (use LRU replacement policy)
 3. fully associative

- ◆ LRU replacement
 - If replacement, replace with a least recently used block
 - In a 2-way set associative cache **LRU replacement** can be implemented with one bit at each set whose value indicates the mostly recently referenced block

Direct Mapped

Block address translation in direct-mapped cache

| Block address | Cache block |
|---------------|---------------------|
| 0 | 0 ($= 0 \bmod 4$) |
| 6 | 2 ($= 6 \bmod 4$) |
| 8 | 0 ($= 8 \bmod 4$) |

Cache contents after each reference - red indicates new entry added

| Address of memory block accessed | Hit or miss | Contents of cache blocks after reference | | | |
|----------------------------------|-------------|------------------------------------------|---|-----------|---|
| | | 0 | 1 | 2 | 3 |
| 0 | miss | Memory[0] | | | |
| 8 | miss | Memory[8] | | | |
| 0 | miss | Memory[0] | | | |
| 6 | miss | Memory[0] | | Memory[6] | |
| 8 | miss | Memory[8] | | Memory[6] | |

5 misses

Two-way Set Associate

Block address translation in a two-way set-associative cache

| Block address | Cache set |
|---------------|---------------------|
| 0 | 0 ($= 0 \bmod 2$) |
| 6 | 0 ($= 6 \bmod 2$) |
| 8 | 0 ($= 8 \bmod 2$) |

Cache contents after each reference - red indicates new entry added

| Address of memory block accessed | Hit or miss | Contents of cache blocks after reference | | | |
|----------------------------------|-------------|------------------------------------------|-----------|-------|-------|
| | | Set 0 | Set 0 | Set 1 | Set 1 |
| 0 | miss | Memory[0] | | | |
| 8 | miss | Memory[0] | Memory[8] | | |
| 0 | hit | Memory[0] | Memory[8] | | |
| 6 | miss | Memory[0] | Memory[6] | | |
| 8 | miss | Memory[8] | Memory[6] | | |

4 misses

Fully Associate

Cache contents after each reference - red indicates new entry added

| Address of memory block accessed | Hit or miss | Contents of cache blocks after reference | | | |
|----------------------------------|-------------|------------------------------------------|-----------|-----------|---------|
| | | Block 0 | Block 1 | Block 2 | Block 3 |
| 0 | miss | Memory[0] | | | |
| 8 | miss | Memory[0] | Memory[8] | | |
| 0 | hit | Memory[0] | Memory[8] | | |
| 6 | miss | Memory[0] | Memory[8] | Memory[6] | |
| 8 | hit | Memory[0] | Memory[8] | Memory[6] | |

3 misses