

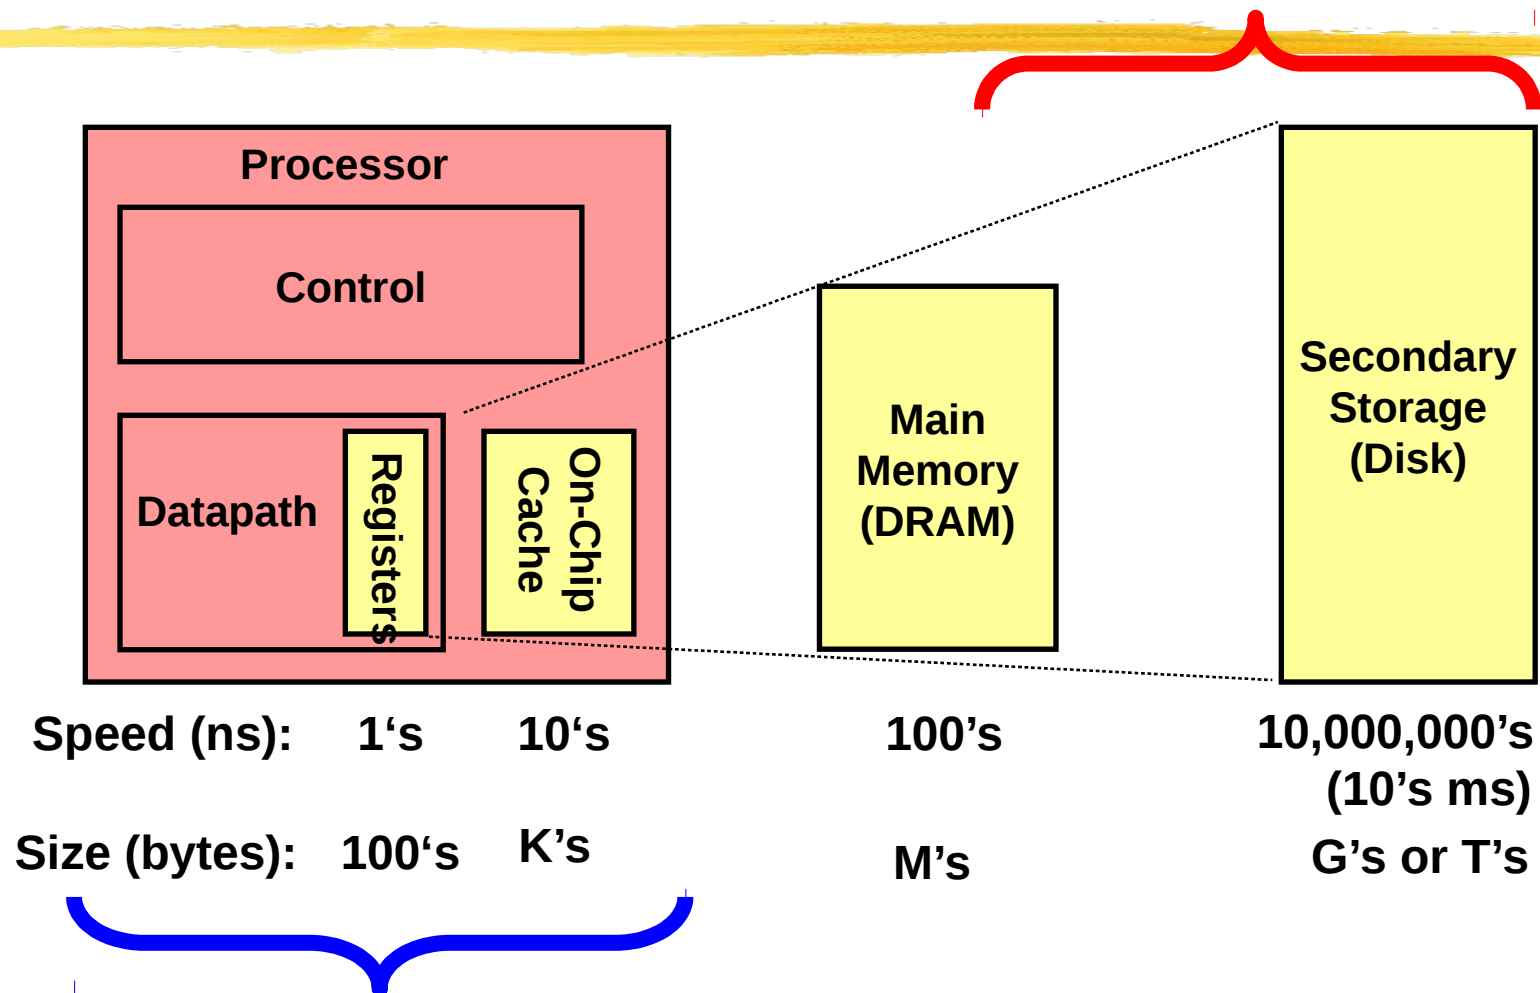
Computer Organization 計算機組織

Virtual Memory



國立成功大學資訊工程學系
105 年度第二學期

We now go to this part

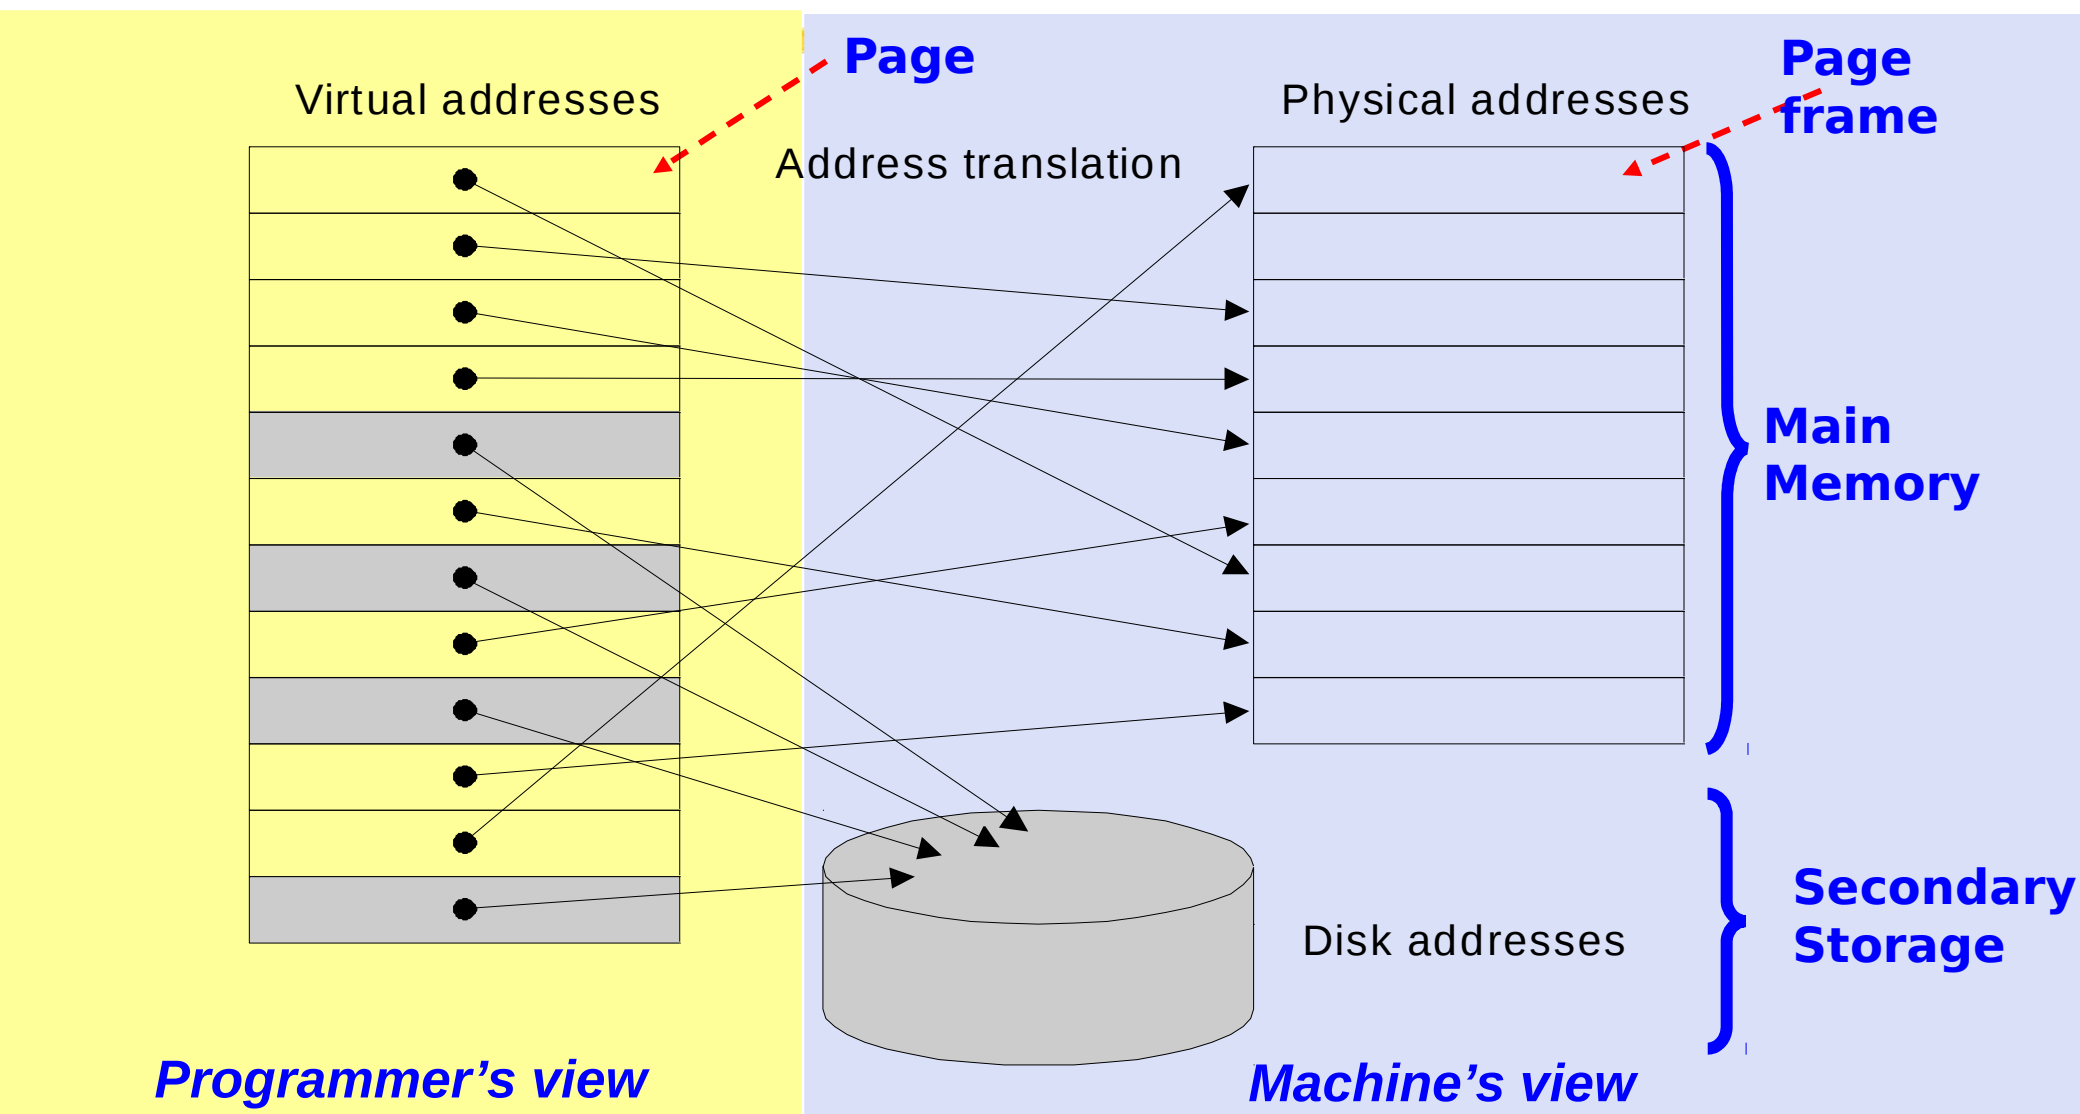


So far the memory system we learn...

Virtual Memory (1/2)

- ◆ **Allow many programs to run at once**
- ◆ **Virtual memory**
 - Every program has its own address space (can be greater than physical memory), starting at address 0
 - only accessible to itself
 - can run anywhere in physical memory
- ◆ **Virtual memory is a technique implementing the translation from virtual space to physical space**
 - OS runs all the time and allocates physical resources

Virtual Memory (2/2)



Mapping of pages from a virtual address to a physical address or disk address

Issues in Virtual Memory

- ◆ **Size of a page**

Page: a basic unit of memory blocks transferring between main memory and secondary storage (e.g., disk)

- ◆ **Placement policy**

where to place a page in *main memory*

- ◆ **Replacement policy**

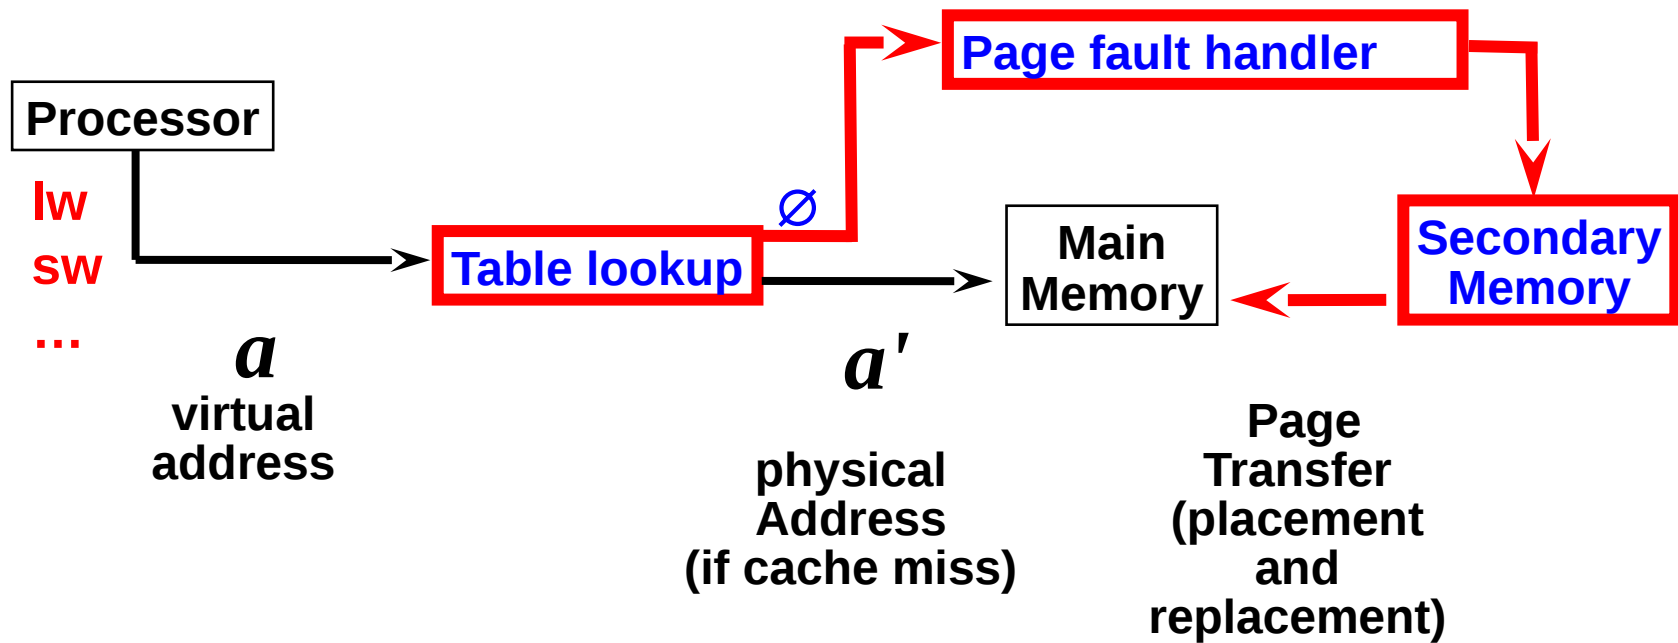
replace pages if *cannot* find any free space in main memory

- ◆ **When to fetch a page?**

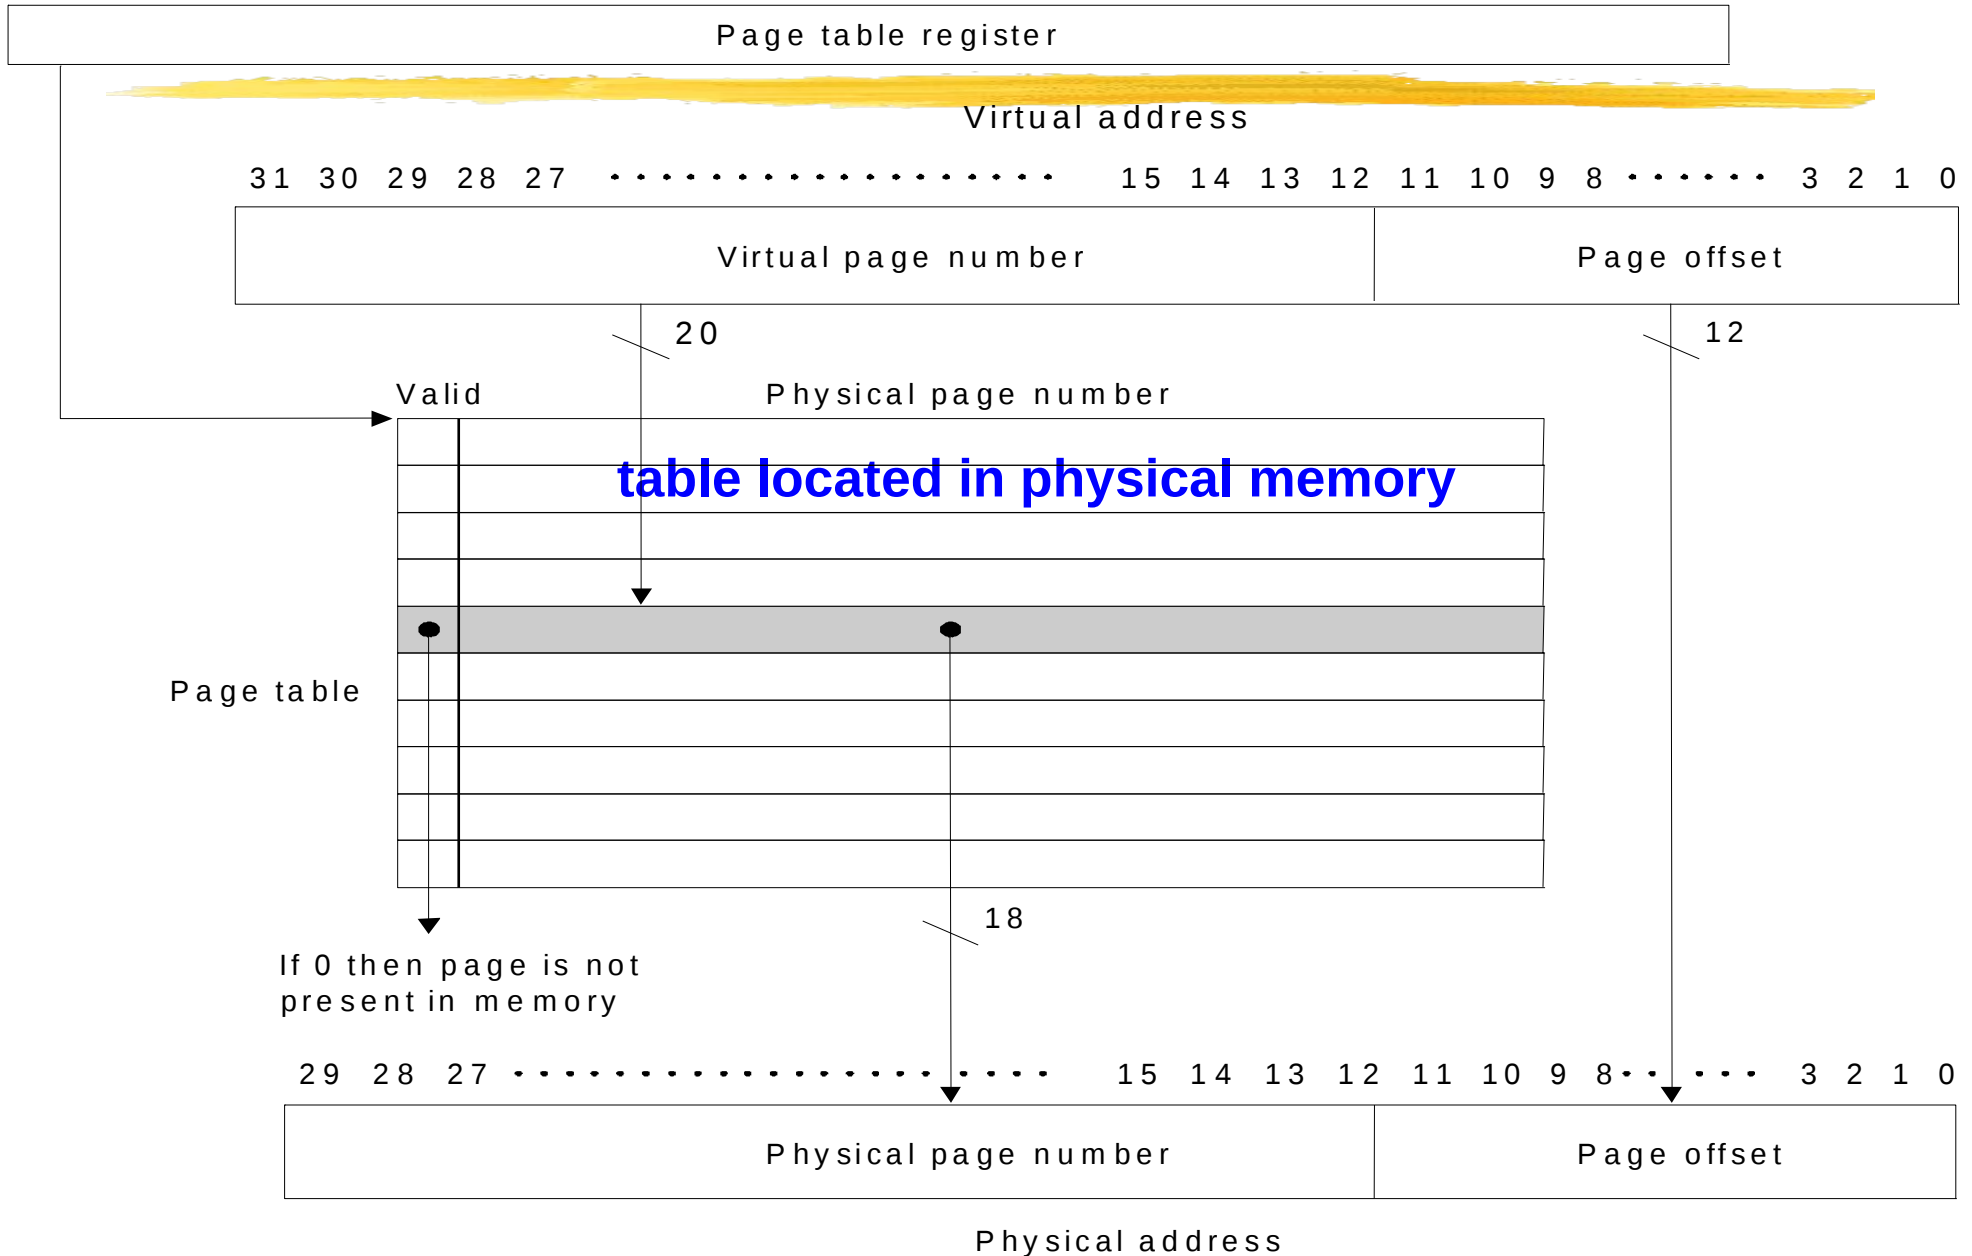
may fetch when we miss

Address Translation: Virtual to Physical

- Given a virtual address a , return the physical address (a') corresponding to a



Paging with 4K-byte Page Size (assume 1GB main memory)



Example Problem

- ◆ Assume:
 - 32-bit virtual address
 - 4 KB page size
 - 4 bytes per page table entry
 - The page table maps page 0 to frame 10, page 1 to frame 8, page 2 to frame 100, page 3 to frame 1, ...
- ◆ Q1: What is the total page table size if we want to be able to access all of the virtual memory?
- ◆ Q2: What is the page no. of the virtual address 8196?
- ◆ Q3: What is the physical address corresponding to the virtual address 8196?

Solution

- ◆ No. of page table entries = address space size / page size = $2^{32} / 2^{12} = 2^{20}$
- ◆ (Maximum) Size of page table = No. of entries \times entry size = $2^{20} \times 4 \text{ bytes} = 4 \text{ MB}$
 - However, a program typically takes a few of pages
- ◆ $8196/4096=2.x$ (i.e. address 8196 is at the 2nd page)
- ◆ Physical address is $100 \times 4096 + 4$

Page Fault

- ◆ Page fault means that page is **not** resident in memory
- ◆ Then, hardware must trap to the operating system so that it can remedy the situation
 - “May” need to pick a page to discard (may write the replaced page to disk)
 - i.e., **page replacement policy**
 - **Note: if a page in “main memory” is updated, then it is “dirty”**
 - Load the requested page in from disk
 - i.e., **placement policy**
 - Update the page table
 - Resume to program so HW will retry and succeed!

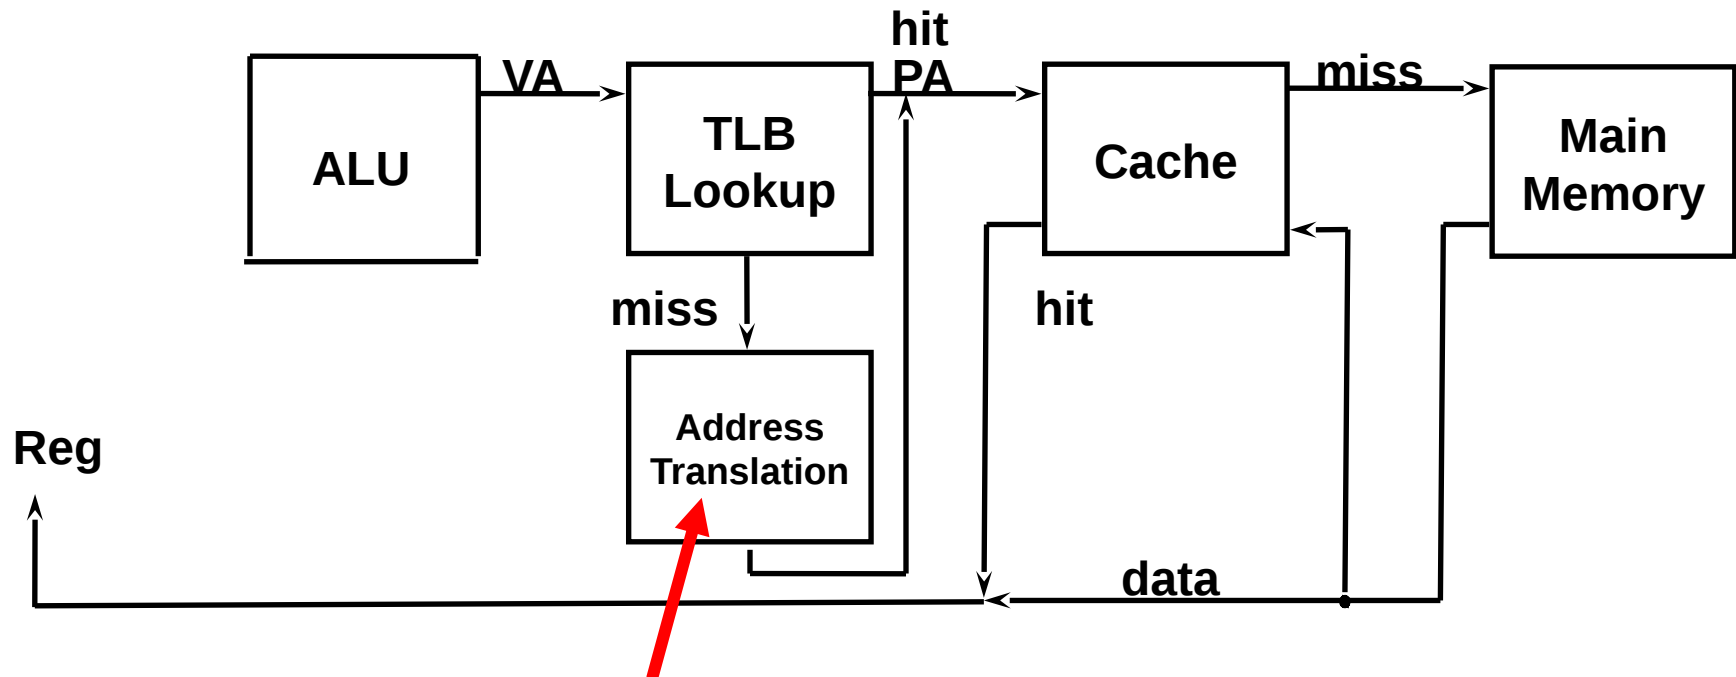
So Far



- ◆ **We assume that we are given a virtual memory address to access the virtual memory system**
- ◆ **We have not integrate virtual memory with processor, i.e., how does a processor send out a virtual memory address due to lw or sw?**
- ◆ **Put everything together in the following...**

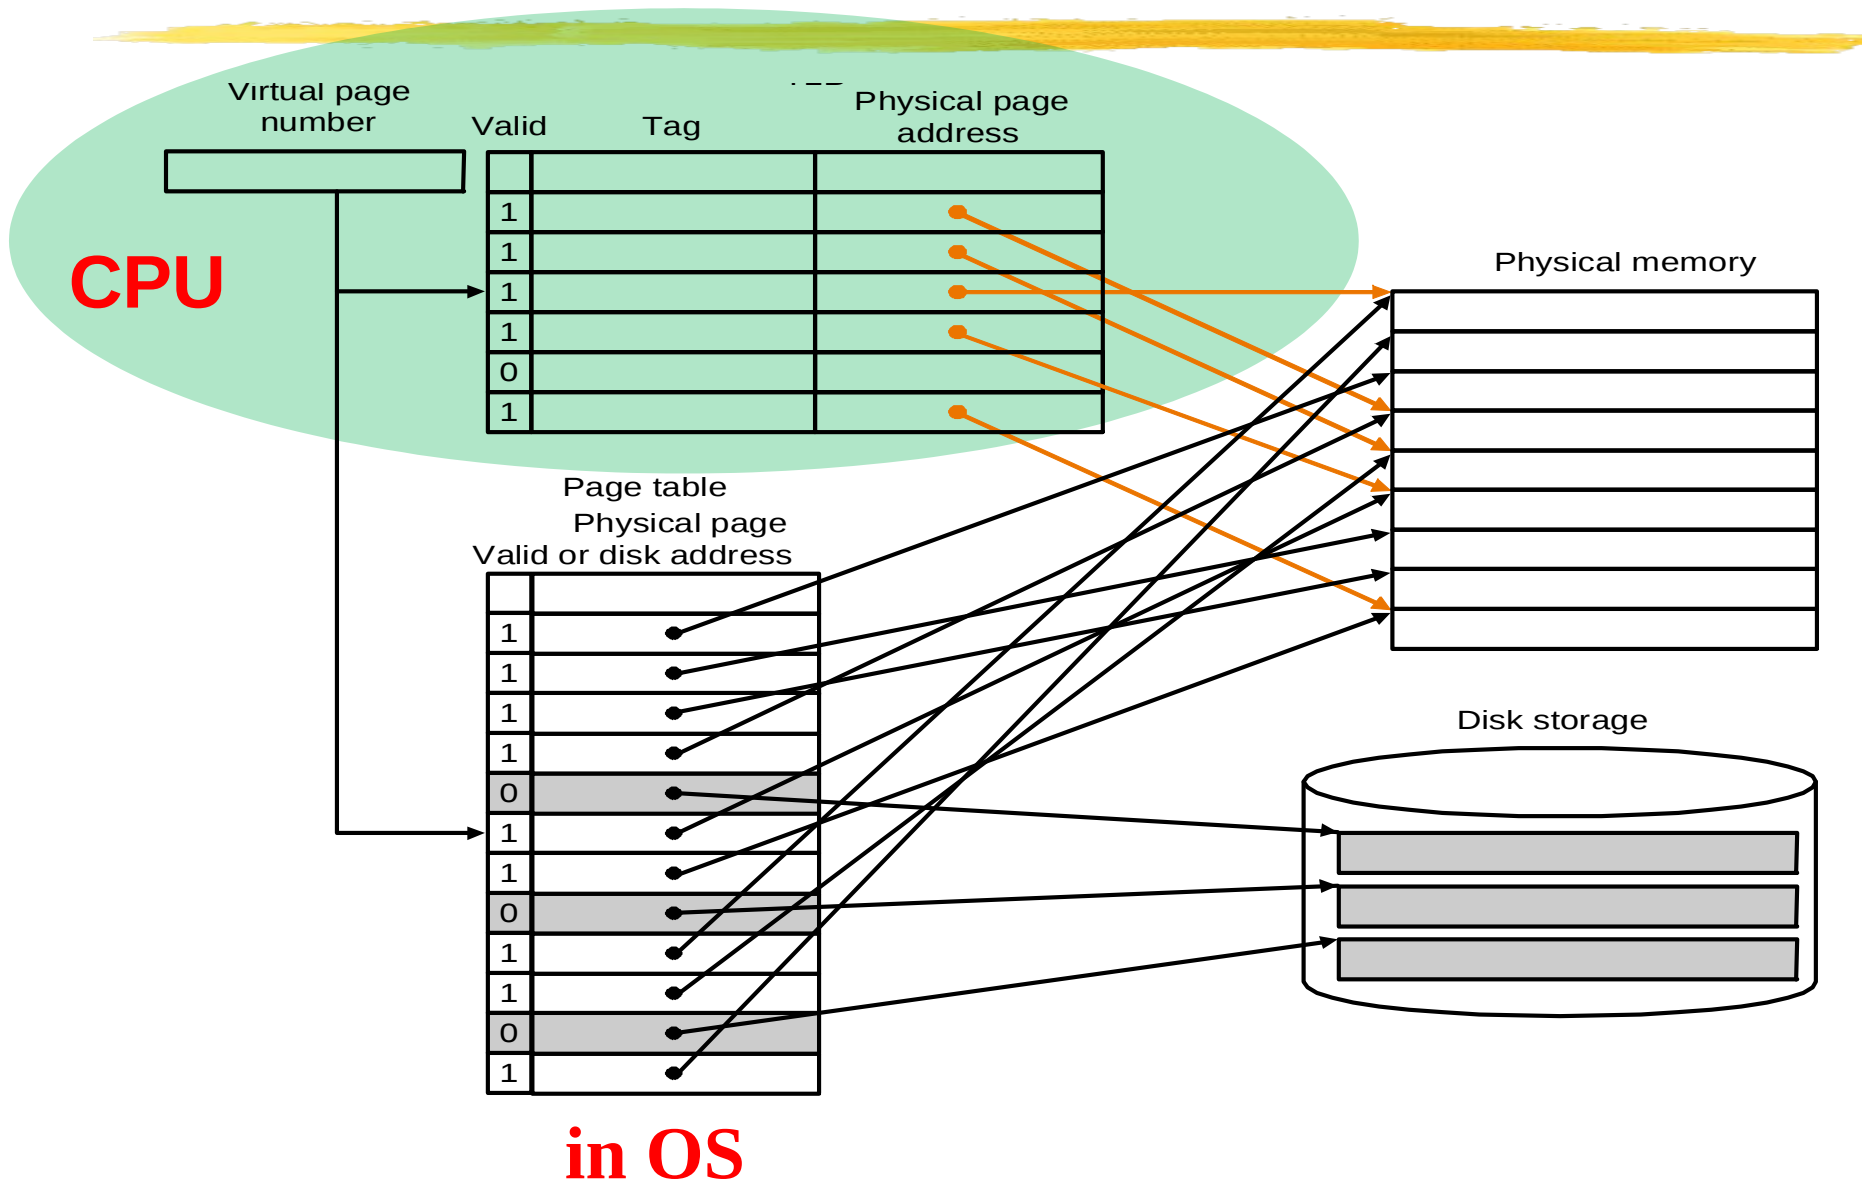
Practical Address Translation

- ◆ To speedup the address translation, typical processor implements a hardware page table cache, namely **Translation Lookaside Buffer (TLB)**



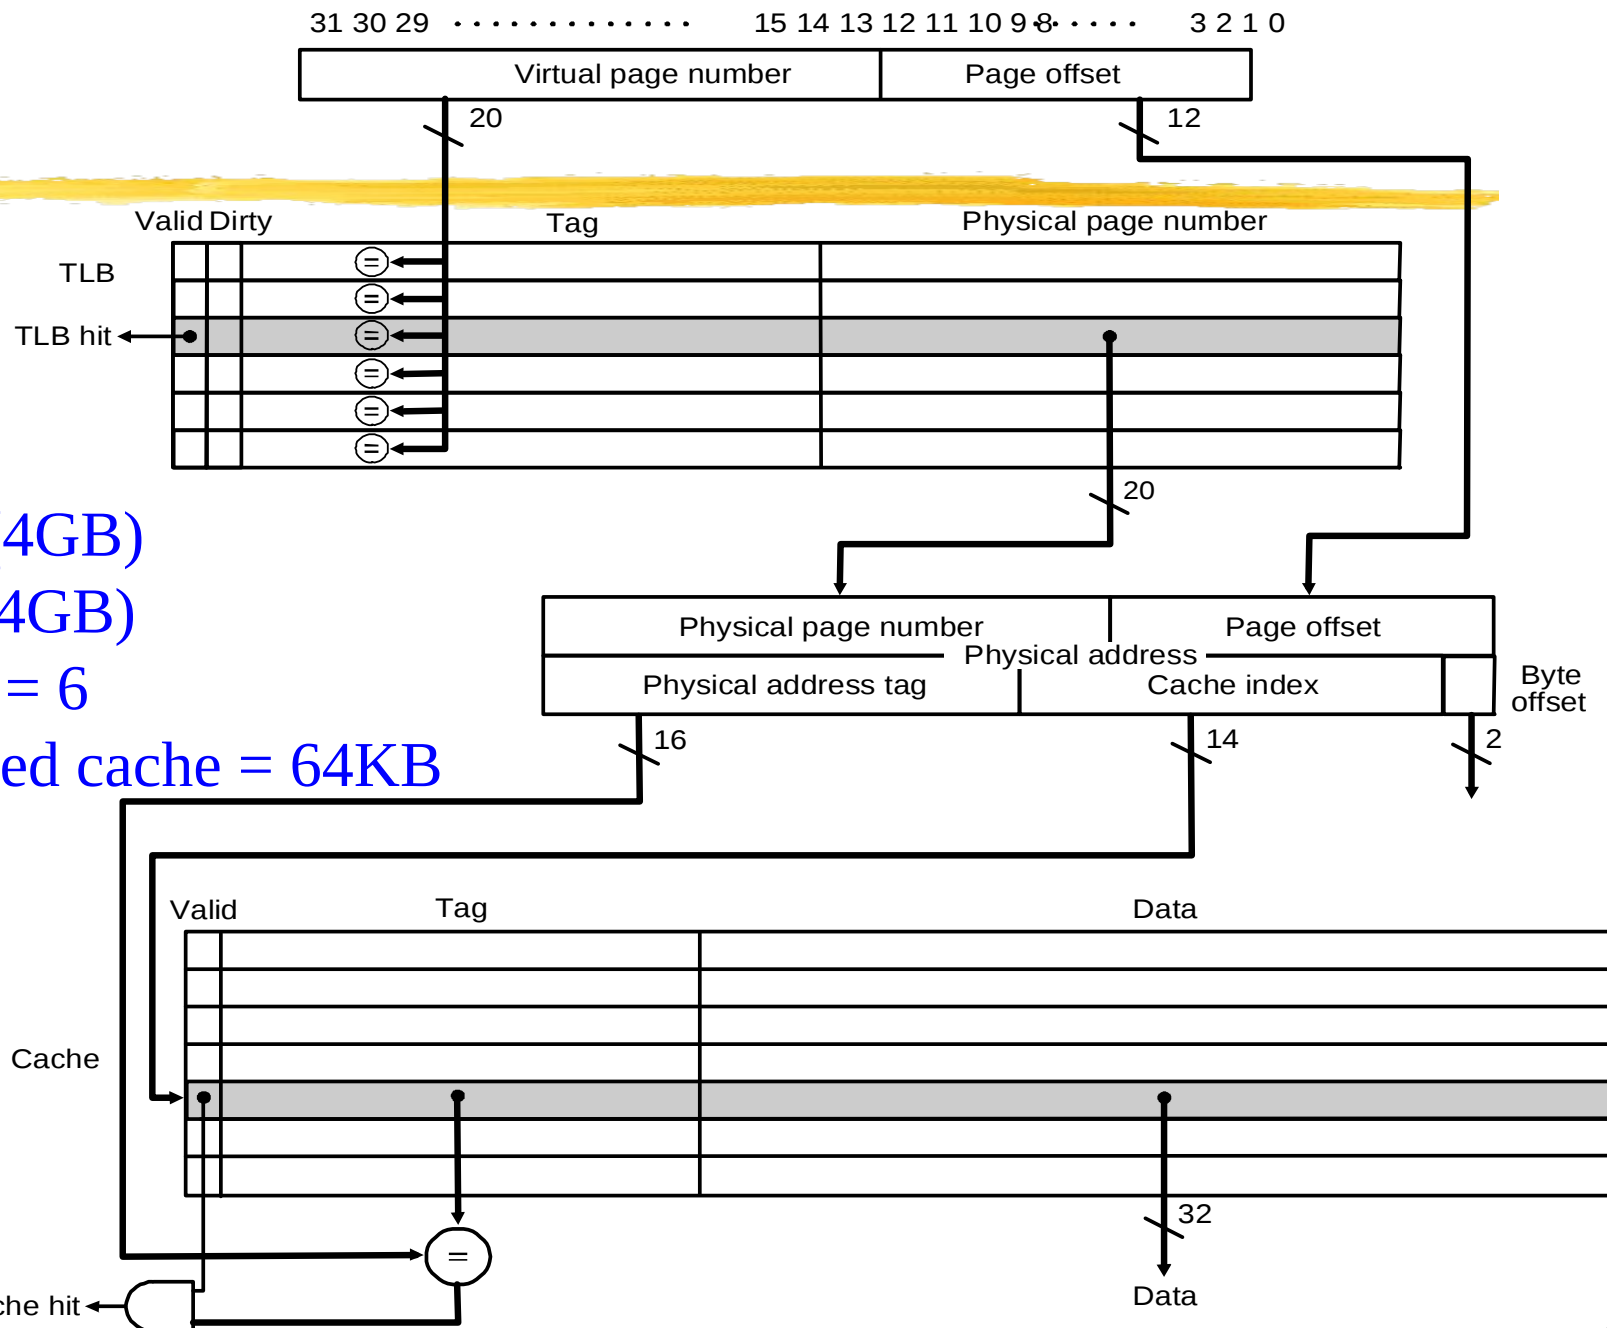
“May” invoke page fault handler

TLB




More on TLB

- ◆ TLB can be organized as fully associative, set associative, or direct mapped
 - TLBs are typically small in size, typically $< 128\sim 256$ entries
- ◆ TLB hit on write:
 - Toggle dirty bit (write back to page table on replacement)
- ◆ TLB miss:
 - If only TLB miss \Rightarrow load “page table entry” into TLB
 - If page fault also \Rightarrow OS exception
- ◆ Inclusion property
 - If L1 is a upper level “data” memory of L2, then “data” in L1 must appear in L2



- 32-bit V.A. (4GB)
- 32-bit P.A. (4GB)
- TLB entries = 6
- Direct mapped cache = 64KB

TLB in Pipeline



inst. flow

TLB + Cache

Cache	TLB	Virtual Memory	Possible? Conditions?
Miss	Hit	Hit	Yes; but page table never checked if TLB hits
Hit	Miss	Hit	TLB miss, but entry found in page table; after retry, data in cache
Miss	Miss	Hit	TLB miss, but entry found in page table; after retry, data miss in cache
Miss	Miss	Miss	TLB miss and is followed by a page fault; after retry, data miss in cache
Miss	Hit	Miss	impossible; not in TLB if page not in memory
Hit	Hit	Miss	impossible; not in TLB if page not in memory
Hit	Miss	Miss	impossible; not in cache if page not in memory