

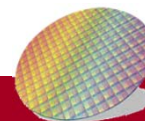


成功大學

National Cheng Kung University

# Chapter 1


Computer Abstractions and Technology



# Course Administration



---

- Course Website:
- See [NCKU Moodle](#) Website
- Homework 1: Due: 3/7 11:55 PM
- Reading: Chapter 1



# The Computer Revolution

- Computers are almost everywhere:
  - PC, Mobile Phone, embedded system, etc
- Makes novel applications feasible
  - Internet
  - Search Engines
  - Vehicle Navigation
  - Human genome project





# Moore's Law

- In 1965, Gordon Moore predicted that the number of transistors that can be integrated on a die would double every 18 to 24 months (i.e., grow exponentially with time).




Gordon Moore




# Classes of Computers

- **Personal** computers
  - General purpose, variety of software
- **Server** computers
  - **Network** based
  - High **capacity**, performance, reliability
  - Range from small servers to building sized
  - Extreme High performance: **Supercomputer**
- **Embedded** computers
  - Hidden as components of systems
  - Stringent power/performance/cost constraints




# PostPC Era


- Personal mobile device (PMD)
  - Battery operated with wireless connectivity to the Internet
  - Likely have touch-sensitive screen
  - Normally do not have keyboard and mouse
  - Different appearance
- Cloud Computing
  - Companies can rent portions of servers so that they can provide software services



# Device Growth



The fast growth of tablet and smart reflect the PostPC era.






## $2^X$ vs. $10^Y$ bytes ambiguity

- 1KB used to have two meanings :  $2^{10}$  Bytes or  $10^3$  bytes
  - Create confusion
  - Add a **binary notation** to remove confusions
  - I.e. Kilo vs. kibi
- **Last column** shows how much **larger** the binary term is than its corresponding decimal term.

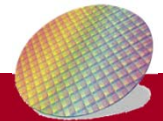
Decimal term	Abbreviation	Value	Binary term	Abbreviation	Value	% Larger
kilobyte	KB	$10^3$	kibibyte	KiB	$2^{10}$	2%
megabyte	MB	$10^6$	mebibyte	MiB	$2^{20}$	5%
gigabyte	GB	$10^9$	gibibyte	GiB	$2^{30}$	7%
terabyte	TB	$10^{12}$	tebibyte	TiB	$2^{40}$	10%
petabyte	PB	$10^{15}$	pebibyte	PiB	$2^{50}$	13%
exabyte	EB	$10^{18}$	exbibyte	EiB	$2^{60}$	15%
zettabyte	ZB	$10^{21}$	zebibyte	ZiB	$2^{70}$	18%
yottabyte	YB	$10^{24}$	yobibyte	YiB	$2^{80}$	21%





# What You Will Learn

- How **programs** are translated into the **machine** language
  - And how the hardware executes them
- The **hardware/software** interface
  - Instructions
- What determines program **performance**
  - How to compare the performance
- How hardware designers **improve performance**
- What are the reasons for the consequences of the recent switch from **sequential processing** to **parallel processing**?



# Eight Great Ideas

Design for  
**Moore's Law**



Use **abstraction**  
to simplify design



Make the **common**  
**case fast**



Performance  
via **parallelism**



Performance  
via **pipelining**



Performance  
via **prediction**



**Hierarchy** of  
memories




**Dependability**  
via redundancy




# Below Your Program

---



- Application software
  - Written in high-level language (HLL)
- System software
  - **Compiler**: translates HLL code to machine code
  - **Operating System**: service code
    - Handling input/output
    - Managing memory and storage
    - Scheduling tasks & sharing resources
- Hardware
  - Processor, memory, I/O controllers





# Levels of Program Code

## High-level language

Level of abstraction closer to problem domain  
Provides for productivity and portability

## Assembly language

Textual representation of instructions

## Hardware representation (Machine code)

Binary digits (bits)  
Encoded instructions and data

High-level  
language  
program  
(in C)

```
swap(int v[], int k)
{int temp;
  temp = v[k];
  v[k] = v[k+1];
  v[k+1] = temp;
```

Compiler


Assembly  
language  
program  
(for MIPS)

```
swap:
  muli $2, $5, 4
  add  $2, $4, $2
  lw   $15, 0($2)
  lw   $16, 4($2)
  sw   $16, 0($2)
  sw   $15, 4($2)
  jr   $31
```

Assembler

Binary machine  
language  
program  
(for MIPS)

```
000000001010000100000000000011000
000000000000110000001100000100001
100011000110001000000000000000000
100011001111001000000000000000100
101011001111001000000000000000000
101011000110001000000000000000100
00000011111000000000000000001000
```



# Components of a Computer

- **Five** components



- Input
- Output
- Memory
- Control
- Datapath

- **Almost Same** components for all kinds of computer

- Desktop, server, embedded

- Input/output includes

- User-interface devices
  - **Display**, keyboard, mouse
- Storage devices
  - **Hard disk**, CD/DVD, flash
- Network adapters
  - For **communicating** with other computers



# Anatomy of a Computer

---


Output device




Network cable

Input device



Input device



# Opening the Box - Macbook





Hard drive Processor Fan with cover Spot for memory DIMMs Spot for battery Motherboard Fan with cover DVD drive



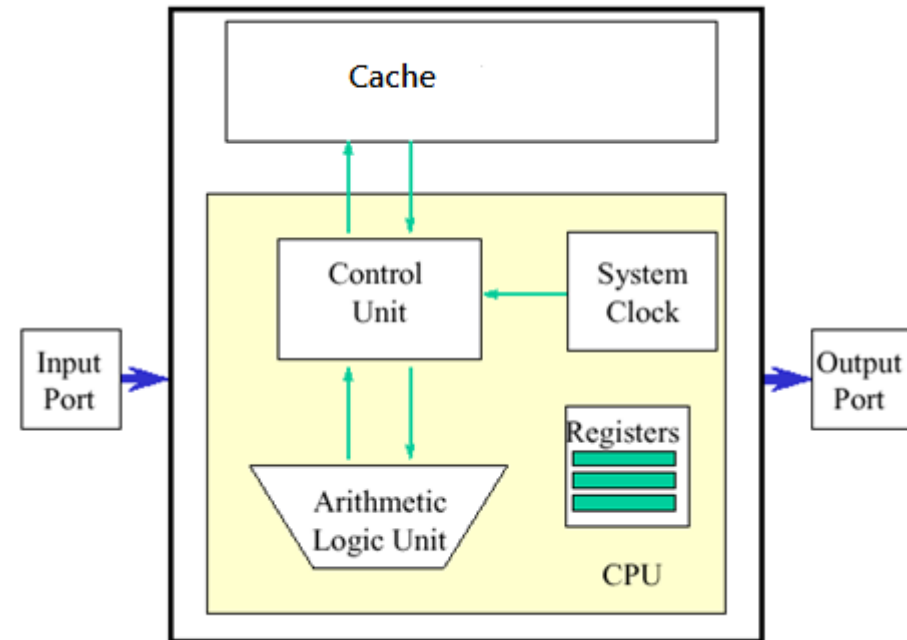


## Opening the Box – iPad 2




# Inside the Processor (CPU)

- **Datapath**: performs operations on data
- **Control**: tells datapath, memory, and I/O devices what to do according to the wishes of the instructions.
- **Cache** memory
  - Small fast SRAM memory for immediate access to data
- Nontechnical definition of **cache**: a safe place for hiding thing

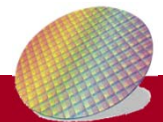


More in Chapter 3, 4 and 5



# Inside the Processor- A5

- Apple A5



# A Safe Place for Data

- **Volatile** main memory (primary memory)
  - Loses instructions and data when **power off**
- **Non-volatile** secondary memory
  - Magnetic disk
  - Flash memory
  - Optical disk (CDROM, DVD), SSD



DRAM



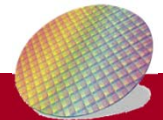
Hard Disk



Flash

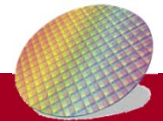


CDROM




# Networks

- Communication and resource sharing
- Local area network (LAN): Ethernet
  - Within a building
- Wide area network (WAN: the Internet)
- Wireless network: WiFi, Bluetooth, etc
- Mobile network: 3G(UMTS, CDMA2000), 3.5G (EVDO, HSDPA), 4G (LTE), 5G...




# Technology Trends

- Electronics technology continues to evolve
  - Increased capacity and performance
  - Reduced cost



DRAM capacity




Year	Technology	Relative performance/cost
1951	Vacuum tube	1
1965	Transistor	35
1975	Integrated circuit (IC)	900
1995	Very large scale IC (VLSI)	2,400,000
2013	Ultra large scale IC	250,000,000,000





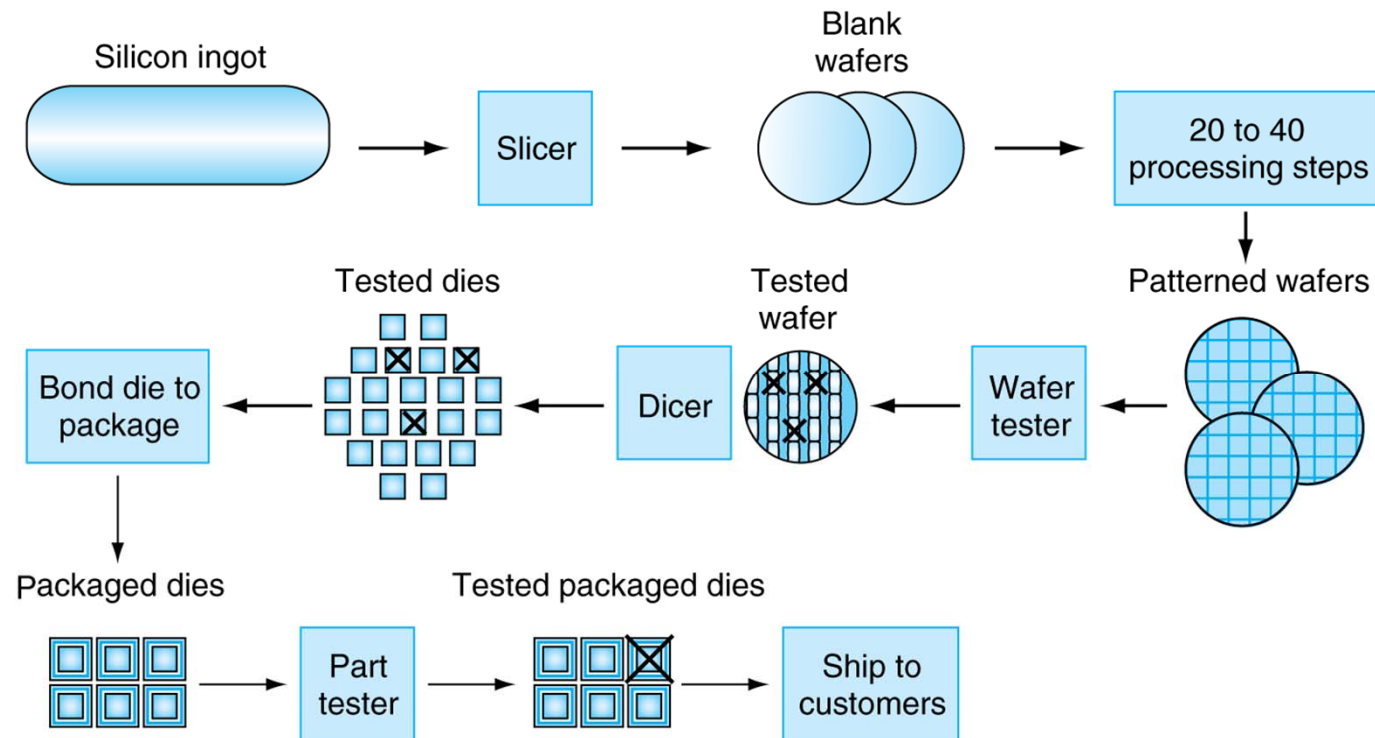
# nMOS and pMOS transistor

- nMOS transistor
  - when input is 1 ( $V_{\text{gate}}=1$ ) => conduct current
- pMOS transistor
  - when input is 1 ( $V_{\text{gate}}=0$ ), => Conduct current
- Three parts in a transistor: **conductor**, **insulator**, and **switch**
- Feature size => normally refer to **the length of the channel**







# Manufacturing ICs




- **Yield**: proportion of working dies per wafer





# Intel Core i7 Wafer




- 300mm wafer, 280 chips, 32nm technology
- Each chip is 20.7 x 10.5 mm






## Analogy to IC Manufacturing



# Integrated Circuit Cost



$$\text{Cost per die} = \frac{\text{Cost per wafer}}{\text{Dies per wafer} \times \text{Yield}}$$


$$\text{Dies per wafer} \approx \text{Wafer area} / \text{Die area}$$

$$\text{Yield} = \frac{1}{(1 + (\text{Defects per area} \times \text{Die area} / 2))^2}$$

Just an approximation

*Cost per die*  $\propto$  *area*<sup>3</sup> or *area*<sup>4</sup>

- **Wafer** cost and area are fixed
- **Defect rate** determined by manufacturing process
- **Die area** determined by **architecture** and **circuit** design



## Yield Example

- Example: Assume a wafer diameter is 15cm, dies per wafer is 90, defect per unit area is 0.018 defects/cm<sup>2</sup>, cost per wafer is 10
  - Find the yield
  - Find the cost per die

- Solutions


$$\text{Wafer area} = \pi \times (d/2)^2 = \pi \times 7.5^2 = 176.7 \text{ cm}^2$$

$$\text{Die area} = 176.7/90 = 1.96 \text{ cm}^2$$

$$\text{Yield} = ?$$

$$\text{Cost per die} = \text{cost per wafer} / (\text{dies per wafer} \times \text{yield})$$



$$\text{Cost per die} = ?$$







# Which airplane has the best performance?


成功



747



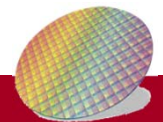
Concorde



# Response Time and Throughput

- **Response** time
  - The time between the start and completion of a task
    - e.g., 5 min to finish a task/transaction
- **Throughput**
  - Total amount of work done in a given time
    - e.g., tasks/transactions/... per hour
- We'll focus on response time for now...
- Performance of a computer X

$$Performance_X = \frac{1}{Execution\ Time_X}$$






## Relative Performance

- Define Performance = 1/Execution Time
- “X is  $n$  time faster than Y”




$$\frac{Performance_X}{Performance_Y} = n = \frac{Execution\ Time_Y}{Execution\ Time_X}$$

- Example: A takes 10s and B takes 15s to run a program, how much faster is A than B?



# Response Time

- **Response** time (or called **elapsed** time)
  - **Total** response time,
  - including all aspects
    - Processing,
    - I/O,
    - OS overhead
    - idle time
  - e.g. A task takes **5** min to finish, including all processing time, I/O, and etc.
  - It is the **time** that users will experience



# CPU Time


- CPU execution time (CPU time)
  - Time spent processing a given job (CPU may process many jobs at the same time)
  - Do not consider I/O time, other jobs' shares
  - Can be divided into user CPU time and system CPU time
    - User CPU time: CPU time spent in the program
    - System CPU time: CPU time spent in OS for this program

Quick questions: Define the following term

Response Time:

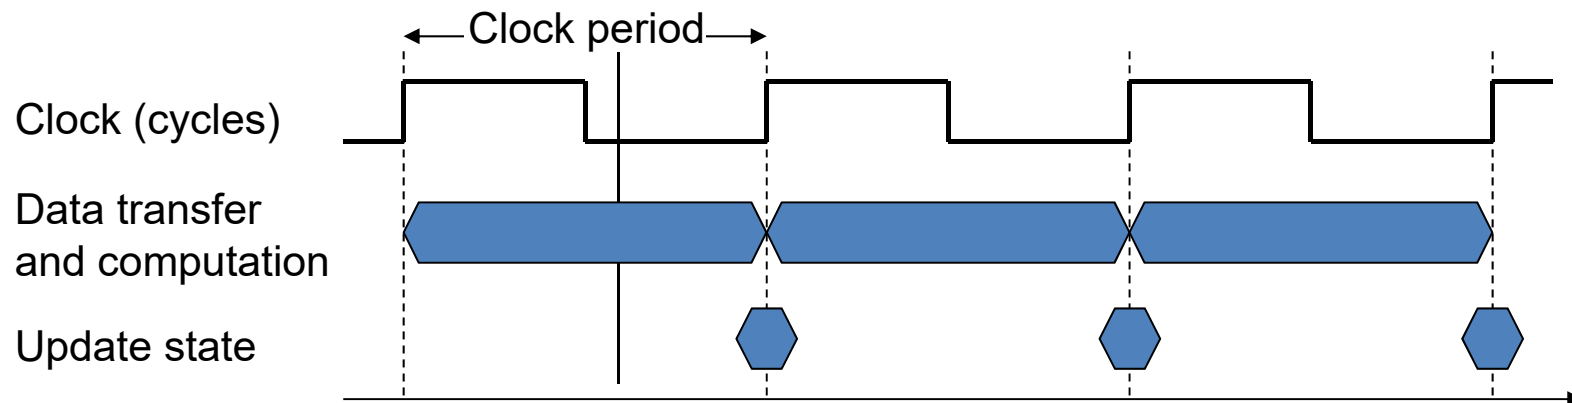
Throughput:

CPU time:




# Review: CPU Clocking

- Operation of digital hardware governed by a constant-rate clock



- Clock **period**: duration of a clock cycle
  - e.g.,  $250\text{ps} = 0.25\text{ns} = 250 \times 10^{-12}\text{s}$
- Clock **frequency** (rate): cycles per second
  - e.g.,  $4.0\text{GHz} = 4000\text{MHz} = 4.0 \times 10^9\text{Hz}$




# CPU Time

$$\begin{aligned}\text{CPU Time} &= \text{CPU Clock Cycles} \times \text{Clock Cycle Time} \\ &= \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}}\end{aligned}$$

- Performance improved by
  - Reducing number of clock cycles
  - Increasing clock rate
  - Hardware designer must often trade off clock rate against cycle count
- Example: a task takes 100 CPU cycles and each cycle is 1ns, what is the CPU Time?

$$\text{CPU Time} = 100 \times 1 = 100 \text{ ns}$$




## CPU Time Example

- Computer A: 2GHz clock, 10s CPU time. Design Computer B. It has  $1.2 \times$  clock cycles. However, CPU time is 6s because it uses faster clock. What is the clock rate of Computer B?

$$\begin{aligned}\text{Clock Cycles}_A &= \text{CPU Time}_A \times \text{Clock Rate}_A \\ &= 10\text{s} \times 2\text{GHz} = 20 \times 10^9\end{aligned}$$

$$\begin{aligned}\text{Clock Rate}_B &= \frac{\text{Clock Cycles}_B}{\text{CPU Time}_B} = \frac{1.2 \times \text{Clock Cycles}_A}{6\text{s}} \\ &= \frac{1.2 \times 20 \times 10^9}{6\text{s}} = \frac{24 \times 10^9}{6\text{s}} = 4\text{GHz}\end{aligned}$$



# Instruction Count and CPI

- CPI: Cycle per instruction
  - Number of cycles each instruction takes to execute


$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}}$$

- Instruction Count for a program determined by **program, ISA and compiler**
- **Redefine CPU time using CPI**

$$\text{CPU Time} = \text{Clock Cycles} \times \text{Clock Cycle Time}$$

$$= \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

$$= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$$





## CPI Example

- Computer A: Cycle Time = 250ps, CPI = 2.0
- Computer B: Cycle Time = 500ps, CPI = 1.2
- Same ISA (same number of instructions)
- Which is faster, and by how much?


$$\begin{aligned}\text{CPU Time}_A &= \text{Instruction Count} \times \text{CPI}_A \times \text{Cycle Time}_A \\ &= 1 \times 2.0 \times 250\text{ps} = 1 \times 500\text{ps}\end{aligned}$$

A is faster...

$$\begin{aligned}\text{CPU Time}_B &= \text{Instruction Count} \times \text{CPI}_B \times \text{Cycle Time}_B \\ &= 1 \times 1.2 \times 500\text{ps} = 1 \times 600\text{ps}\end{aligned}$$

$$\frac{\text{CPU Time}_B}{\text{CPU Time}_A} = \frac{1 \times 600\text{ps}}{1 \times 500\text{ps}} = 1.2$$

1.2 times faster



## CPI in More Detail


- If different **instruction** classes take different numbers of **cycles**

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

### ■ **Weighted** average CPI

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left( \text{CPI}_i \times \frac{\text{Instruction Count}_i}{\text{Instruction Count}} \right)$$

Relative frequency



## CPI Example

- Compiled code sequences using instructions in classes A, B, C. CPI\_A=1, CPI\_B=2, CPI\_C=3, find average **CPI for instruction sequence 1 and 2**


Instruction Class	A	B	C
Inst. Count in sequence 1	2	1	2
Inst. Count in sequence 2	4	1	1

- Sequence 1: IC = **5**

- Clock Cycles  
 $= 2 \times 1 + 1 \times 2 + 2 \times 3$   
 $= 10$
- Avg. CPI =  $10/5 = 2.0$

- Sequence 2: IC = **6**

- Clock Cycles  
 $= 4 \times 1 + 1 \times 2 + 1 \times 3$   
 $= 9$
- Avg. CPI =  $9/6 = 1.5$




# Performance Summary

## The BIG Picture


$$\text{CPU Time} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}$$

- Performance depends on
  - Algorithm: affects IC, possibly CPI
  - Programming language: affects IC, CPI
  - Compiler: affects IC, CPI
  - Instruction set architecture: affects IC, CPI,  $T_c$



# Power Trends

- CPU clock rate and power has increased rapidly for decades




- In CMOS IC technology, energy of single transition

$$\text{Energy} = \frac{1}{2} CV^2$$

- Power of single transition

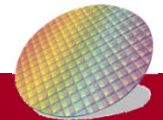
$$\text{Power} = \frac{1}{2} CV^2 f$$

Frequency increase 1000x  
Voltage decrease from 5V-> 1V,  
Power increase 30x




## Reducing Power

- Suppose a new CPU has 85% of capacitive load of old CPU, and its voltage and frequency is reduced by 15%
  - What is  $\frac{\text{Power}_{\text{new}}}{\text{Power}_{\text{old}}}$  ?
- The power wall
  - Hard to reduce voltage further
  - Hard to remove more heat
- How else can we improve performance?




# Uniprocessor Performance



Constrained by power, instruction-level parallelism, memory latency

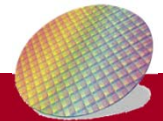
The Sea Change: The Switch to Multiprocessors





# Multiprocessors


- **Multicore** microprocessors
  - **More** than one **processor** per chip
- Requires explicitly **parallel programming**
  - Compared to Instruction level parallelism
    - Hardware executes multiple instructions at once
    - Hidden from the programmer
  - Why parallel programming is hard to do?
    - Programming for both **correctness** and **performance**
    - Load balancing
    - Optimizing communication and synchronization



# SPEC CPU Benchmark

- **Benchmark:** Programs used to measure performance
  - Supposedly typical of actual workload
- Standard Performance Evaluation Corp (SPEC)
  - Develops benchmarks for CPU, I/O, Web, ...
- SPEC CPU2006
  - **Time** to execute a selection of programs
    - Negligible I/O, so focuses on CPU performance
  - Normalize relative to reference machine
  - Summarize as geometric mean of performance ratios
    - CINT2006 (integer) and CFP2006 (floating-point)

$$\sqrt[n]{\prod_{i=1}^n \text{Execution time ratio}_i}$$



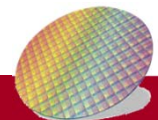
# CINT2006 for Opteron X4 2356

**Ref time:** Time required by a reference computer

Name	Description	IC×10 <sup>9</sup>	CPI	Tc (ns)	Exec time	Ref time	SPEC ratio
perl	Interpreted string processing	2,118	0.75	0.40	637	9,777	15.3
bzip2	Block-sorting compression	2,389	0.85	0.40	817	9,650	11.8
gcc	GNU C Compiler	1,050	1.72	0.47	24	8,050	11.1
mcf	Combinatorial optimization	336	10.00	0.40	1,345	9,120	6.8
go	Go game (AI)	1,658	1.09	0.40	721	10,490	14.6
hmmer	Search gene sequence	2,783	0.80	0.40	890	9,330	10.5
sjeng	Chess game (AI)	2,176	0.96	0.48	37	12,100	14.5
libquantum	Quantum computer simulation	1,623	1.61	0.40	1,047	20,720	19.8
h264avc	Video compression	3,102	0.80	0.40	993	22,130	22.3
omnetpp	Discrete event simulation	587	2.94	0.40	690	6,250	9.1
astar	Games/path finding	1,082	1.79	0.40	773	7,020	9.1
xalancbmk	XML parsing	1,058	2.70	0.40	1,143	6,900	6.0
Geometric mean							11.7

15.3 times faster

$$\text{CPU Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$




# SPECRatio of a benchmark suite

Benchmarks	Ultra 5 (sec)	Opteron (Sec)	SPECRatio
Wupwise	1600	51.5	31.06
Swim	3100	125.0	24.73
Mgrid	1800	98.0	18.37
Applu	2100	94.0	22.34
Mesa	1400	64.6	21.69
Galgel	2900	86.4	33.57
Art	2600	92.4	28.13
Equake	1300	72.6	17.92
Facerec	1900	73.6	25.80
Amp	2200	136.0	16.14
Lucas	2000	88.8	22.52
Fma3d	2100	120.0	17.48
sixtrack	1100	123.0	8.95
apsi	2600	150.0	17.36
G. M.			20.886

If a benchmark suite contains many benchmarks, the overall SPECRatio is the **geometric mean** of each SPECRatio

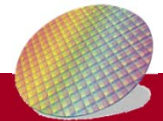
$$\begin{aligned} SPECRatio_{overall} &= \sqrt[n]{\prod_{i=1}^n SPECRatio_i} \\ &= \sqrt[14]{31.06 \times 24.73 \dots \times 17.36} \\ &= 20.86 \end{aligned}$$



# SPEC Power Benchmark


- Power consumption of server at different workload levels
  - Performance: ssj\_ops/sec
  - Power: Watts (Joules/sec)

$$\text{Overall ssj\_ops per Watt} = \left( \sum_{i=0}^{10} \text{ssj\_ops}_i \right) / \left( \sum_{i=0}^{10} \text{power}_i \right)$$



## SPECpower\_ss2008 for X4


Target Load %	Performance (ssj_ops/sec)	Average Power (Watts)
100%	231,867	295
90%	211,282	286
80%	185,803	275
70%	163,427	265
60%	140,160	256
50%	118,324	246
40%	92,035	233
30%	70,500	222
20%	47,126	206
10%	23,066	180
0%	0	141
Overall sum	1,283,590	2,605
$\sum \text{ssj\_ops} / \sum \text{power}$		493



# Pitfall

- Improving an aspect of a computer and expecting a **proportional** improvement in **overall** performance
  - Example: **multiplication** accounts for 80s of 100s.  
Improve **multiplication** by **5x** and expect overall performance to improve **5x**, which is **20s**
  - => **Wrong**
- **Amdahl's Law**: Possible improvement is limited by the amount that the improved feature is used

$$T_{\text{improved}} = \frac{T_{\text{affected}}}{\text{improvement factor}} + T_{\text{unaffected}}$$





# Amdahl's Law

- Amdahl's Law: Possible improved is limited by the amount that the improved feature is used


$$T_{\text{improved}} = \frac{T_{\text{affected}}}{\text{improvement factor}} + T_{\text{unaffected}}$$

- Example: **multiplication** accounts for 80s of 100s. How much improvement in **multiplication** performance to get **4x** overall?

$$25 = \frac{80}{n} + 20$$

$$n = 16$$

**multiplication** needs to be **16x** faster




# Pitfall: use MIPS as a Performance Metric

- MIPS: Millions of Instructions Per Second

$$\begin{aligned} \text{MIPS} &= \frac{\text{Instruction count}}{\text{Execution time} \times 10^6} \\ &= \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}} \times 10^6 = \frac{\text{Clock rate}}{\text{CPI} \times 10^6} \end{aligned}$$

- Question: A and B has different instruction set. A runs 5 MIPS, and B run 4 MIPS, which one is faster?

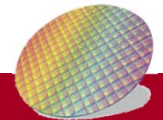
We don't know because we don't know how much time each instruction takes



# Pitfall: MIPS as a Performance Metric

---

- MIPS: Millions of Instructions Per Second
  - Doesn't account for
    - Differences in **ISAs** between computers
    - Differences in **complexity** between instructions
  - Note: CPI varies between programs on a given CPU, between various instruction sets
- See example next slide



# MIPS Example

Measurement	Computer A	Computer B
Instr. Count	10 billion	8 billion
Clock rate	4 GHz	4GHz
CPI	1.0	1.1

$$\text{MIPS} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}$$

- Which computer has **higher MIPS**?

$$\text{MIPS}_A = 4 / (1.0 \times 10^6)$$

$$\text{MIPS}_B = 4 / (1.1 \times 10^6)$$


A has **larger** MIPS

- Which is **faster**?

$$\text{CPU time}_A = 10 \times 10^9 \times 1.0 / 4 \times 10^9 = 2.5$$

$$\text{CPU time}_B = 8 \times 10^9 \times 1.1 / 4 \times 10^9 = 2.2$$

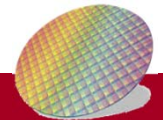
B is **faster**



## Concluding Remarks

- **Cost/performance** is improving
  - Due to underlying technology development
- **Hierarchical** layers of abstraction
  - In both hardware and software
- **Instruction set** architecture
  - The hardware/software interface
- **Execution** time: the best performance measure
- **Power** is a limiting factor
  - Use parallelism to improve performance

## Read Chapter 1





成功大學

National Cheng Kung University

# Backup slides

