




成功大學

National Cheng Kung University

# Chapter 3 Part 2

## Arithmetic for Computers

### -Floating Point



成功大

- ```
float a; // single precision
double b; //double precision
```

Can't use  
integer to  
represent



# Floating Point Standard- IEEE Std 754-1985

- Single precision - 32-bit

single: 8 bits

single: 23 bits


Significand=1  
+fraction

| S | Exponent | Fraction |
|---|----------|----------|
|---|----------|----------|

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

$$x = (-1)^S \times (\text{Significand}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0  $\Rightarrow$  non-negative, 1  $\Rightarrow$  negative)
- Normalized number**  $\pm 1.xxxxxxx_2 \times 2^{yyyy}$ 
  - Always has a leading **1**, so no need to represent it explicitly (**hidden** bit)
- Exponent**: excess representation: actual exponent + **Bias**
  - Ensures exponent is unsigned
  - Single precision: Bias = **127**, Double precision: Bias = **1023**



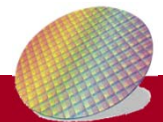


## Floating-Point Example – single-precision

What number is represented by the following **single-precision** float?

$x = 11000000101000...00_2$  (32-bit, single precision)


- $S = 1$
- Fraction =  $01000...00_2$  –
- Exponent =  $10000001_2 = 129$
- $x = (-1)^1 \times (1 + .01_2) \times 2^{(129 - 127)}$   
 $= (-1) \times (1 + 1/4) \times 2^2$   
 $= -5.0$



## Floating-Point Example

- Represent  $-0.75$  in **single**-precision floating point
  - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
  - $S = ?$
  - Fraction = ?
  - Exponent = ?

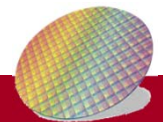
Hidden 1 is not represented



## Floating-Point Example

- Represent  $-0.75$  in **single**-precision floating point
  - $-0.75 = -(1/4 + 1/2) = (-1)^1 \times 1.1_2 \times 2^{-1}$
  - $S = 1$
  - Fraction = **1000...00**      **Hidden 1** is not represented
  - Exponent =  $-1 + \text{Bias} = 126 = 01111110_2$

Answer: **1011111101000...00**




Represent  $3.4375 \times 10^{-1}$  in **single**-precision floating point

$$\begin{aligned} 3.4375 \times 10^{-1} &= 0.3475 \\ &= 0.0101100 = 1.0110000000 \times 2^{-2} \end{aligned}$$

$$\begin{aligned} 0.34375 * 2 &= 0.6875 \dots (0) \\ 0.8750 * 2 &= 1.375 \dots (1) \\ 0.375 * 2 &= 0.75 \dots (0) \\ 0.75 * 2 &= 1.5 \dots (1) \\ 0.5 * 2 &= 1 \end{aligned}$$

- $S = 0$
- Fraction = 011000...00
- Exponent =  $-2 + \text{Bias } (127) = 125 = 01111101_2$

Answer: 001111101011000...00




# Why uses bias (excess presentation) in the exponents



- Easier to compare which exponent is larger
  - Just need to check the bit from left to right

| 8 bits |          | Bias=127 |          |          |
|--------|----------|----------|----------|----------|
| 127    | 01111111 | 254      | 11111110 |          |
| 126    | 01111110 | 253      | 11111101 |          |
|        |          | ....     |          |          |
| .....  |          | .....    |          |          |
| 1      | 00000001 | 128      | 10000000 |          |
| 0      | 00000000 | 127      | 01111111 |          |
| -1     | 11111111 | ....     |          |          |
| ....   |          |          |          |          |
| -126   | 10000010 | 1        | 00000001 |          |
| -127   | 10000001 | 0        | 00000000 | reserved |
| -128   | 10000000 | 255      | 11111111 | reserved |








# Floating Point Standard- IEEE Std 754-1985

- Double precision (64-bit)

double: 11 bits


double: 52 bits



$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

$$x = (-1)^S \times (\text{Significand}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0  $\Rightarrow$  non-negative, 1  $\Rightarrow$  negative)
- Normalized number  $\pm 1.xxxxxxx_2 \times 2^{yyyy}$ 
  - Have hidden 1  $\text{Fraction} = \text{Significand} - 1$
- Exponent: excess representation: actual exponent + Bias
  - Ensures exponent is unsigned
  - Double: Bias = 1023






## Floating-Point Example – double-precision

- What number is represented by the following double float?

$x = 1011111111011000...00_2$  (64-bit)

- $S = 1$
  - Fraction =  $1000...00_2$
  - Exponent =  $01111111101_2$
- $x = (-1)^1 \times (1 + .1_2) \times 2^{(1021 - 1023)}$   
 $= (-1) \times (1 + 1/2) \times 2^{-2}$   
 $= -3/8$



## Floating-Point Example

- Represent  $-0.75$  in **double**-precision floating point

- $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$

$$0.75 * 2 = 1.5 \text{ ..... } 1$$


$$0.5 * 2 = 1.0 \text{ ..... } 1$$

- $S = ?$

- Fraction = ?

Hidden 1 is not represented

- Exponent = ?



## Floating-Point Example

- Represent  $-0.75$  in **double**-precision floating point

- $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$

$$0.75 * 2 = 1.5 \text{ ..... } 1$$

$$0.5 * 2 = 1.0 \text{ ..... } 1$$


- $S = 1$

$$0.75_{10} = 0.11_2$$

- Fraction =  $1000...00_2$  Hidden 1 is not represented

- Exponent =  $-1 + \text{Bias} = -1 + 1023 = 1022_{10} = 0111111110_2$

Ans:  $1011111111101000...00$



# Floating point - Half-precision

- Half precision

5 bits

10 bits

| S | Exponent | Fraction |
|---|----------|----------|
|---|----------|----------|


$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- Bias = 15

Represent -0.75 in half-precision floating point

- $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
- $S = 1$
- Fraction =  $1000...00_2$
- Exponent =  $-1 + 15 = 14 = 01110_2$


Ans:  $101110\ 1000000000_2$



# IEEE 754 Encoding of FP number

- Exp.=0 and Fract.=0  $\Rightarrow$  0
- Exp.=0 and Fract.  $\neq$  0  $\Rightarrow$  denormalized number (discuss later)
- Exp.=111..111 and Fract.= 0  $\Rightarrow \pm\infty$  (discuss later)
- Exp.=111...111 and Fract. $\neq$ 0  $\Rightarrow$  Non a Number (NaN) (discuss later)

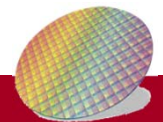
| Single precision |          | Double precision |          | Object represented          |
|------------------|----------|------------------|----------|-----------------------------|
| Exponent         | Fraction | Exponent         | Fraction |                             |
| 0                | 0        | 0                | 0        | 0                           |
| 0                | Nonzero  | 0                | Nonzero  | $\pm$ denormalized number   |
| 1–254            | Anything | 1–2046           | Anything | $\pm$ floating-point number |
| 255              | 0        | 2047             | 0        | $\pm$ infinity              |
| 255              | Nonzero  | 2047             | Nonzero  | NaN (Not a Number)          |



# Denormalized Numbers

- (Review) Smallest normalized value
  - 00000001 00000000.....0000
  - Fraction: 000...00  $\Rightarrow$  significand = 1.0
  - Exponent = 1 – 127 = –126
  - Smallest value =  $1.0 \times 2^{-126}$
- How to represent number smaller than  $1.0 \times 2^{-126}$ ?
- E.g.  $0.5 \times 2^{-126} \Rightarrow$  Use denormalized number

|   |          |          |
|---|----------|----------|
| S | Exponent | Fraction |
|---|----------|----------|



## Denormalized Numbers (32-bit)

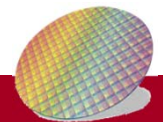
- Exponent = 00000000
- Fraction  $\Rightarrow$  **hidden** bit is 0 (**not 1**)

$$x = (-1)^s \times (\text{Fraction}) \times 2^{-126}$$

$0.5 \times 2^{-126}$  : Exponent = **0 00000000**  
: Fraction = 10000000000000...000

$$0.5 \times 2^{-126} = \mathbf{0\ 00000000\ 10000000000000...000}$$

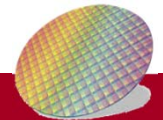
- Allow for gradual **underflow**, with diminishing precision
- Denormalized with fraction = 000...0





## Special number: Infinities and NaNs

- Exponent =  $111\dots 1$ , Fraction =  $000\dots 0$ 
  - $\pm\infty$
  - Can be used in subsequent calculations, avoiding need for overflow check
  - E.g.  $F+(+\infty)=+\infty$  , or  $F/\infty=0$
- Exponent =  $111\dots 1$ , Fraction  $\neq 000\dots 0$ 
  - Not-a-Number (NaN)
  - Indicates illegal or undefined result
    - e.g.,  $0.0 / 0.0$
  - Can be used in subsequent calculations




## Example

- Smallest **positive single** precision normalized number

$$1.000000000...00000_2 \times 2^{-126}$$

| S | Exp        | Fraction                     |
|---|------------|------------------------------|
| - | -----      | -----                        |
| 0 | 0000 0001  | 0000 0000 0000 0000 0000 000 |
|   | $2^{-126}$ |                              |

- Smallest **positive single** precision denormalized no.  
(Hint: Fraction is 23-bit)



## Example

- Smallest **positive single** precision normalized number


$$1.000000000...00000_2 \times 2^{-126}$$

| S | Exp        | Fraction                     |
|---|------------|------------------------------|
| - | -----      | -----                        |
| 0 | 0000 0001  | 0000 0000 0000 0000 0000 000 |
|   | $2^{-126}$ | $\times 1.0$                 |

- Smallest **positive single** precision denormalized no.

(Hint: Fraction is 23-bit)

| S | Exp        | Fraction                     |
|---|------------|------------------------------|
| - | -----      | -----                        |
| 0 | 0000 0000  | 0000 0000 0000 0000 0000 001 |
|   | $2^{-126}$ | $\times 2^{-23}$             |



# Floating-Point Addition

- Consider a 4-digit decimal example

$$- 9.999 \times 10^1 + 1.610 \times 10^{-1}$$

- 1. **Align** decimal points
  - Shift number with smaller exponent

$$9.999 \times 10^1 + 0.016 \times 10^1$$

- 2. Add significands


$$9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$$

- 3. Normalize result & check for over/underflow

$$1.0015 \times 10^2$$

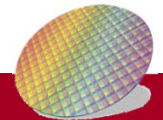
- 4. Round and **renormalize** if necessary

$$1.002 \times 10^2$$



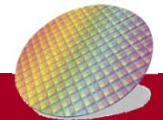
# Floating-Point Addition

- Now consider a 4-digit binary example
  - $1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$  ( $0.5 + -0.4375$ )
- 1. **Align** binary points
  - Shift number with smaller exponent
$$1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$$
- 2. Add significands
$$1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$$
- 3. Normalize result & check for over/underflow
$$1.000_2 \times 2^{-4}, \text{ with no over/underflow}$$
- 4. Round and renormalize if necessary
$$1.000_2 \times 2^{-4} \text{ (no change)} = 0.0625$$



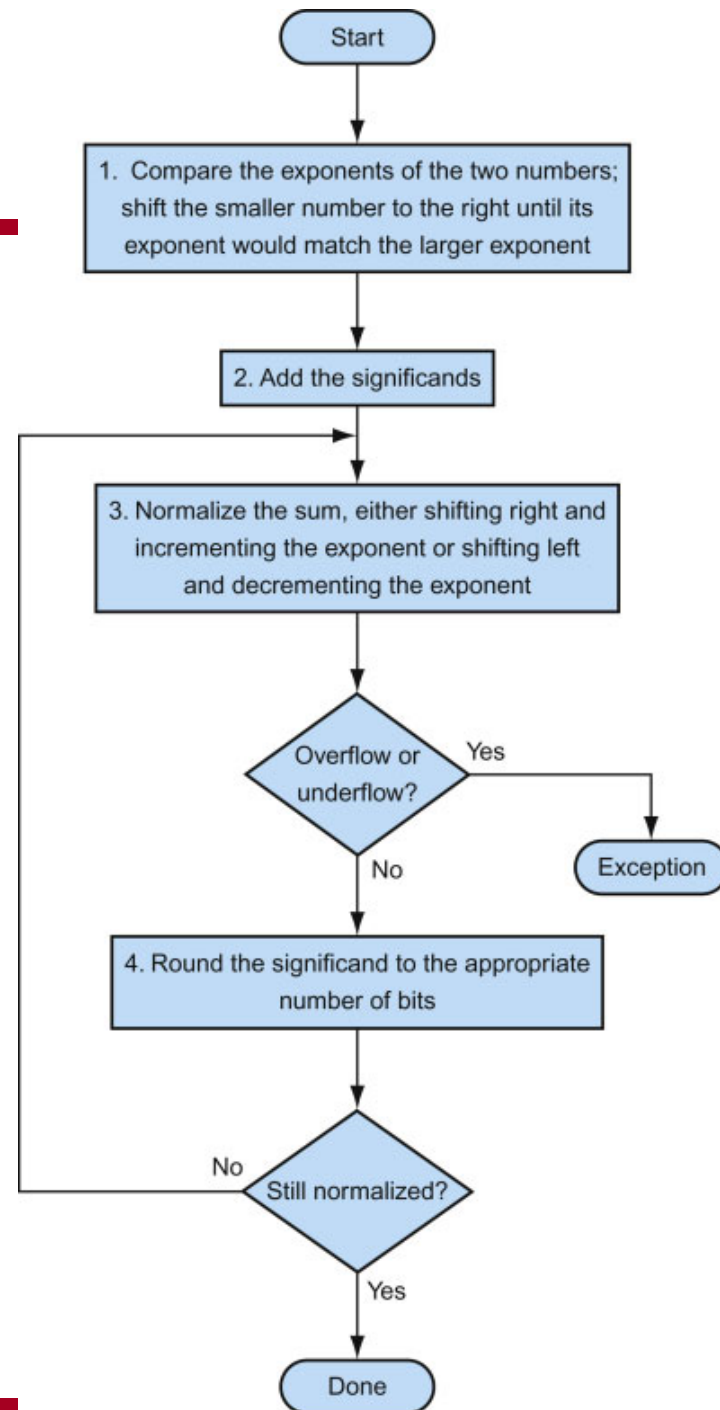
## FP Adder Hardware

- Much more complex than integer adder
  - Steps includes **shift** exponents and fraction, add fraction, ..., etc.
- Doing it in one clock cycle would take **too long**
  - Much longer than integer operations
  - Slower clock would penalize all instructions
- FP adder usually takes **several cycles**
  - Can be pipelined (see Chapter 4 about pipeline)



# FP addition flow

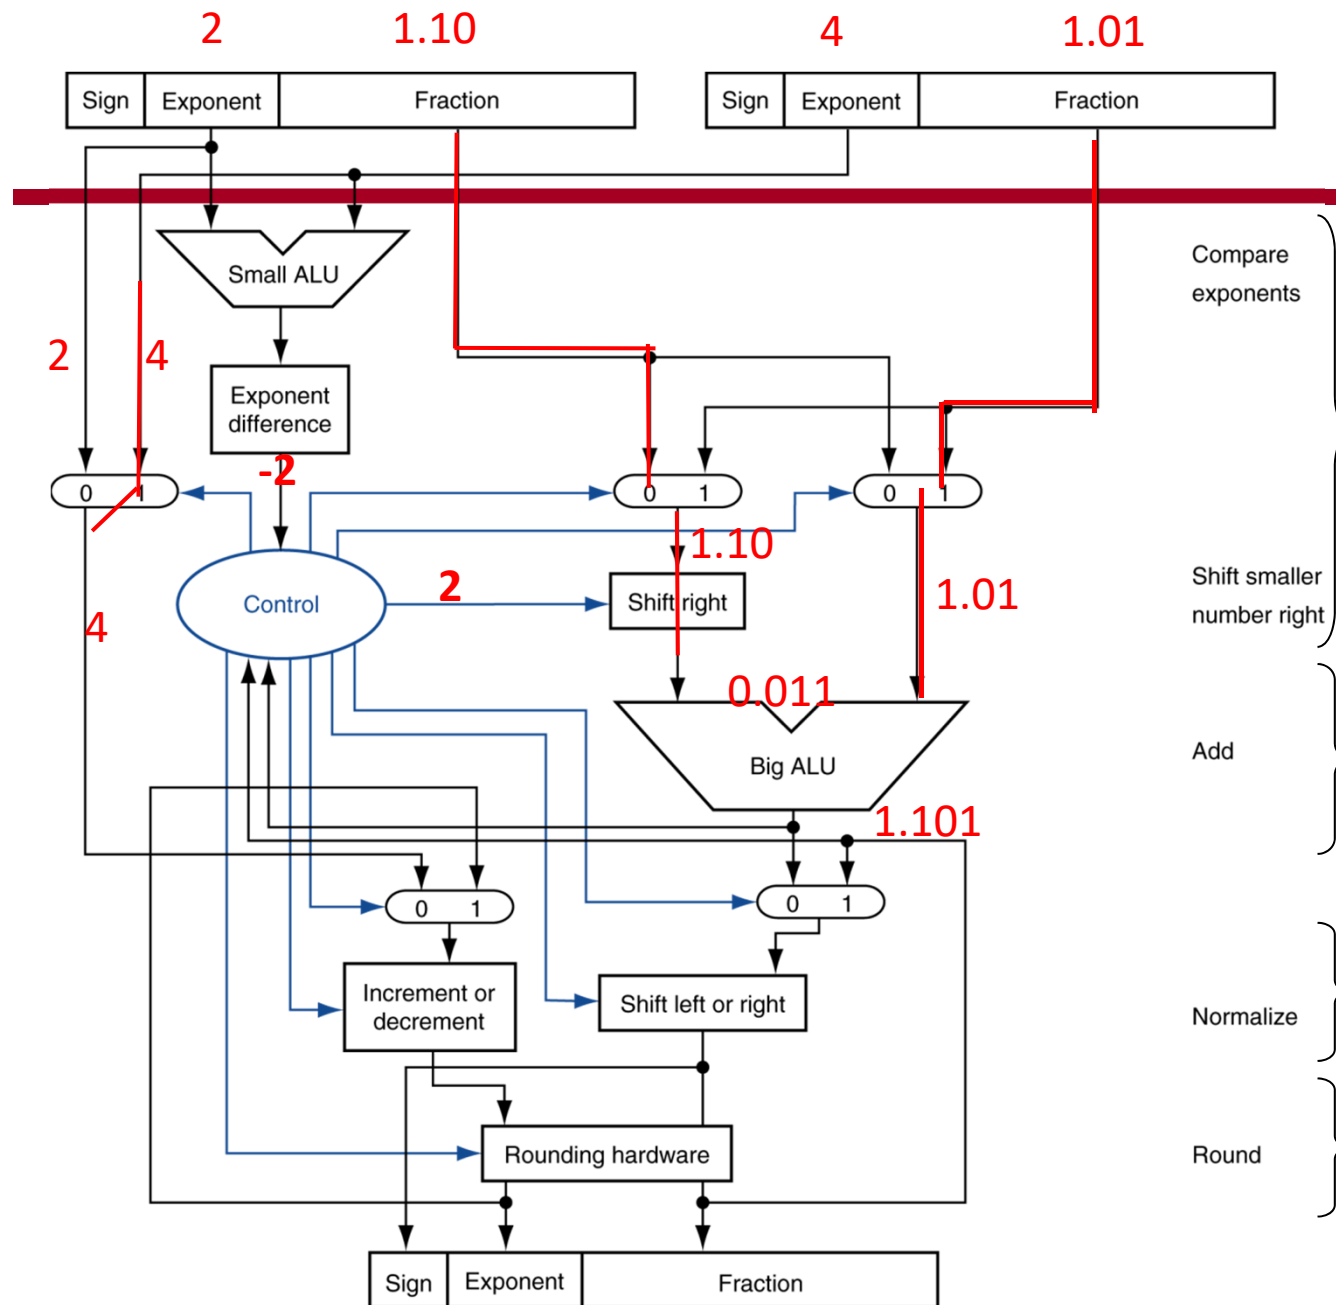
Floating-point Addition. The normal path is to execute steps 3 and 4 once, but if rounding causes the sum to be unnormalized, we must repeat step 3.



# FP Adder Hardware



成功大學



Compare exponents

Step 1: Align binary points

Shift smaller number right

Add

Step 2: Add significands

Normalize

Step 3: Normalize result & check for over/underflow

Round

Step 4: Round and renormalize if necessary

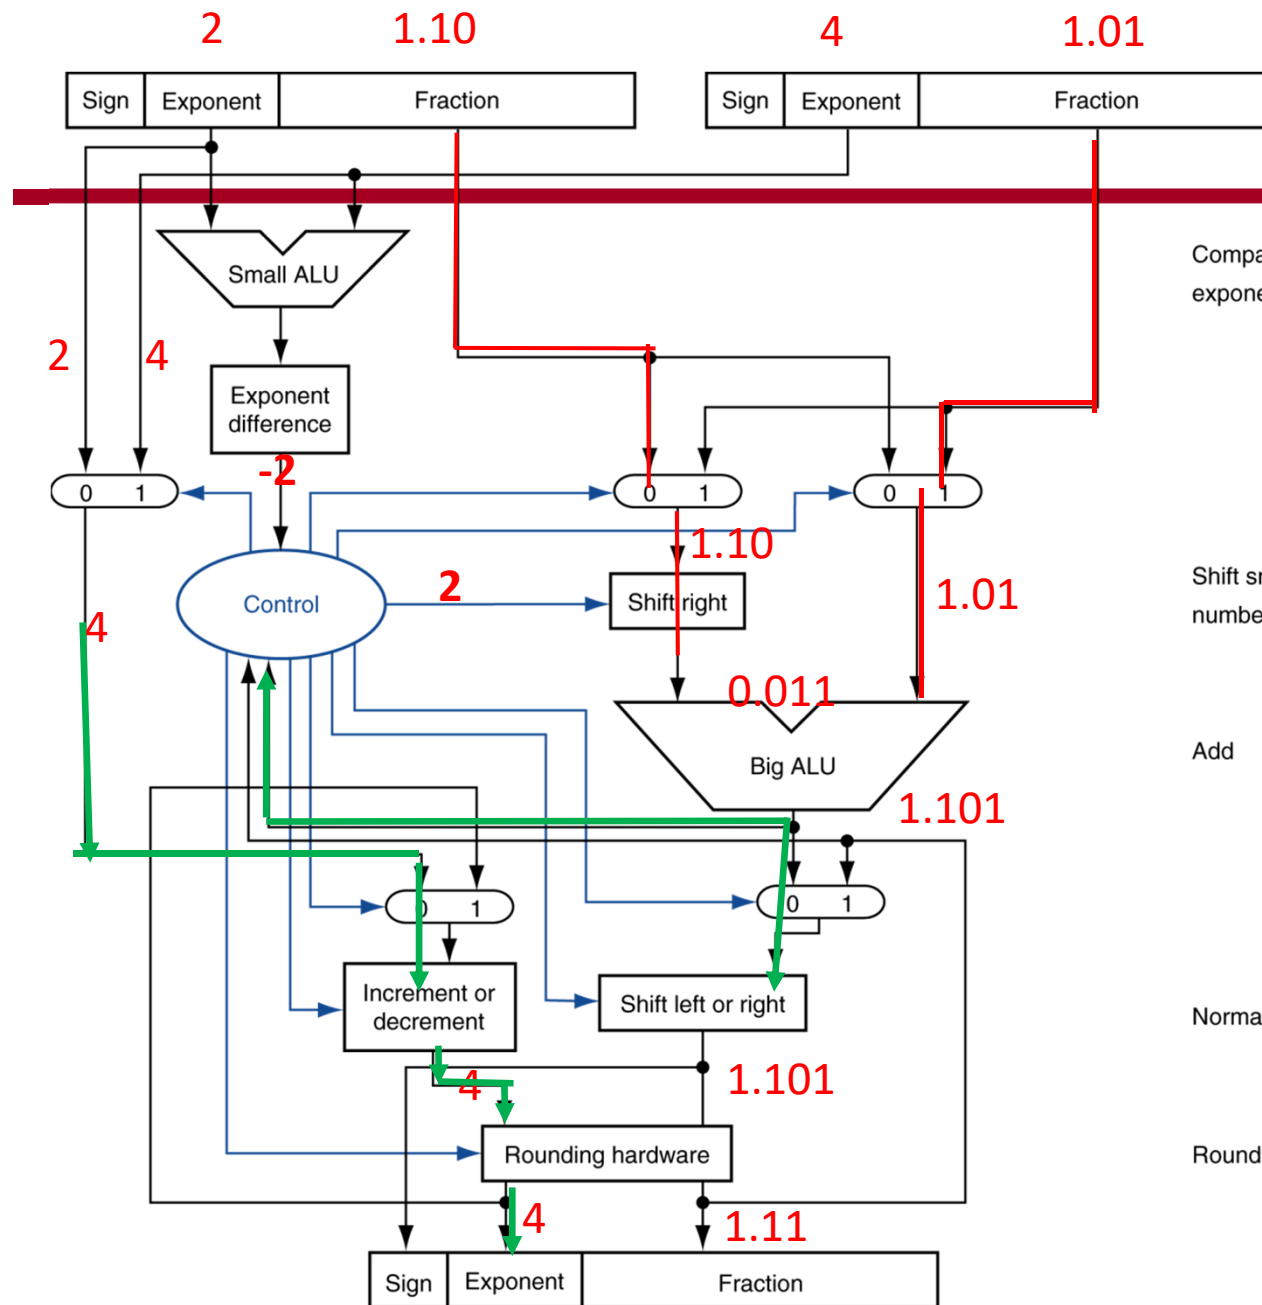




# FP Adder Hardware



成功大學



Compare exponents

Step 1: Align binary points

Shift smaller number right

Add


Step 2: Add significands

Normalize

Step 3: Normalize result & check for over/underflow

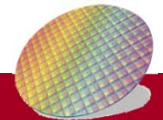
Round

Step 4: Round and renormalize if necessary



# Floating-Point Multiplication


- Consider a 4-digit decimal example
  - $1.110 \times 10^{10} \times 9.200 \times 10^{-5}$
- 1. Add exponents
  - For biased exponents, subtract **bias** from sum
  - New exponent =  $10 + -5 = 5$
- 2. Multiply significands
  - $1.110 \times 9.200 = 10.212 \Rightarrow 10.212 \times 10^5$
- 3. Normalize result & check for over/underflow
  - $1.0212 \times 10^6$
- 4. Round and renormalize if necessary
  - $1.021 \times 10^6$
- 5. Determine sign of result from signs of operands
  - $+1.021 \times 10^6$



# Floating-Point Multiplication

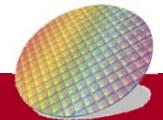
- Now consider a 4-digit binary example  
 $1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$  ( $0.5 \times -0.4375$ )
- 1. Add exponents
  - Unbiased:  $-1 + -2 = -3$
  - Biased:  $(-1 + 127) + (-2 + 127) - 127 = -3 + 254 - 127$
- 2. Multiply significands
  - $1.000_2 \times 1.110_2 = 1.110_2 \Rightarrow 1.110_2 \times 2^{-3}$
- 3. Normalize result & check for over/underflow
  - $1.110_2 \times 2^{-3}$  (no change) with no over/underflow
- 4. Round and renormalize if necessary
  - $1.110_2 \times 2^{-3}$  (no change)
- 5. Determine sign:  $+ve \times -ve \Rightarrow -ve$ 
  - $-1.110_2 \times 2^{-3} = -0.21875$

Remove one bias




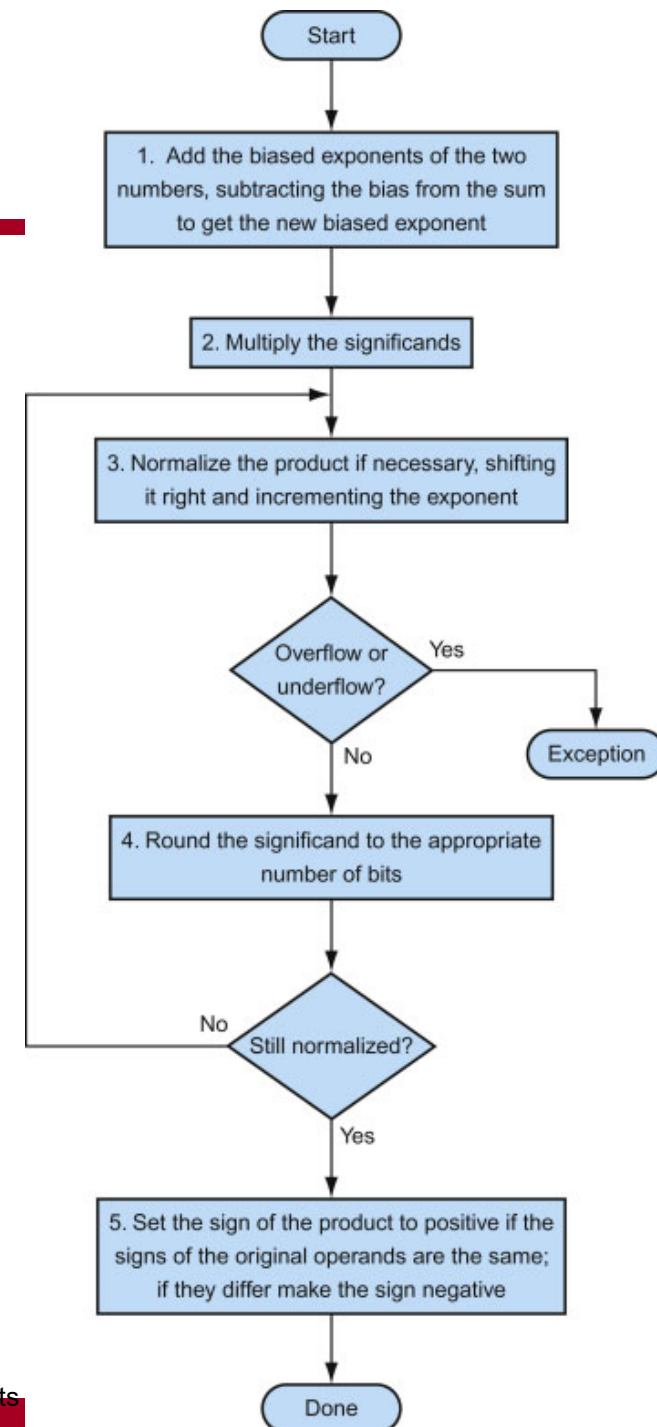
# FP Arithmetic Hardware

- FP multiplier is of similar complexity to **FP** adder
  - But do **multiplication** for significands instead of an **addition**
- FP arithmetic hardware usually does
  - **Addition, subtraction, multiplication, division, reciprocal, square-root**
  - **FP  $\leftrightarrow$  integer conversion**
- Operations usually takes **several cycles**
  - Can be pipelined (See Chapter 4)



# FP Multiplication

The normal path is to execute steps 3 and 4 once, but if rounding causes the sum to be unnormalized, we must repeat step 3.



# Improve Accuracy

- IEEE Std 754 specifies additional rounding control
  - Extra bits of precision (**guard, round, sticky**)
- **Guard & round** bits: two **extra** (hidden) bits on the right during intermediate additions
  - Improve precision

Consider the addition  $2.56 \times 10^0 + 2.34 \times 10^2 = 2.3656$


Without **guard** and **round** bit

$$0.02 \times 10^2 + 2.34 \times 10^2 = 2.36 \times 10^2$$

With **guard** and **round** bit

$$0.0256 \times 10^2 + 2.3400 \times 10^2 = 2.3656 \times 10^2 = 2.37 \times 10^2$$

↑  
closer to accurate answer




## Improve Accuracy: sticky bit

- **Sticky bit**: one **bit** is set when there are nonzero bits to the right of the round bit.
  - Allow computer to see the difference between  $0.50000..0_{10}$  and  $0.50000..1_{10}$
- Without Sticky bit

$2.34500000000001$  will be stored as  $2.345$
- With Sticky bit

$2.34500000000001$  will be stored as  $2.345$  and sticky bit = 1
- Used for rounding

$2.345$  with sticky bit = 1 is larger than  $2.345$



## Rounding: Round to nearest even

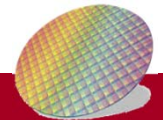
**GRS** are three bits are only used while doing calculations and aren't stored in the floating-point variable **before** or **after** the calculations.

### GRS - Action

**0xx** - round down = do nothing (x means any bit value, 0 or 1)

**100** - this is a **tie**: round up if the fraction's bit just before **G** is 1, else round down(=do nothing)

|                       | GRS                                       |
|-----------------------|-------------------------------------------|
| <b>101</b> - round up | <b>01 100</b> $\Rightarrow$ <b>10 000</b> |
| <b>110</b> - round up | <b>00 100</b> $\Rightarrow$ <b>00 000</b> |
| <b>111</b> - round up |                                           |





3.29 Calculate the sum of  $2.6125 \times 10^1$  and  $4.150390625 \times 10^{-1}$  by hand, assuming both are stored in the 16-bit half precision. Assume 1 guard, 1 round bit, and 1 sticky bit and round to the nearest even.

$$\text{3.29 } 2.6125 \times 10^1 + 4.150390625 \times 10^{-1}$$

$$2.6125 \times 10^1 = 26.125 = 11010.001 = 1.1010001000 \times 2^4$$

$$4.150390625 \times 10^{-1} = .4150390625 = .011010100111 = 1.1010100111 \times 2^{-2}$$

For the second number, shift binary point 6 to the left to align exponents,

$$1.1010100111 \times 2^{-2} = 0.0000011010100111 \times 2^{-4}$$

GR

$$1.1010001000 \ 00$$


$$0.0000011010 \ 10 \ 0111 \text{ (Guard bit = 1, Round bit = 0, Sticky bit = 1)}$$

-----

$$1.1010100010 \ 101 \text{ (Guard bit = 1, Round bit = 0, Sticky bit = 1)}$$

In this case, the extra bit (G,R,S) is more than half of the least significant bit (0). Thus, the value is rounded up.

$$1.1010100011 \times 2^4 = 11010.100011 \times 2^0 = 26.546875 = 2.6546875 \times 10^1$$



## Fallacy: Right Shift and Division

- Left shift by  $i$  places multiplies an integer by  $2^i$  and thus right shift divides by  $2^i$

Correct for unsigned number, incorrect for signed number


- For signed number, this is correct

$$00001011_2 \gg 2 = 00000010_2 \quad (11/4=2)$$

- For signed integers, this is **incorrect**

– e.g.,  $-5 / 4 = -1 \dots -1$

$$11111011_2 \gg 2 = 00111110_2 = 62 \text{ not } -1$$




## Pitfall: FP addition is not associative

- Is  $(x+y)+z$  equal to  $x+(y+z)$  ??? **NO**

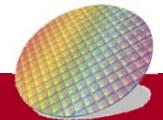
|   |           | $(x+y)+z$ | $x+(y+z)$ |
|---|-----------|-----------|-----------|
| x | -1.50E+38 | 0.00E+00  | -1.50E+38 |
| y | 1.50E+38  |           | 1.50E+38  |
| z | 1.0       | 1.0       |           |
|   |           | 1.00E+00  | 0.00E+00  |

- Parallel Programs may interleave operations in unexpected orders
  - Assumptions of **associativity may fail**
- Need to validate parallel programs under varying degrees of parallelism



## Concluding Remarks

- Bits have no inherent meaning
  - Interpretation depends on the instructions applied
- Computer representations of numbers
  - Finite range and precision
- ISAs support arithmetic
  - Signed and unsigned integers
  - Floating-point approximation to reals
- Bounded range and precision
  - Operations can overflow and underflow





成功大學

National Cheng Kung University

## Backup slides

